

Computer Vision




Edge Detection

Dr. Mrinmoy Ghorai

Indian Institute of Information Technology
Sri City, Chittoor





Today's Agenda: Edge Detection



Previous classes: Image Filtering

Spatial domain

- Smoothing, sharpening, measuring texture



Frequency domain

- Denoising, sampling

Today's class

- Detecting edges





- Finding straight lines




Why finding edges is important?

- Cues for 3D shape
- Group pixels into objects or parts
- Shape analysis
- Recover geometry and viewpoint



Origin of Edges




Edges are caused by a variety of factors


Closeup of edges




Closeup of edges



Closeup of edges



Closeup of edges



Closeup of edges




Closeup of edges



Characterizing edges


- An edge is a place of rapid change in the image intensity function

image




Characterizing edges

- An edge is a place of rapid change in the image intensity function




Characterizing edges

- An edge is a place of rapid change in the image intensity function




Characterizing edges



- An edge is a place of rapid change in the image intensity function






Intensity profile





Intensity profile



Intensity profile




With a little Gaussian noise





Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



Effects of noise



- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



Where is the edge?

Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal





Where is the edge?

The larger the noise the stronger the response


Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal




- What can we do about it?


Solution: smooth first



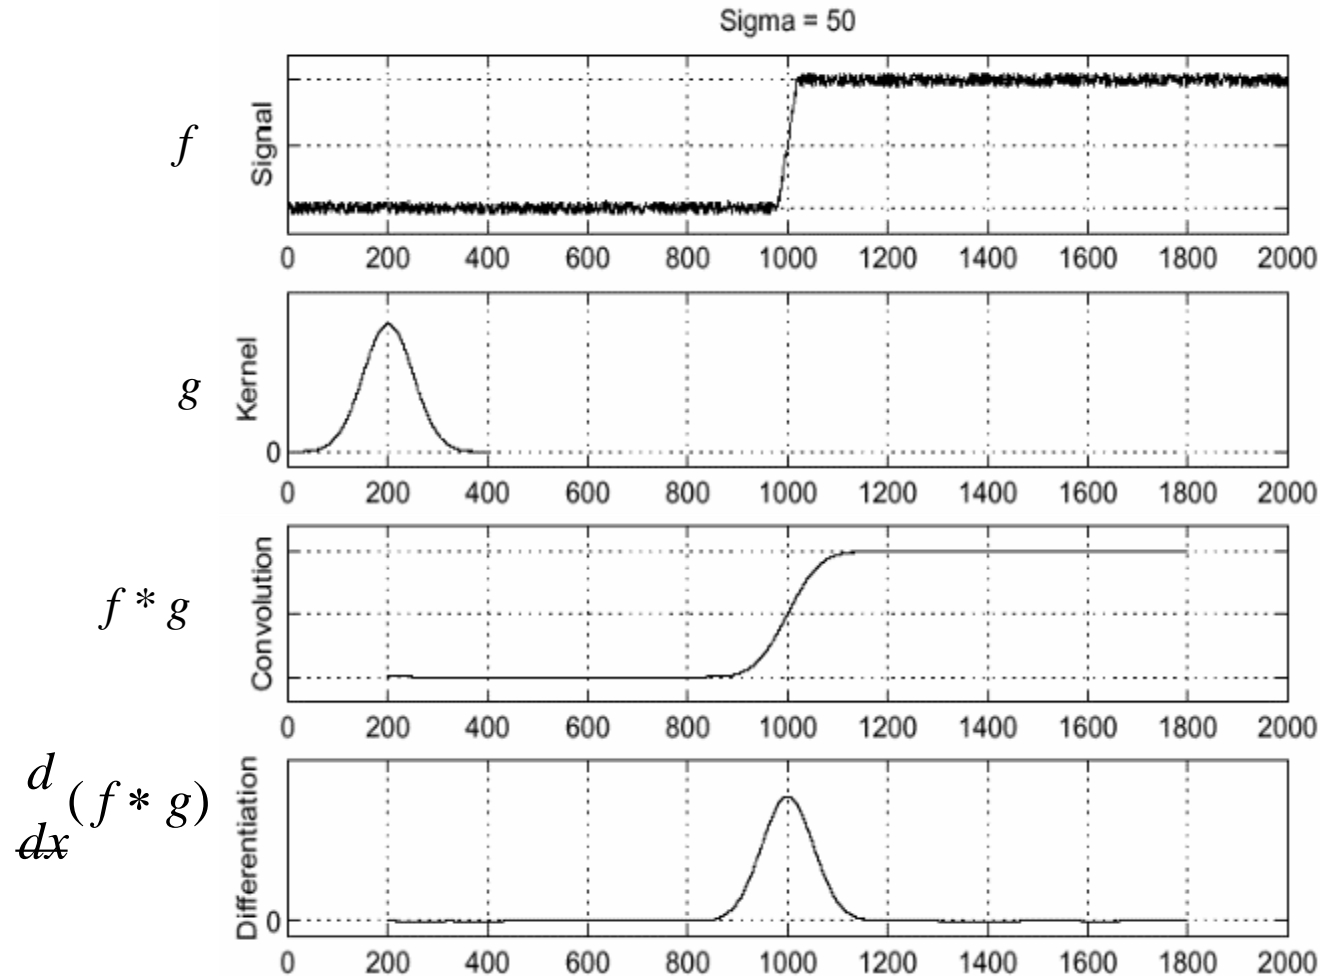
Solution: smooth first



Solution: smooth first



Solution: smooth first




- To find edges, look for peaks in $\frac{d}{dx}(f * g)$

Derivative theorem of convolution



- Differentiation is convolution, and convolution is associative:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

- This saves us one operation:






Gaussian Kernel




$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$


Gaussian filters




$\sigma = 1$ pixel



$\sigma = 5$ pixels




$\sigma = 10$ pixels




$\sigma = 30$ pixels

Derivative of Gaussian filter



* [1 0 -1] =




derivative of Gaussian (x)

Gaussian



$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$


Derivative of Gaussian filter



x-direction



y-direction



Designing an edge detector

- Criteria for a good edge detector:
 - **Good detection:**
 - find all real edges, ignoring noise or other artifacts

Designing an edge detector

- Criteria for a good edge detector:
 - **Good detection:**
 - find all real edges, ignoring noise or other artifacts
 - **Good localization**
 - detect edges as close as possible to the true edges
 - return one point only for each true edge point

Canny edge detector

- The most widely used edge detector

A computational approach to edge detection

[J Canny](#) - IEEE Transactions on pattern analysis and machine ..., 1986 - ieeexplore.ieee.org

Abstract: This paper describes a computational approach to edge detection. The success of the approach depends on the definition of a comprehensive set of goals for the computation of edge points. These goals must be precise enough to delimit the desired behavior of the detector while making minimal assumptions about the form of the solution. We define detection and localization criteria for a class of edges, and present mathematical forms for ...

[Cited by 27743](#) [Related articles](#) [All 27 versions](#) [Import into BibTeX](#) [Save](#) [More](#)


J. Canny, [***A Computational Approach To Edge Detection***](#), IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

Example





input image (“Lena”)


Derivative of Gaussian filter




x-direction



y-direction



Compute Gradients (DoG)





Input Image



X-Derivative of Gaussian

Compute Gradients (DoG)



Input Image





X-Derivative of Gaussian



Y-Derivative of Gaussian

Compute Gradients (DoG)




Input Image



X-Derivative of Gaussian



Y-Derivative of Gaussian



Gradient Magnitude


Get Orientation at Each Pixel

- Get orientation




$$\text{theta} = \text{atan2}(\text{gy}, \text{gx})$$

Non-maximum suppression for each orientation






At q , we have a maximum if the value is larger than those at both p and at r .
Interpolate to get these values.



Sidebar: Interpolation options

- 'nearest'
 - Copy value from nearest known
 - Very fast but creates blocky edges
- 'bilinear'
 - Weighted average from four nearest known pixels
 - Fast and reasonable results
- 'bicubic' (default)
 - Non-linear smoothing over larger area
 - Slower, visually appealing, may create negative pixel values



Before Non-max Suppression



After non-max suppression




Hysteresis thresholding

- Threshold at low/high levels to get weak/strong edge pixels
- Do connected components, starting from strong edge pixels



Hysteresis thresholding

- Check that maximum value of gradient value is sufficiently large
 - drop-outs? use **hysteresis**
 - use a high threshold to start edge curves and a low threshold to continue them.



Final Canny Edges



Canny edge detector

1. Filter image with x, y derivatives of Gaussian

Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient

Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
 - Thin multi-pixel wide “ridges” down to single pixel width

Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
 - Thin multi-pixel wide “ridges” down to single pixel width
4. Thresholding and linking (hysteresis):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them

Canny edge detector

1. Filter image with x, y derivatives of Gaussian
 2. Find magnitude and orientation of gradient
 3. Non-maximum suppression:
 - Thin multi-pixel wide “ridges” down to single pixel width
 4. Thresholding and linking (hysteresis):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
- **MATLAB:** `edge (image, 'canny')`

Effect of σ (Gaussian kernel spread/size)



original



Canny with $\sigma = 1$





Canny with $\sigma = 2$

The choice of σ depends on desired behavior

- large σ detects large scale edges
- small σ detects fine features

Why edges?



Reduce dimensionality of data

Preserve content information


Useful in applications such as:

- object detection

- structure from motion

- tracking

Why **not** edges?




But, not that useful, **why?**

Difficulties:


1. Modeling assumptions
2. Parameters
3. Multiple sources of information
(brightness, color, texture, ...)
4. Real world conditions

Is edge detection even well defined?


Canny edge detection




1. smooth



2. gradient




3. thresh, suppress, link





Canny is optimal w.r.t. some model.



Canny edge detection




1. smooth



2. gradient




3. thresh, suppress, link



And yet...

Canny difficulties

1. Modeling assumptions
Step edges, junctions, etc.
2. Parameters
Scales, threshold, etc.
3. Multiple sources of information
Only handles brightness
4. Real world conditions
Gaussian iid noise? Texture...




Learning to detect boundaries



image



human segmentation



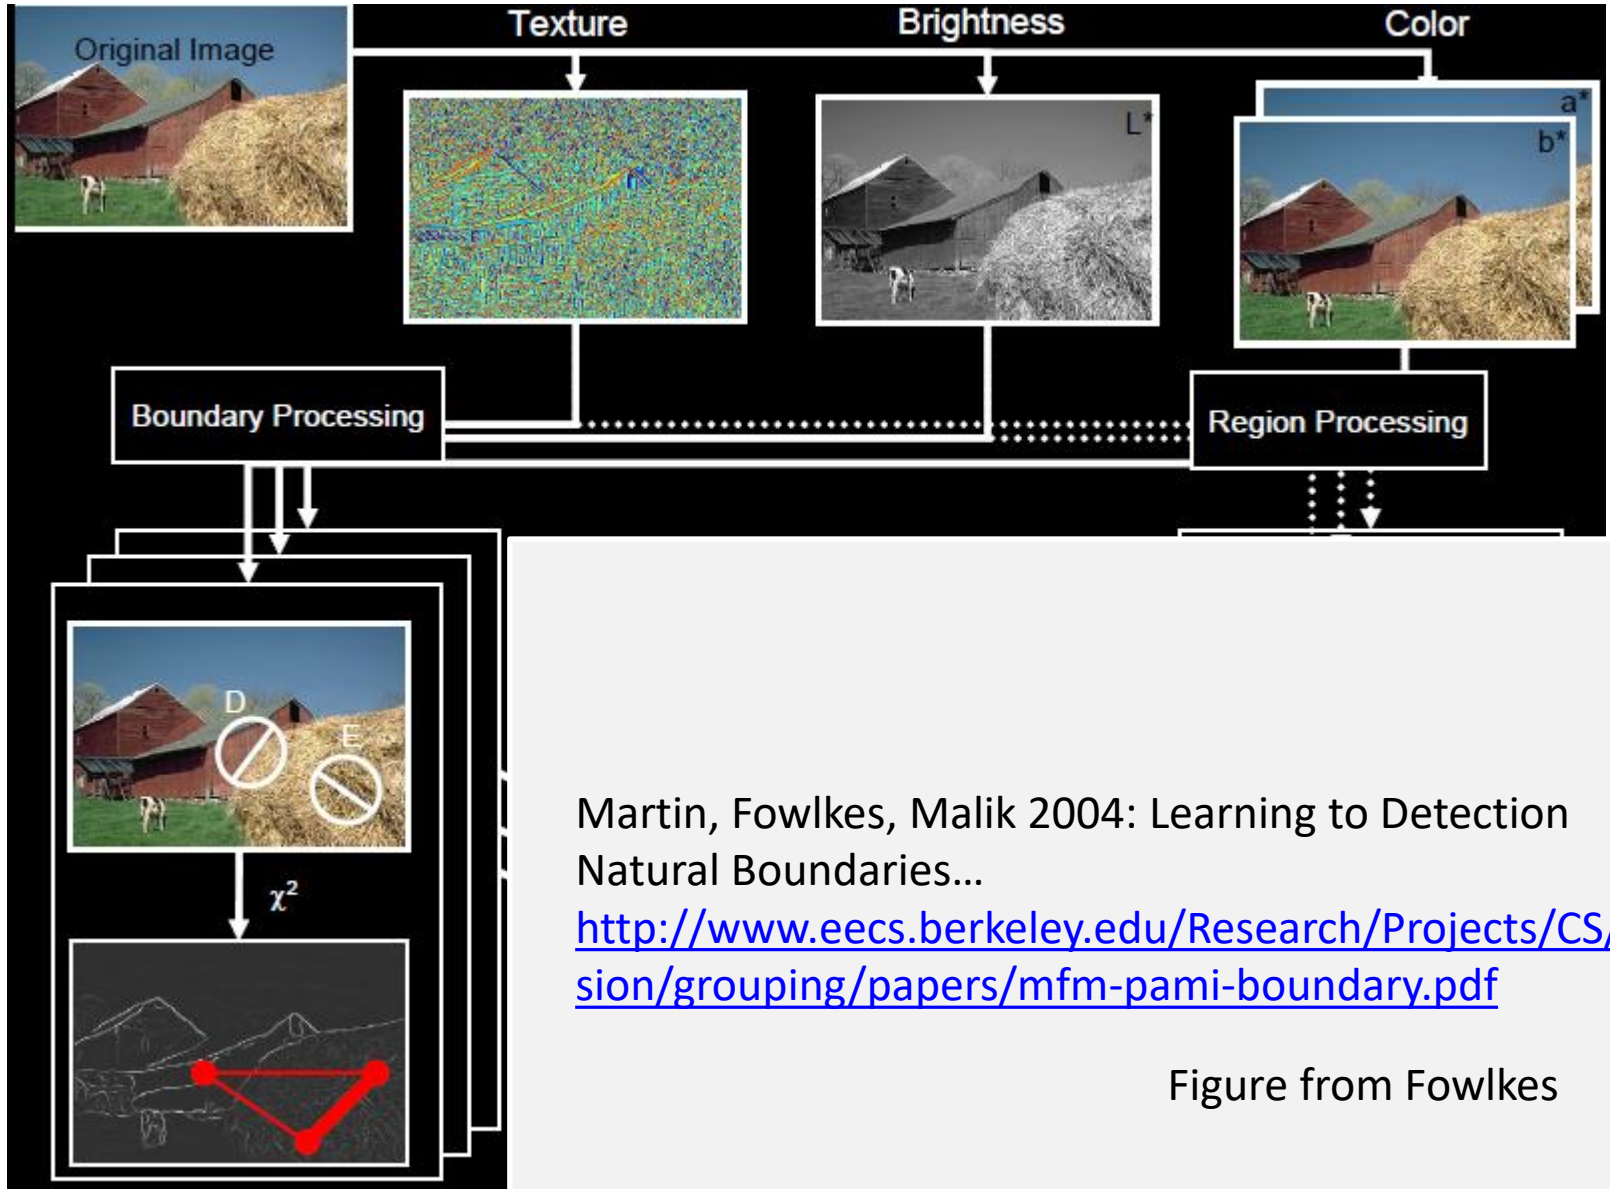
gradient magnitude



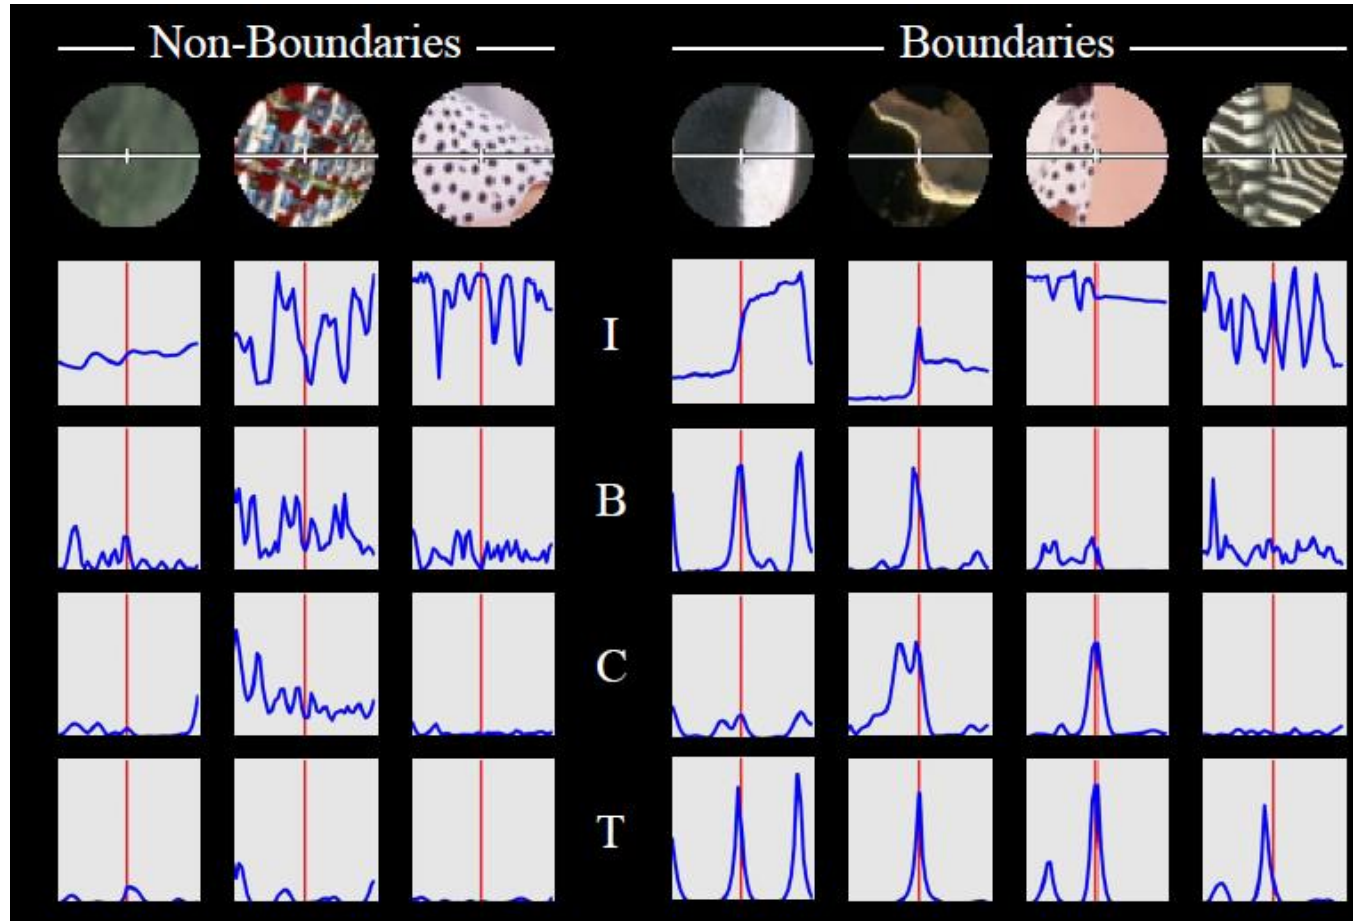
Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

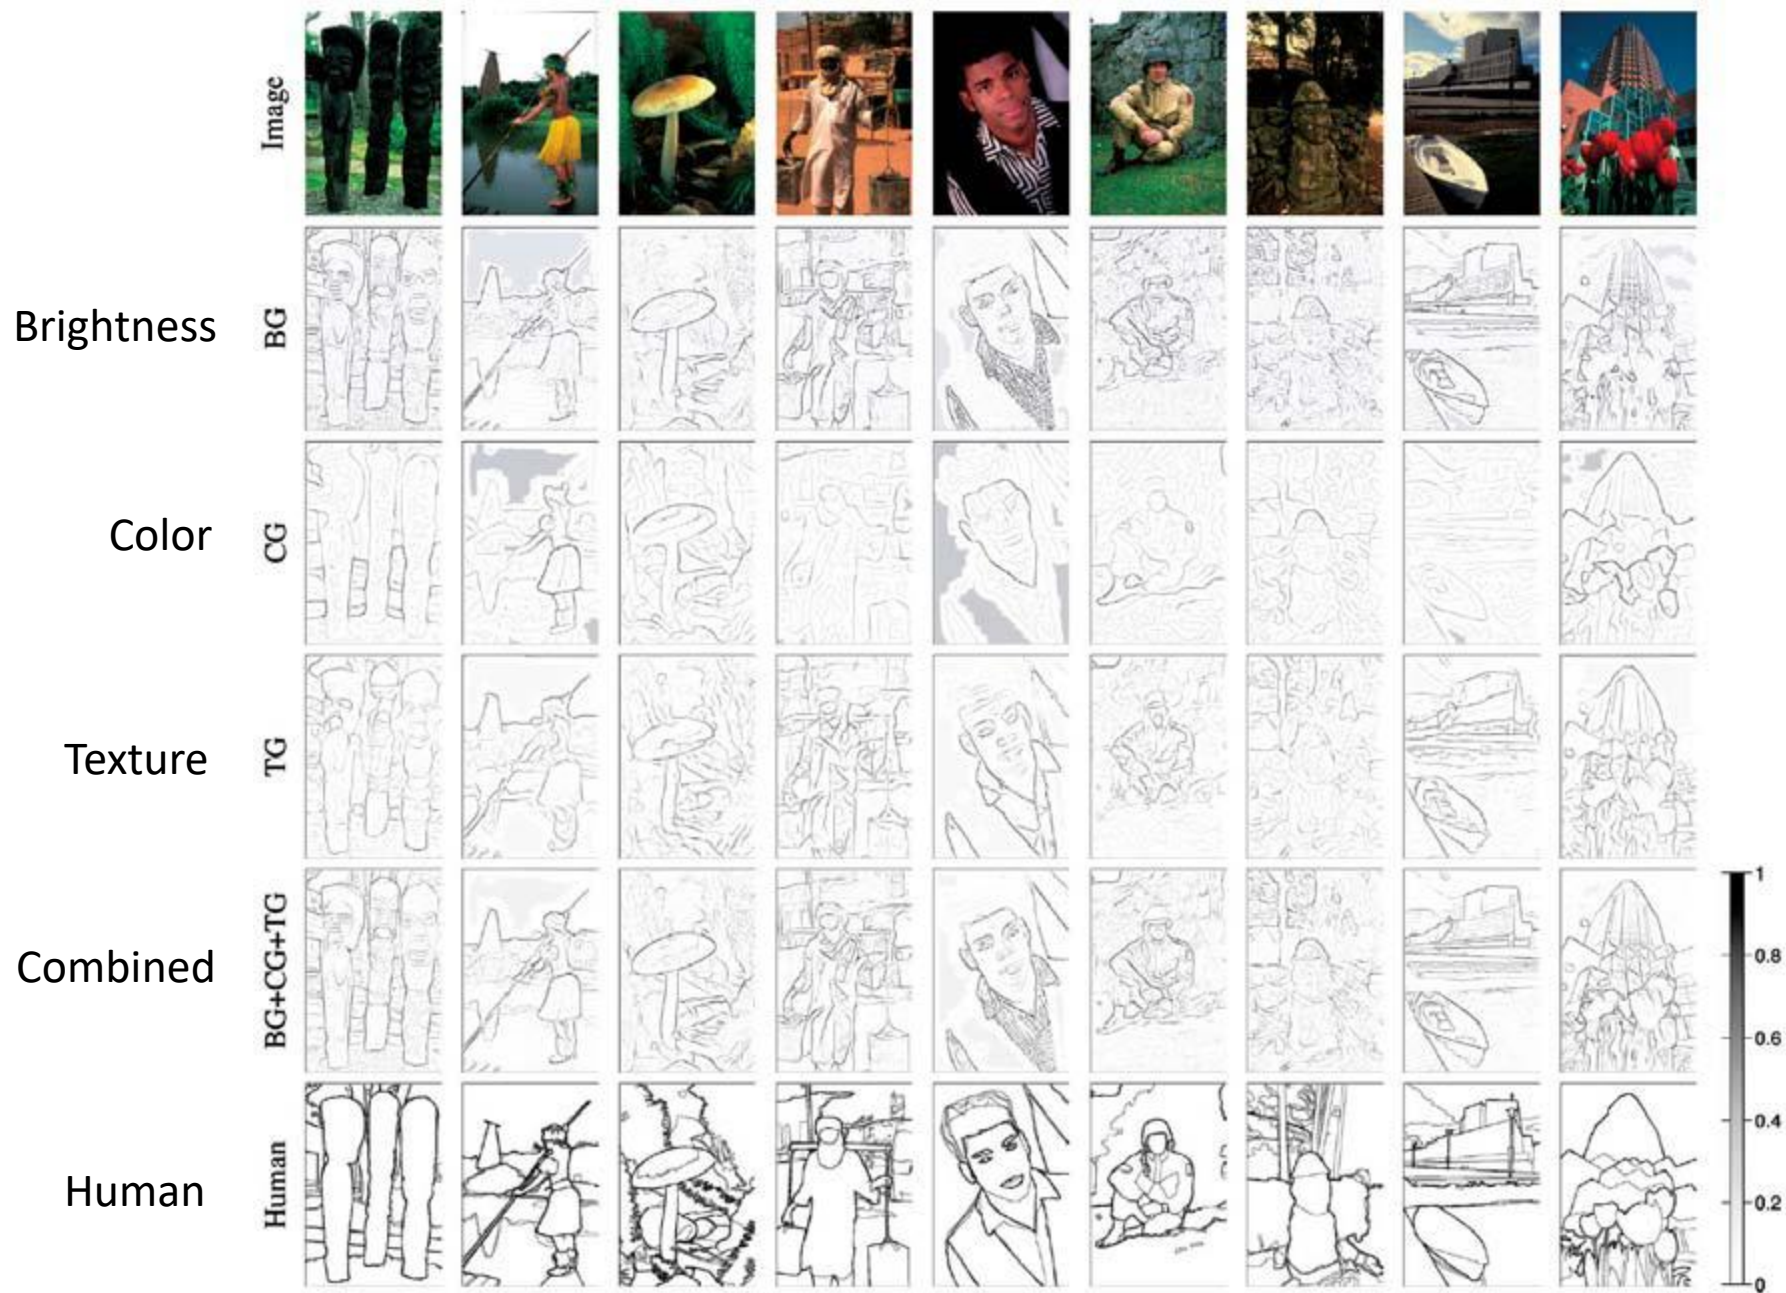
pB boundary detector



pB Boundary Detector



- Estimate Posterior probability of boundary passing through centre point based on local patch based features
- Using a Supervised Learning based framework



Features

Brightness oriented energy,

$$\text{OE}_{\theta,\sigma} = (I * f_{\theta,\sigma}^e)^2 + (I * f_{\theta,\sigma}^o)^2$$


Gaussian second derivative

Gradients computed from two disc halves:

Brightness gradient

Color gradient

Texture gradient




Texture features



(a)

Filterbank (13 filters)




(b)

Universal textons (64)



(c)

Image



(d)

Texton map (color-coded)

Martin, Fowlkes, Malik, 2004: Berkeley (Pb) edge detector

Localization


- edges (due to large filters) are poorly localized; double peaks
- Improve Localization by using derived feature
- Divide by distance to nearest maximum

$$\hat{f}(x) = \tilde{f}(x) \cdot \left(\frac{-f''(x)}{|f'(x)| + \epsilon} \right)$$



where $f(x)$ is feature and the estimated distance to the nearest maximum of $f(x)$ is



$$d(x) = -|f'(x)|/f''(x)$$


Results



Results







For more:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/bsds/bench/html/108082-color.html>

State of edge detection

- Local edge detection is mostly solved
 - Intensity gradient, color, texture
- Often used in combination with object detectors or region classifiers
- Deep learning approach is more common nowadays

Finding straight lines



Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x,y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$

Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x,y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$
2. Assign each edge to one of 8 directions

Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x,y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$
2. Assign each edge to one of 8 directions
3. For each direction d , get edgelets:
 - find connected components for edge pixels with directions in $\{d-1, d, d+1\}$

Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x,y directions)
 - Compute: $\theta = \arctan(g_y / g_x)$
2. Assign each edge to one of 8 directions
3. For each direction d , get edgelets:
 - find connected components for edge pixels with directions in $\{d-1, d, d+1\}$
4. Compute straightness and θ of edgelets using eig of x,y 2^{nd} moment matrix of their points

Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x,y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$
2. Assign each edge to one of 8 directions
3. For each direction d , get edgelets:
 - find connected components for edge pixels with directions in $\{d-1, d, d+1\}$
4. Compute straightness and theta of edgelets using eig of x,y 2nd moment matrix of their points

$$\mathbf{M} = \begin{bmatrix} \sum (x - \mu_x)^2 & \sum (x - \mu_x)(y - \mu_y) \\ \sum (x - \mu_x)(y - \mu_y) & \sum (y - \mu_y)^2 \end{bmatrix} \quad [v, \lambda] = \text{eig}(\mathbf{M})$$

Larger eigenvector
↓

$$\theta = \text{atan2}(v(2,2), v(1,2))$$
$$\text{conf} = \lambda_2 / \lambda_1$$

Finding line segments using connected components

1. Compute canny edges
 - Compute: g_x, g_y (DoG in x,y directions)
 - Compute: $\theta = \text{atan}(g_y / g_x)$
2. Assign each edge to one of 8 directions
3. For each direction d , get edgelets:
 - find connected components for edge pixels with directions in $\{d-1, d, d+1\}$
4. Compute straightness and theta of edgelets using eig of x,y 2nd moment matrix of their points

$$\mathbf{M} = \begin{bmatrix} \sum (x - \mu_x)^2 & \sum (x - \mu_x)(y - \mu_y) \\ \sum (x - \mu_x)(y - \mu_y) & \sum (y - \mu_y)^2 \end{bmatrix} \quad [v, \lambda] = \text{eig}(\mathbf{M})$$

Larger eigenvector
↓

$$\theta = \text{atan2}(v(2,2), v(1,2))$$
$$\text{conf} = \lambda_2 / \lambda_1$$


5. Threshold on straightness, store segment

Canny lines \rightarrow ... \rightarrow straight edges



Things to remember

- Canny edge detector =
smooth \rightarrow derivative \rightarrow thin \rightarrow threshold \rightarrow link
- Pb: learns weighting of
gradient, color, texture
differences
- Straight line detector =
canny + gradient orientations \rightarrow orientation binning
 \rightarrow linking \rightarrow check for straightness



Acknowledgements

- Thanks to the following researchers for making their teaching/research material online
 - Forsyth
 - Steve Seitz
 - Noah Snavely
 - J.B. Huang
 - Derek Hoiem
 - D. Lowe
 - A. Bobick
 - S. Lazebnik
 - K. Grauman
 - R. Zaleski

Thank you: Question?