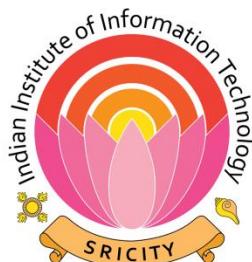


Computer Vision

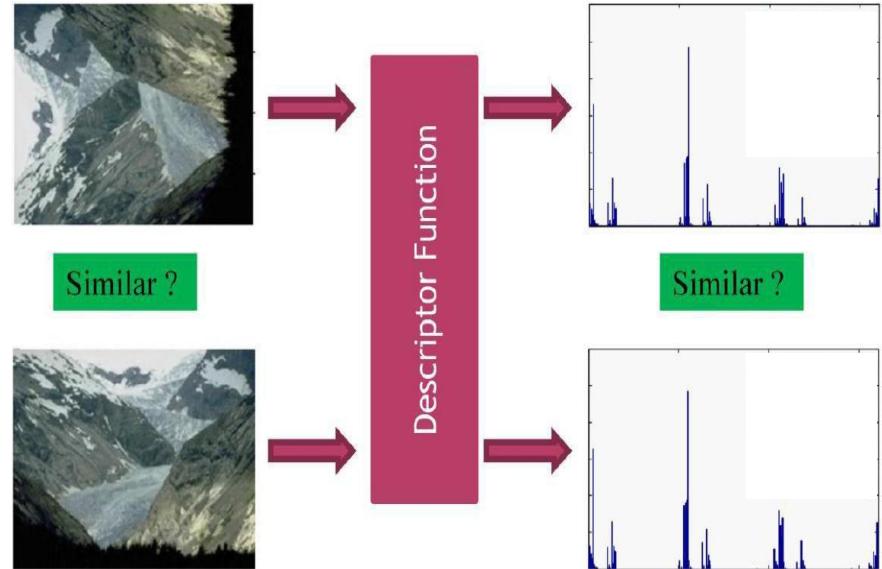
Region Detection & Local Descriptors

Dr. Mrinmoy Ghorai

**Indian Institute of Information Technology
Sri City, Chittoor**

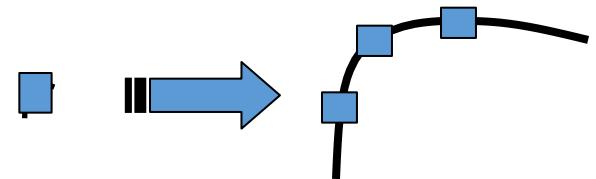


Region Detection & Local Descriptors



Previous Class

- Keypoint detection: repeatable and distinctive
 - Corners, Harris
 - Invariant to scale, rotation, etc.
- Harris Corner Detection
 - Rotation Invariant
 - Partial Intensity Change Invariant
 - *Not Invariant to Scale*

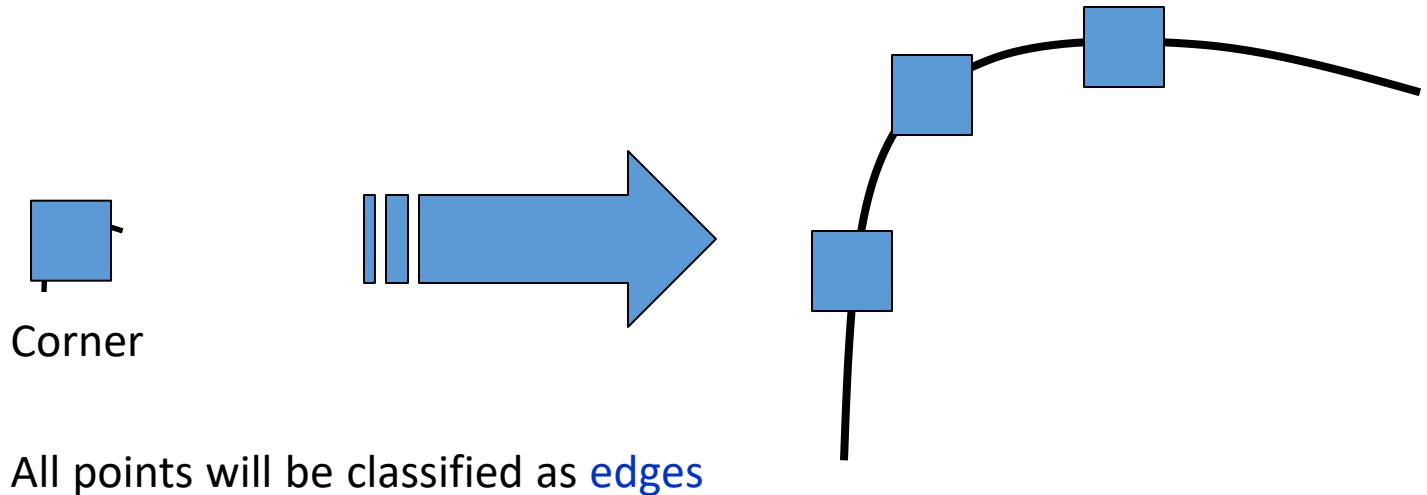


Today's class

- Scale Invariance
- Region Detection
- Local Descriptors
- Image Matching

Harris Detector: Invariance Properties

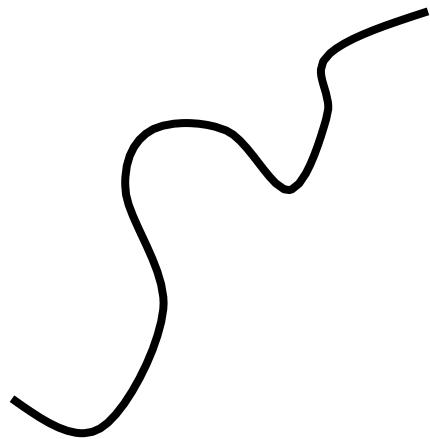
- Scaling



Not invariant to scaling

Scale invariant detection

Suppose you're looking for corners

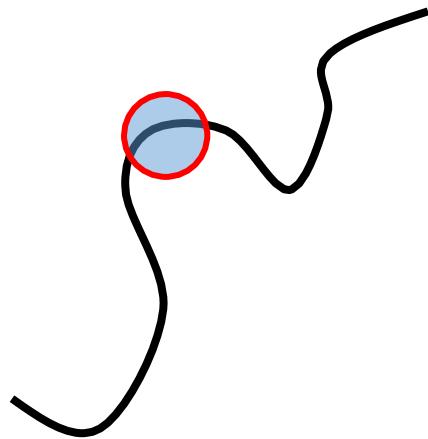


Key idea: find scale that gives local maximum of f

- in both position and scale
- One definition of f : the Harris operator

Scale invariant detection

Suppose you're looking for corners

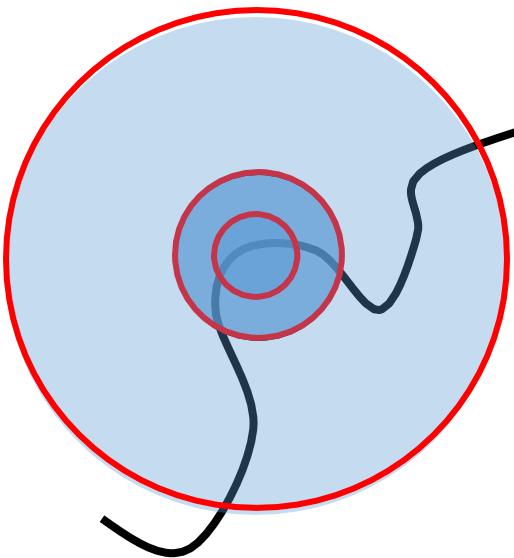


Key idea: find scale that gives local maximum of f

- in both position and scale
- One definition of f : the Harris operator

Scale invariant detection

Suppose you're looking for corners



Key idea: find scale that gives local maximum of f

- in both position and scale
- One definition of f : the Harris operator

Automatic Scale Selection

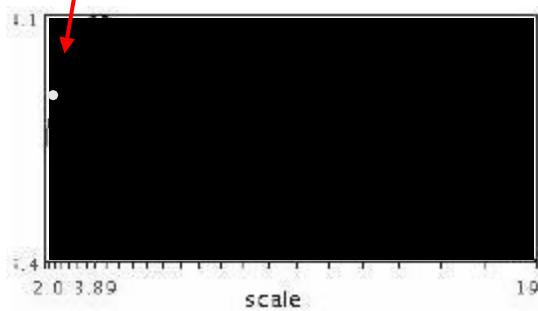


$$f(I_{i_1 \dots i_m}(x, \sigma)) = f(I_{i_1 \dots i_m}(x', \sigma'))$$

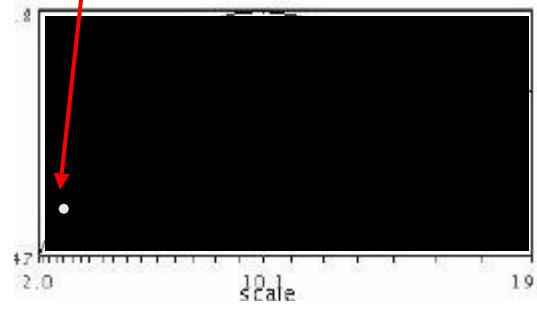
How to find corresponding patch sizes?

Automatic Scale Selection

- Function responses for increasing scale (scale signature)



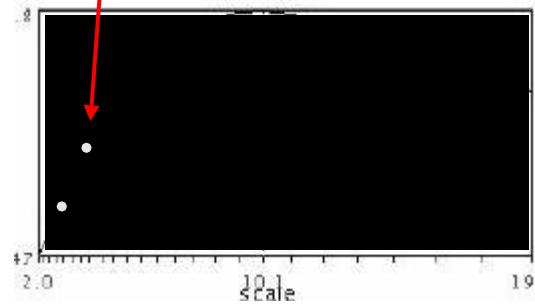
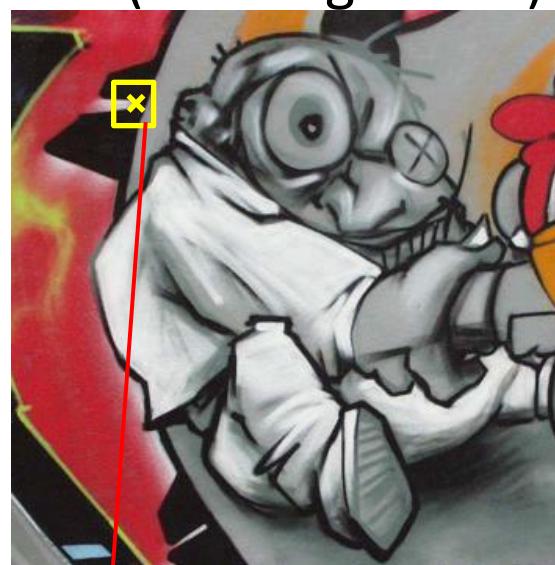
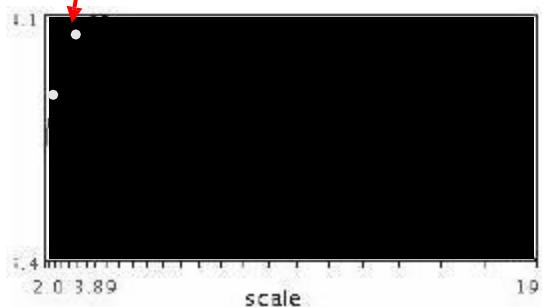
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma))$$

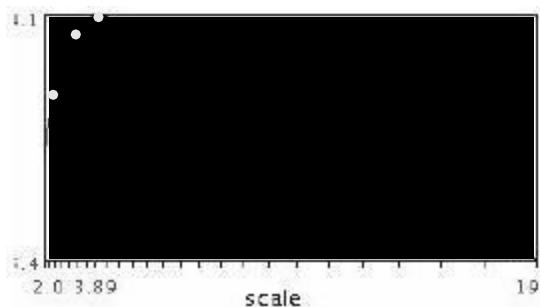
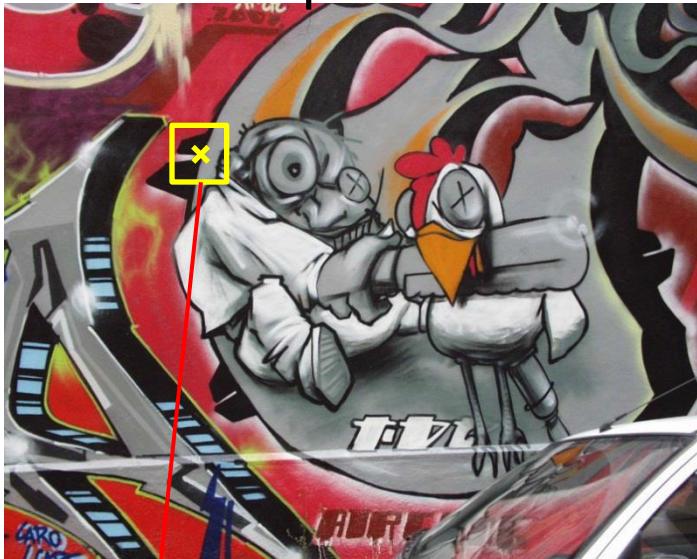
Automatic Scale Selection

- Function responses for increasing scale (scale signature)

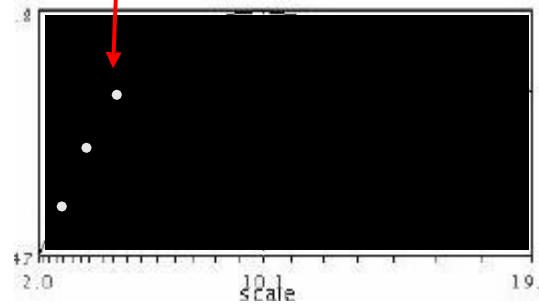


Automatic Scale Selection

- Function responses for increasing scale (scale signature)



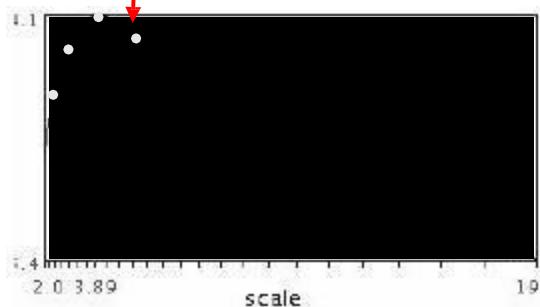
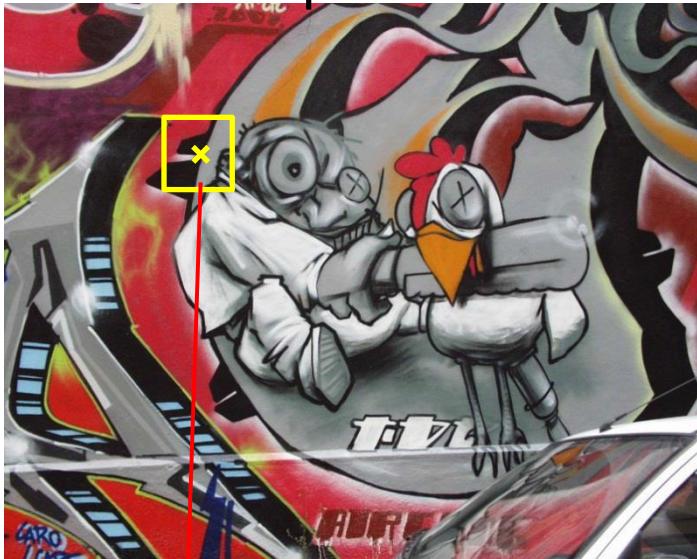
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



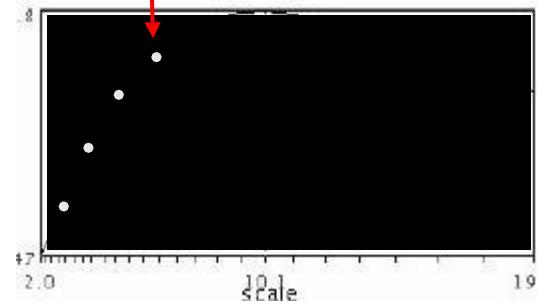
$$f(I_{i_1 \dots i_m}(x', \sigma))$$

Automatic Scale Selection

- Function responses for increasing scale (scale signature)



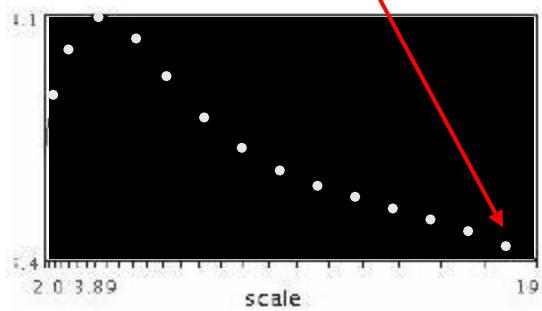
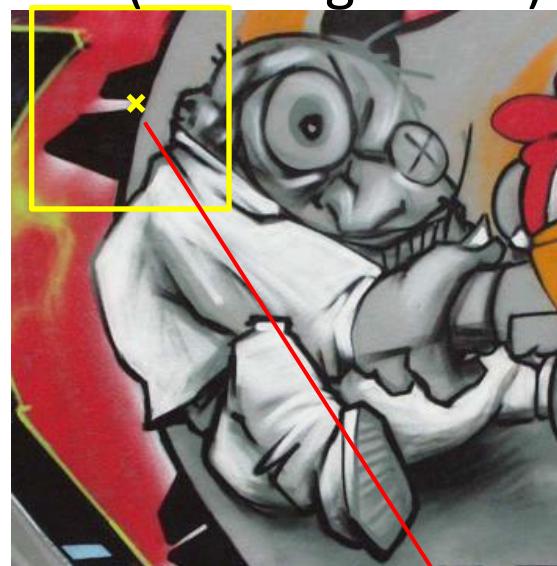
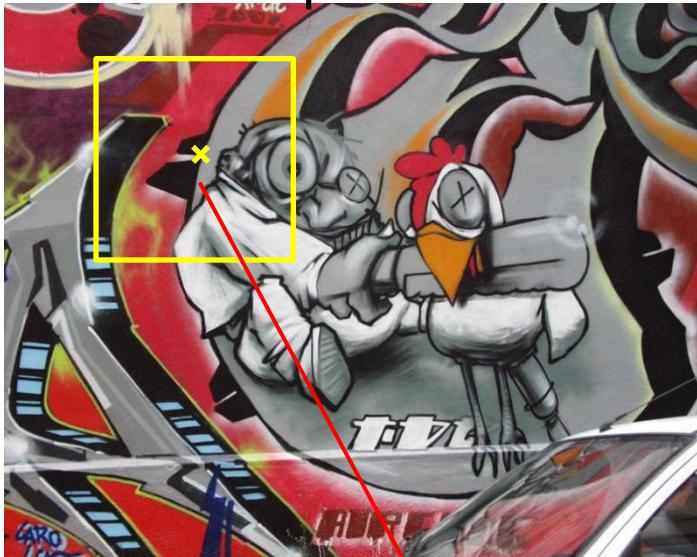
$$f(I_{i_1...i_m}(x, \sigma))$$



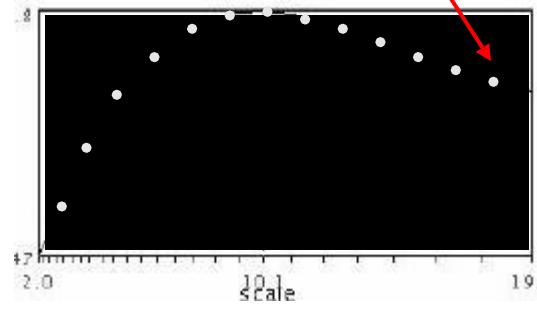
$$f(I_{i_1...i_m}(x', \sigma))$$

Automatic Scale Selection

- Function responses for increasing scale (scale signature)



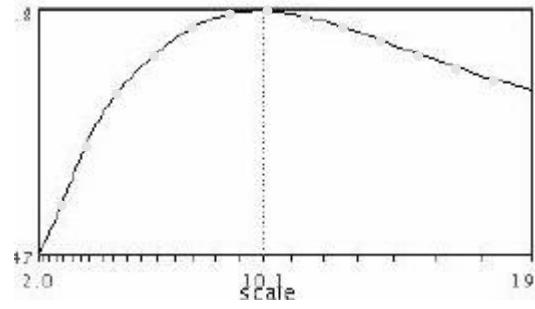
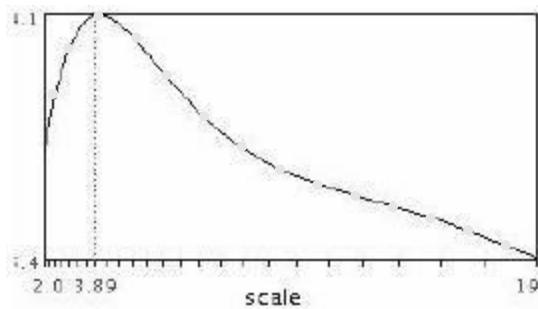
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma))$$

Automatic Scale Selection

- Function responses for increasing scale (scale signature)



Implementation

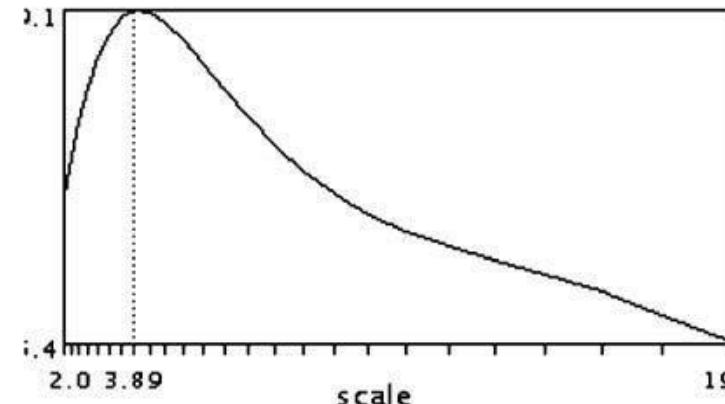
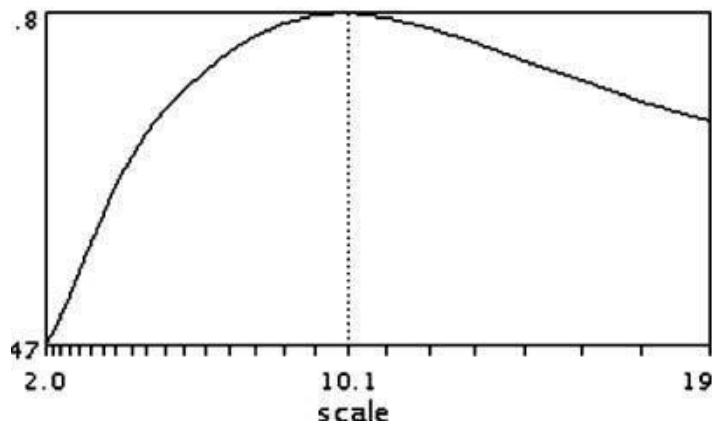
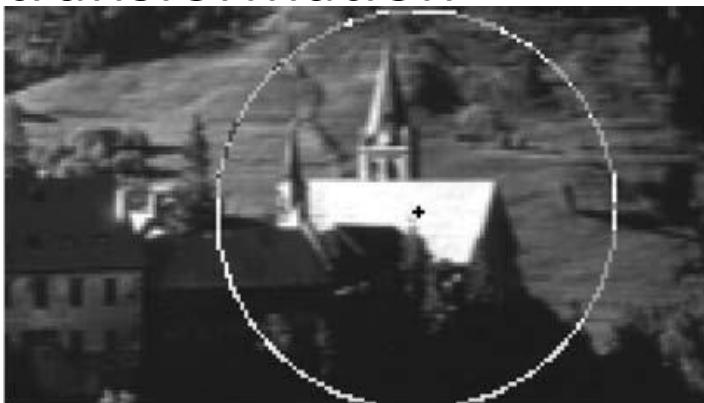
- Instead of computing f for larger and larger windows, we can implement using a fixed window size with a Gaussian pyramid



(sometimes need to create in-between levels, e.g. a $\frac{3}{4}$ -size image)

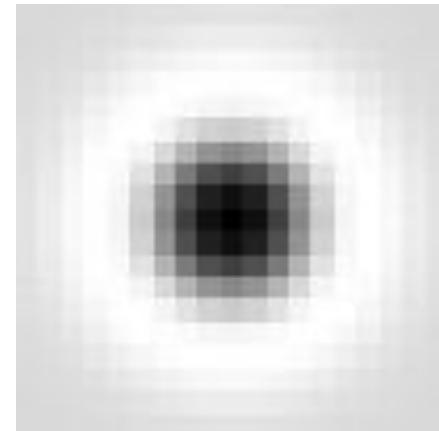
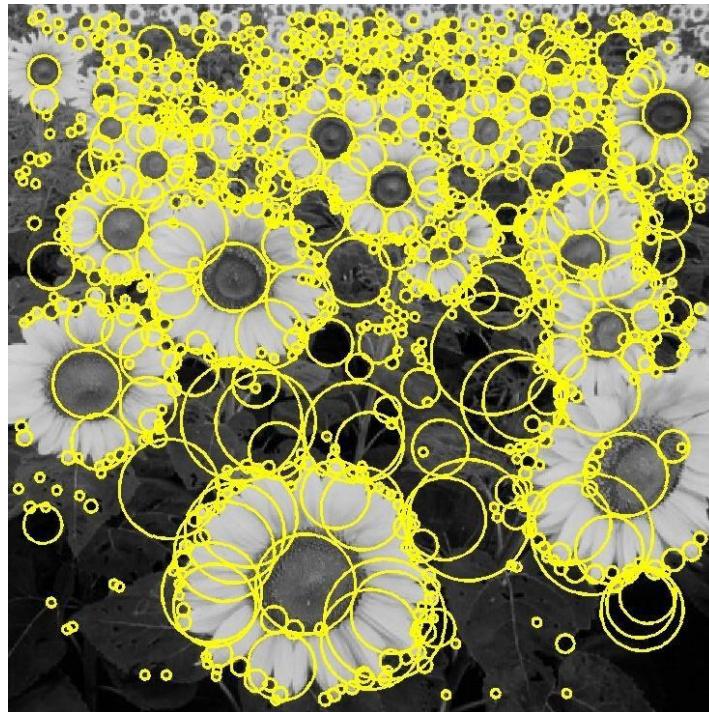
Keypoint detection with scale selection

- We want to extract keypoints with characteristic scale that is *covariant* with the image transformation



Basic idea

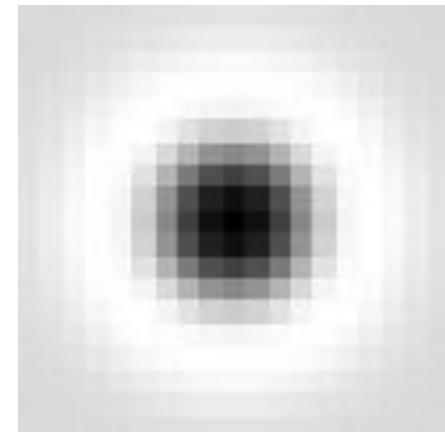
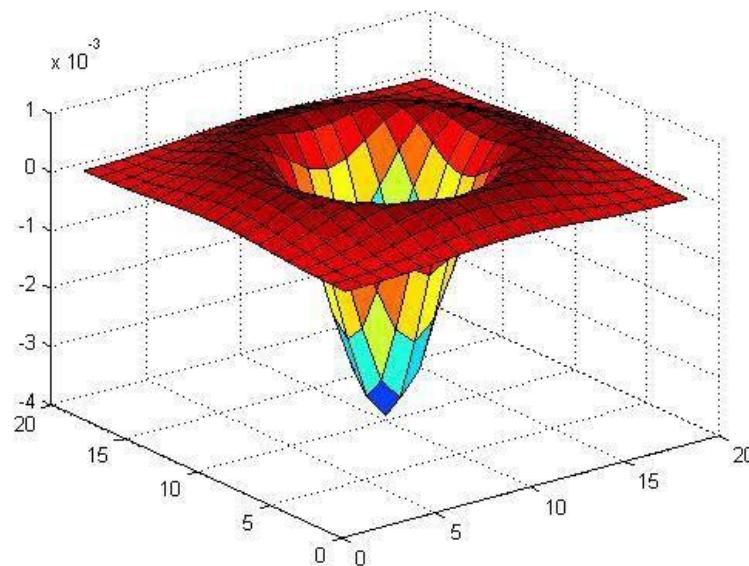
- Convolve the image with a “blob filter” at multiple scales and look for extrema of filter response in the resulting *scale space*



T. Lindeberg. [Feature detection with automatic scale selection.](#)
IJCV 30(2), pp 77-116, 1998.

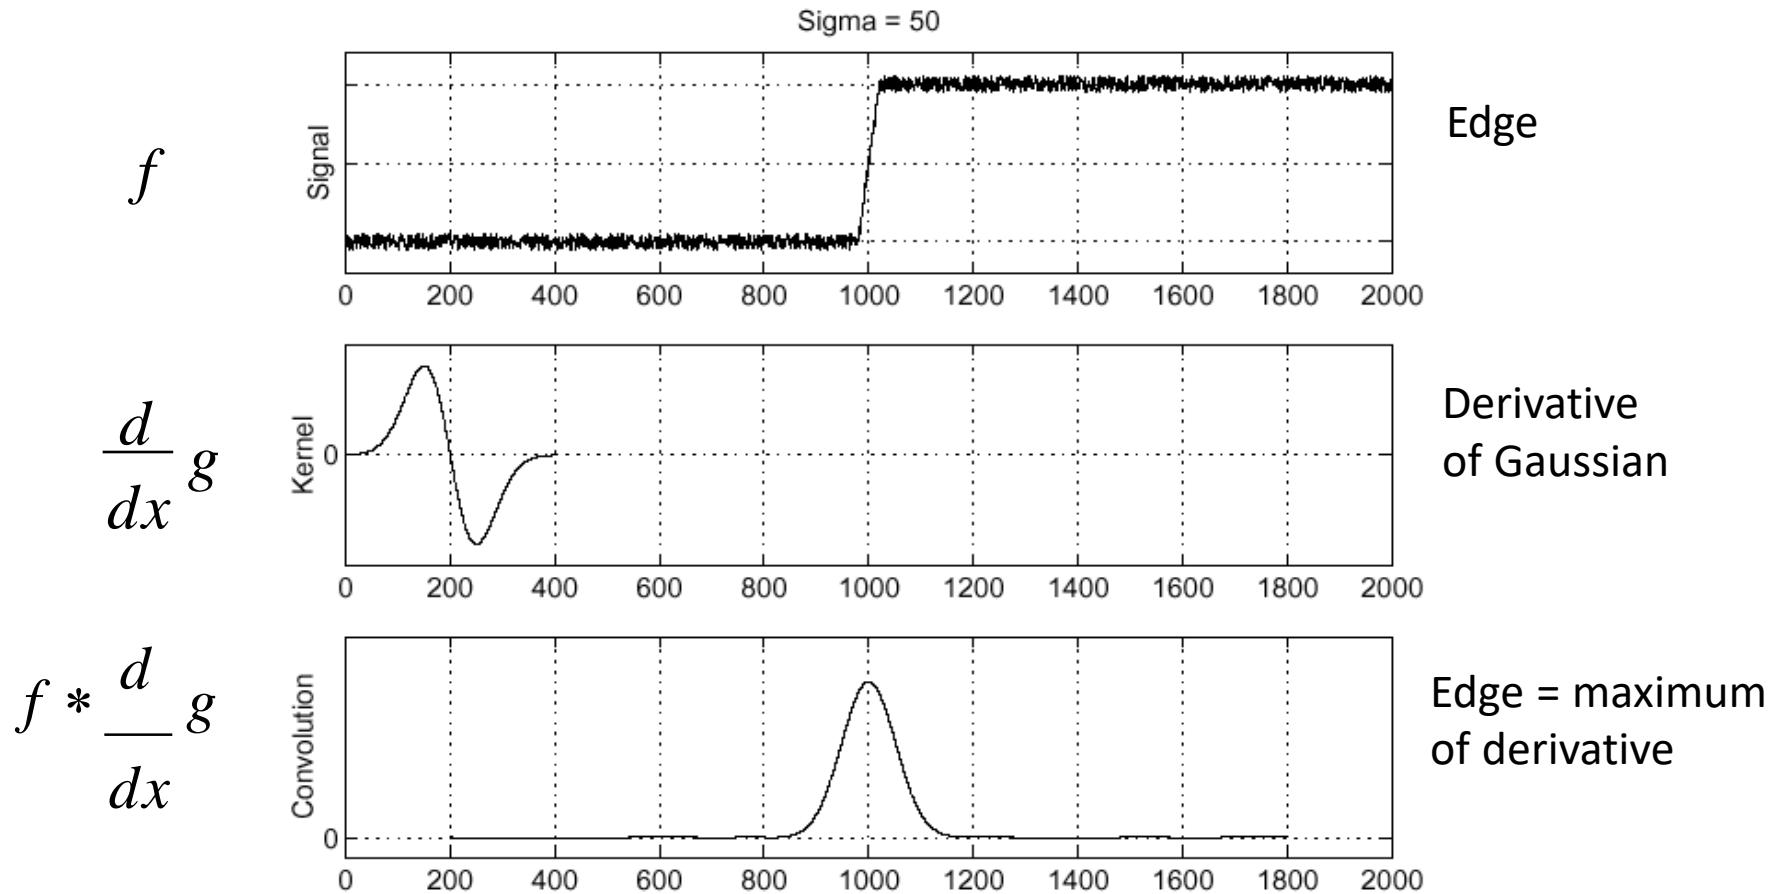
Blob filter

- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D

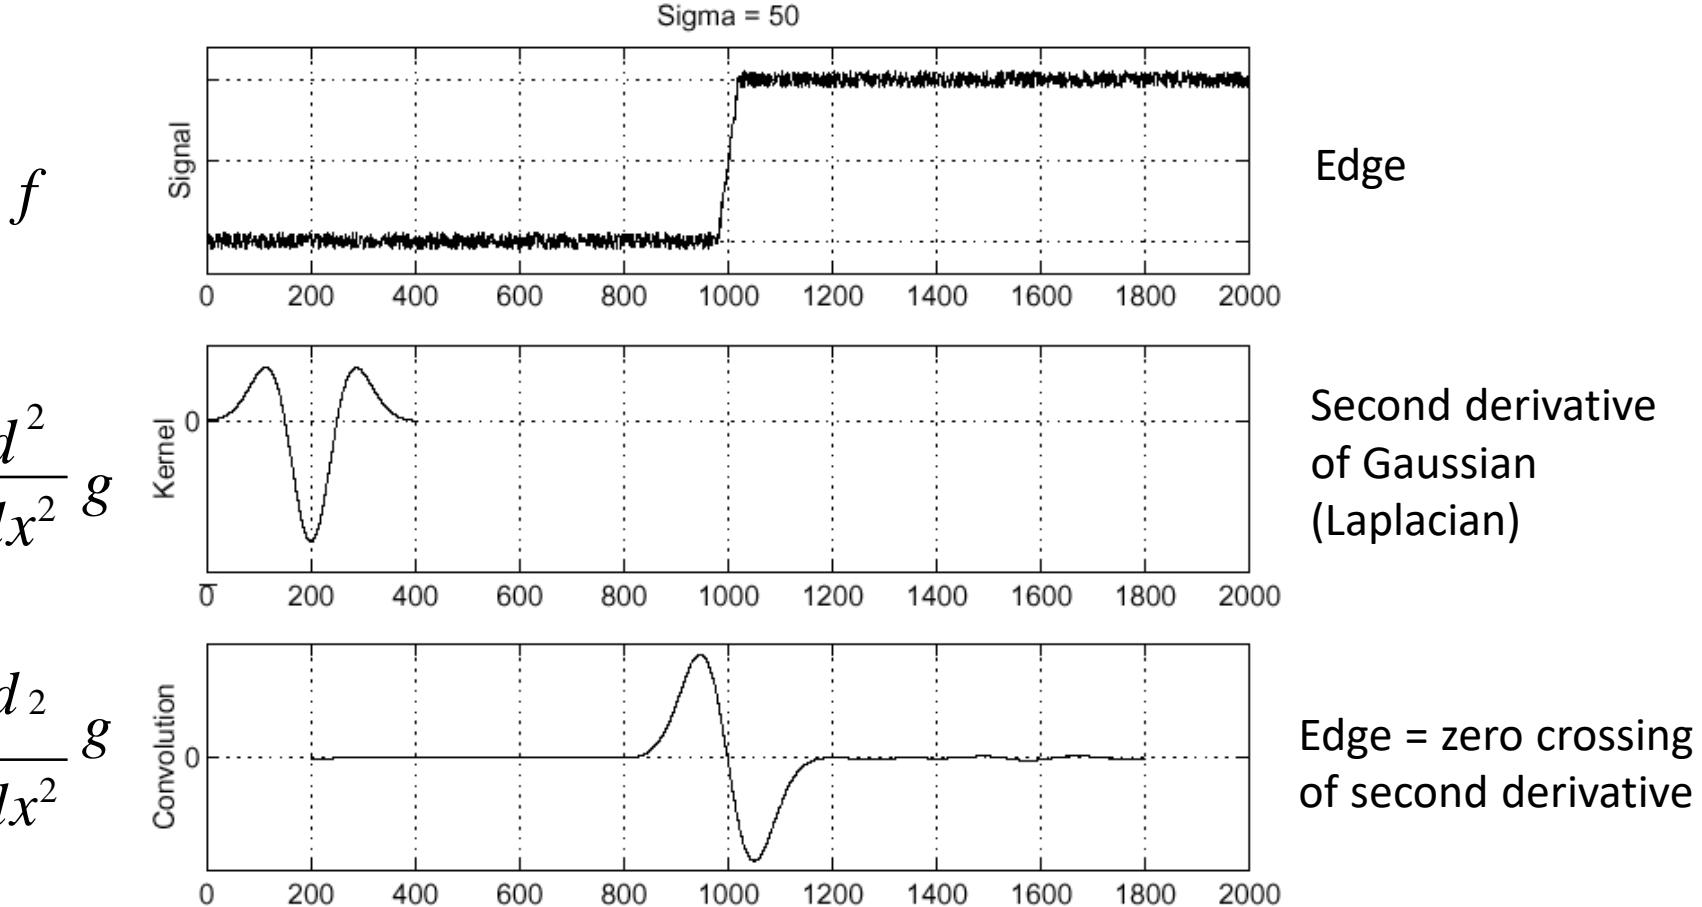


$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

Recall: Edge detection

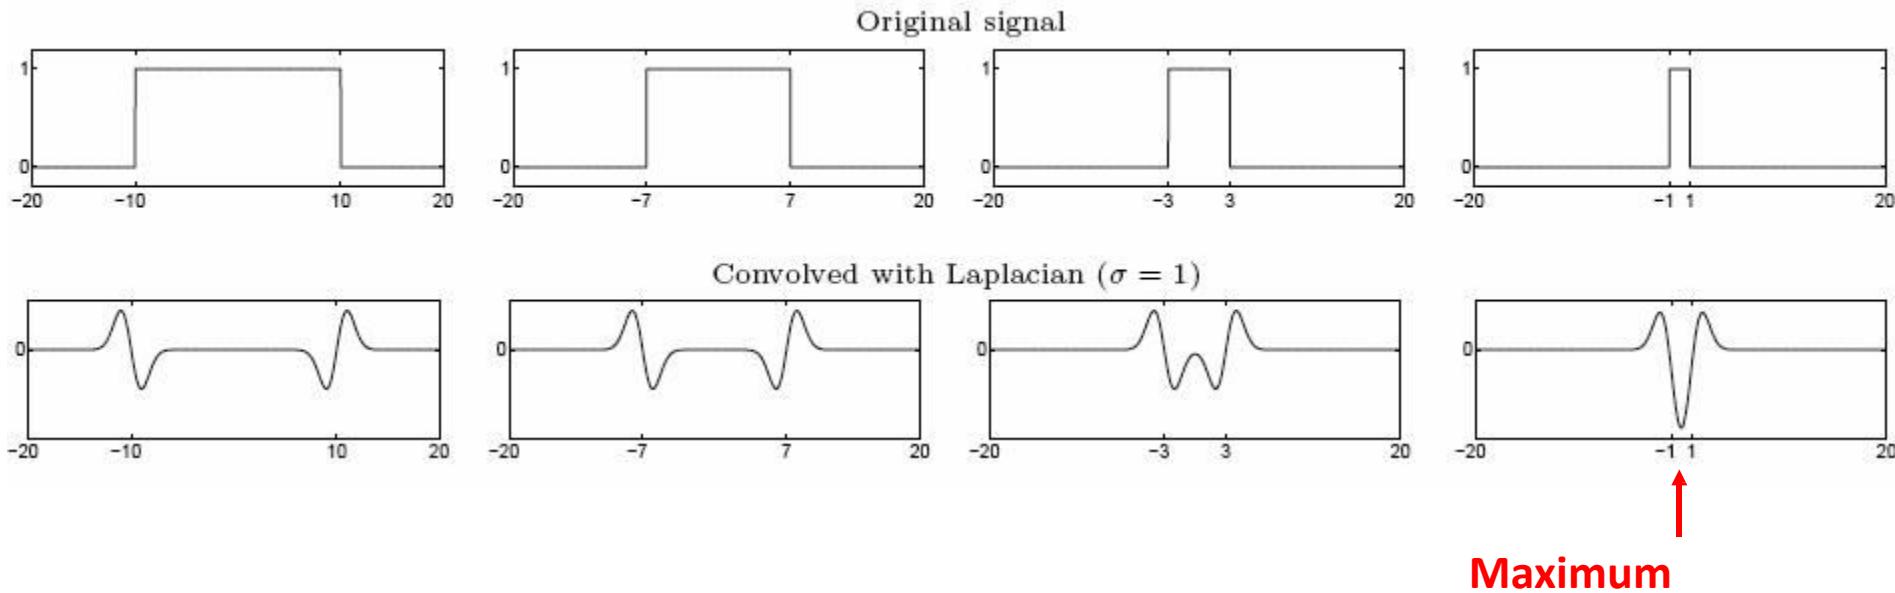


Edge detection, Take 2



From edges to blobs

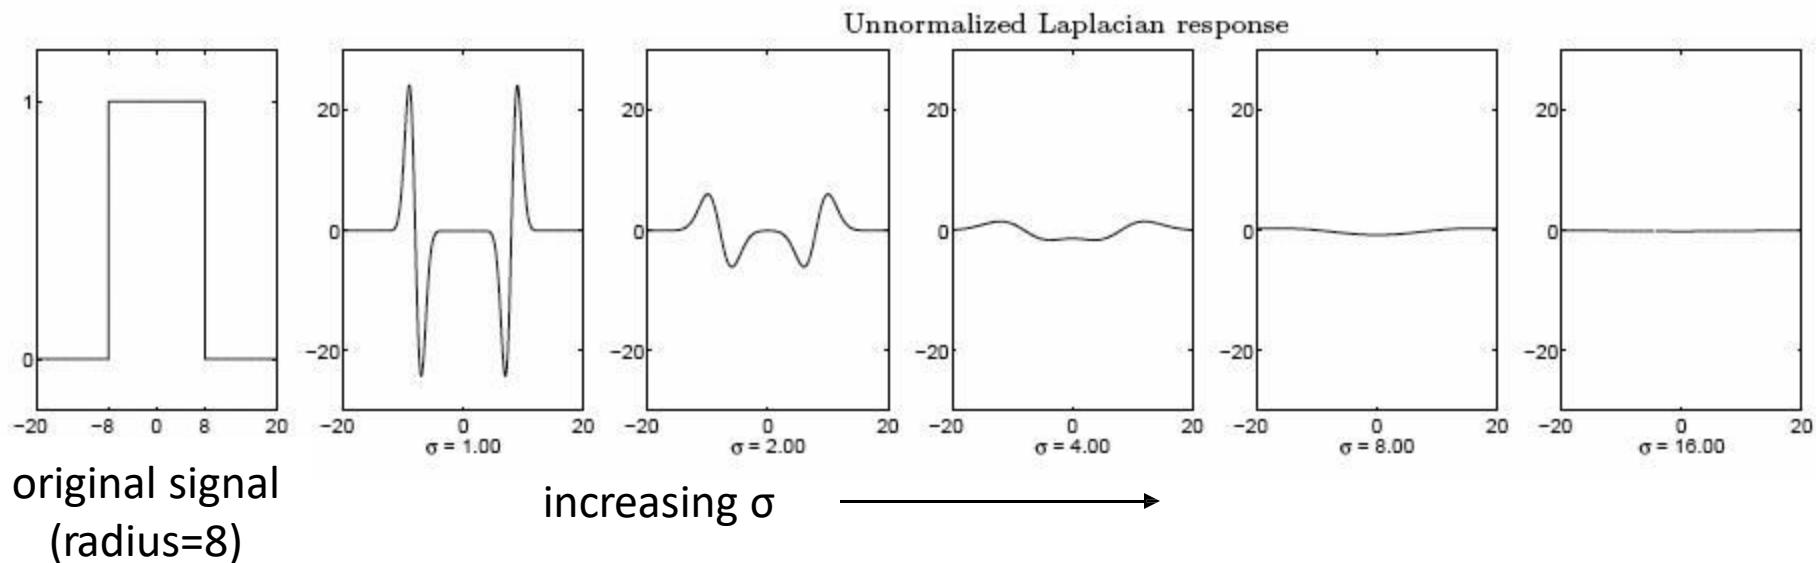
- Edge = ripple
- Blob = superposition of two ripples



Spatial selection: the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is “matched” to the scale of the blob

Scale selection

- We want to find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response
- However, Laplacian response decays as scale increases:

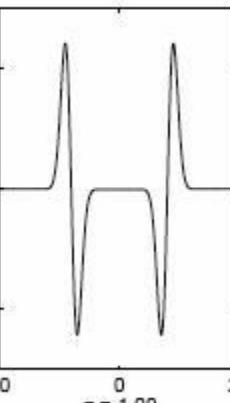
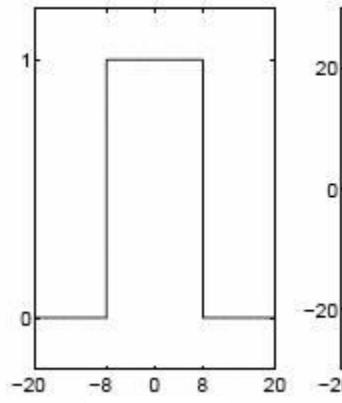


Scale normalization

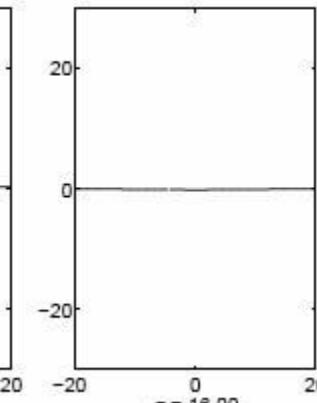
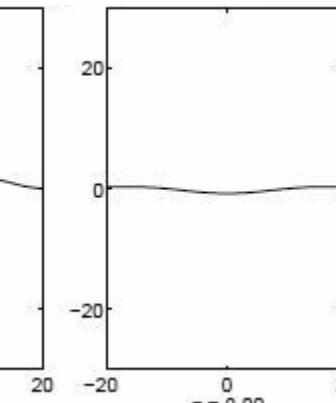
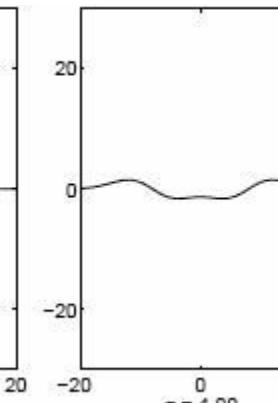
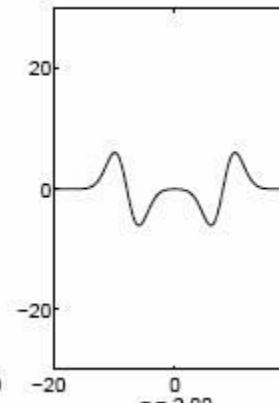
- The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases
- To keep response the same (scale-invariant), must multiply Gaussian derivative by σ
- Laplacian is the second Gaussian derivative, so it must be multiplied by σ^2

Effect of scale normalization

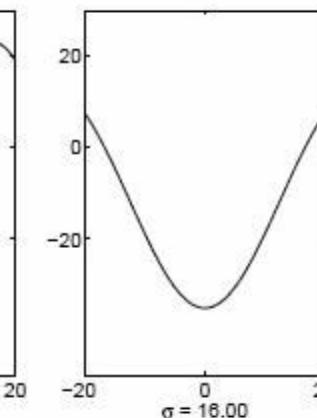
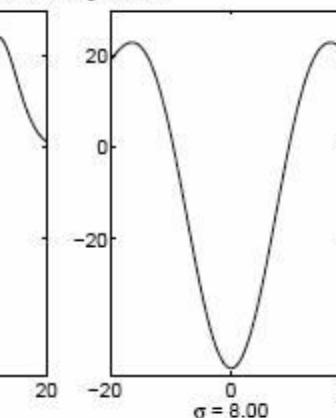
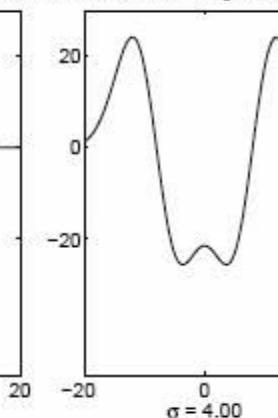
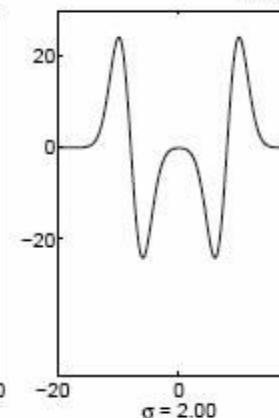
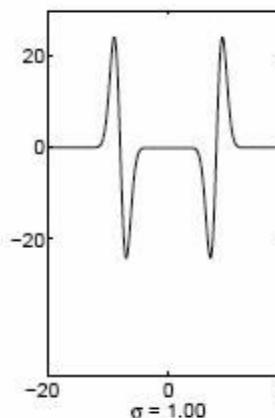
Original signal



Unnormalized Laplacian response



Scale-normalized Laplacian response

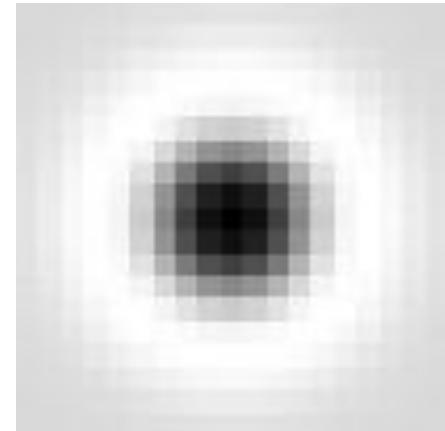
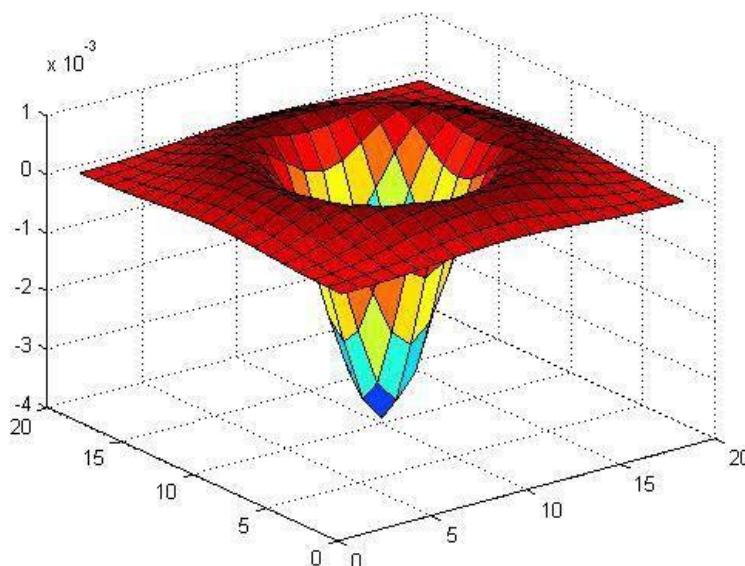


maximum

Blob detection in 2D

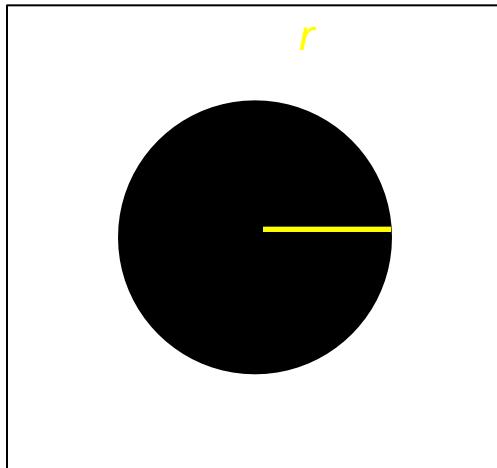
- *Scale-normalized Laplacian of Gaussian:*

$$\nabla_{\text{norm}}^2 g = \sigma^2 \left(\frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

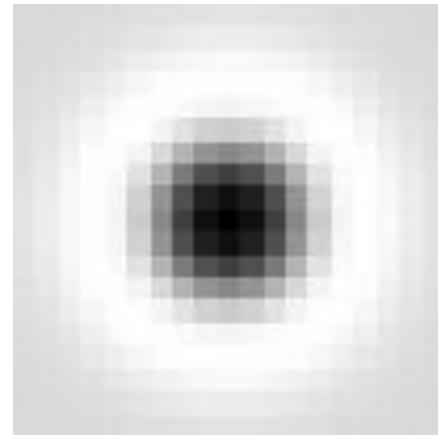


Blob detection in 2D

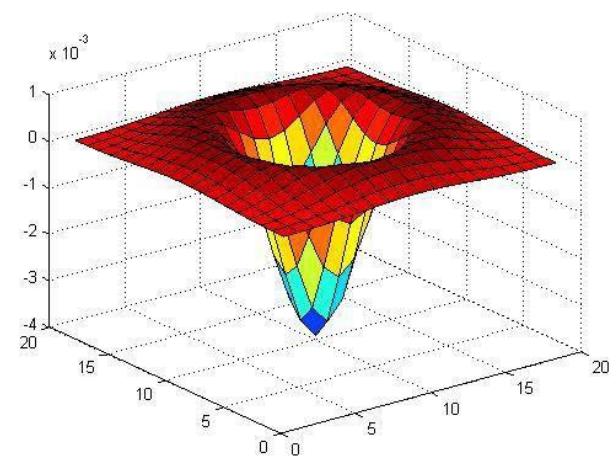
- At what scale does the Laplacian achieve a maximum response to a binary circle of radius r ?



image

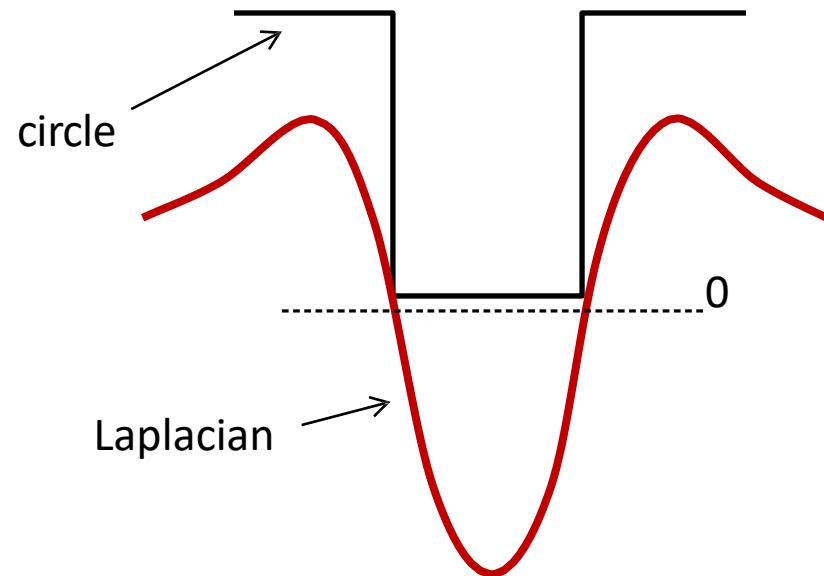
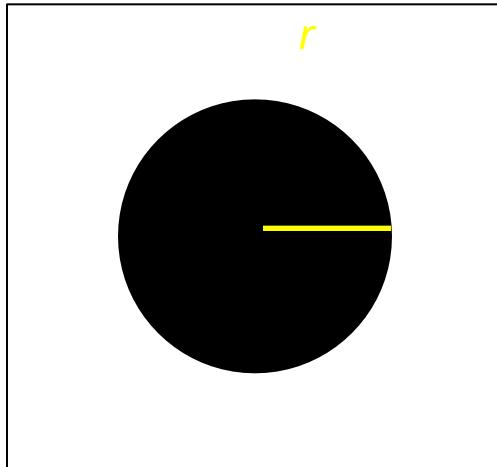


Laplacian



Blob detection in 2D

- At what scale does the Laplacian achieve a maximum response to a binary circle of radius r ?
- To get maximum response, the zeros of the Laplacian have to be aligned with the circle

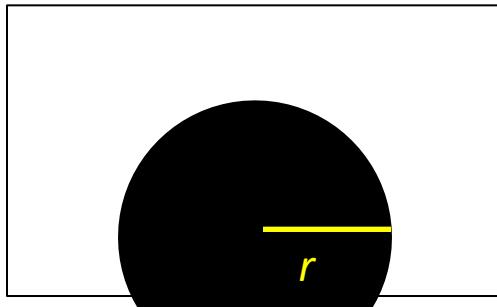


image

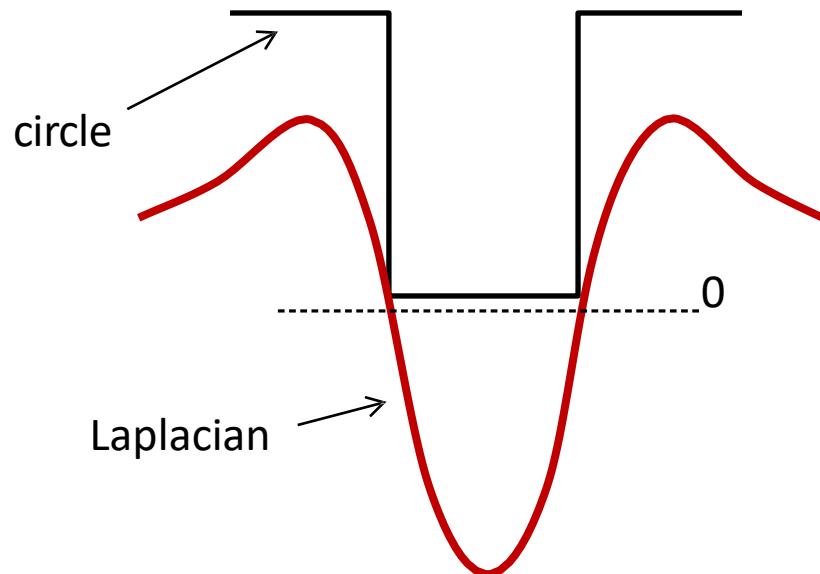
Blob detection in 2D

- At what scale does the Laplacian achieve a maximum response to a binary circle of radius r ?
- To get maximum response, the zeros of the Laplacian have to be aligned with the circle
- The Laplacian is given by (up to scale):

$$(x^2 + y^2 - 2\sigma^2) e^{-(x^2+y^2)/2\sigma^2}$$



image

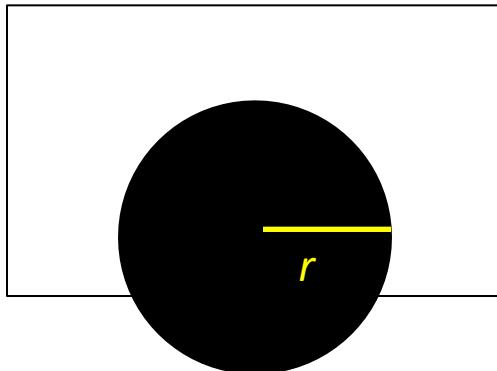


Blob detection in 2D

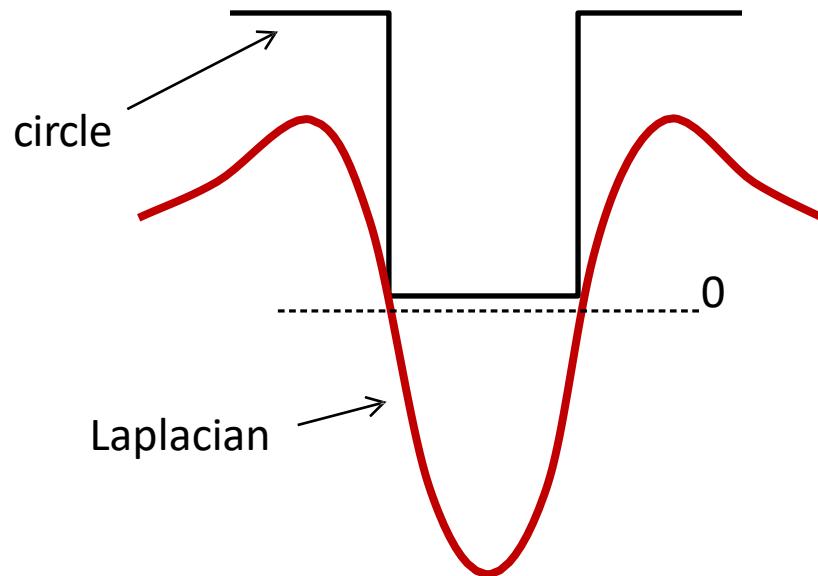
- At what scale does the Laplacian achieve a maximum response to a binary circle of radius r ?
- To get maximum response, the zeros of the Laplacian have to be aligned with the circle
- The Laplacian is given by (up to scale):

$$(x^2 + y^2 - 2\sigma^2) e^{-(x^2+y^2)/2\sigma^2}$$

- Therefore, the maximum response occurs at $\sigma = r / \sqrt{2}$.



image



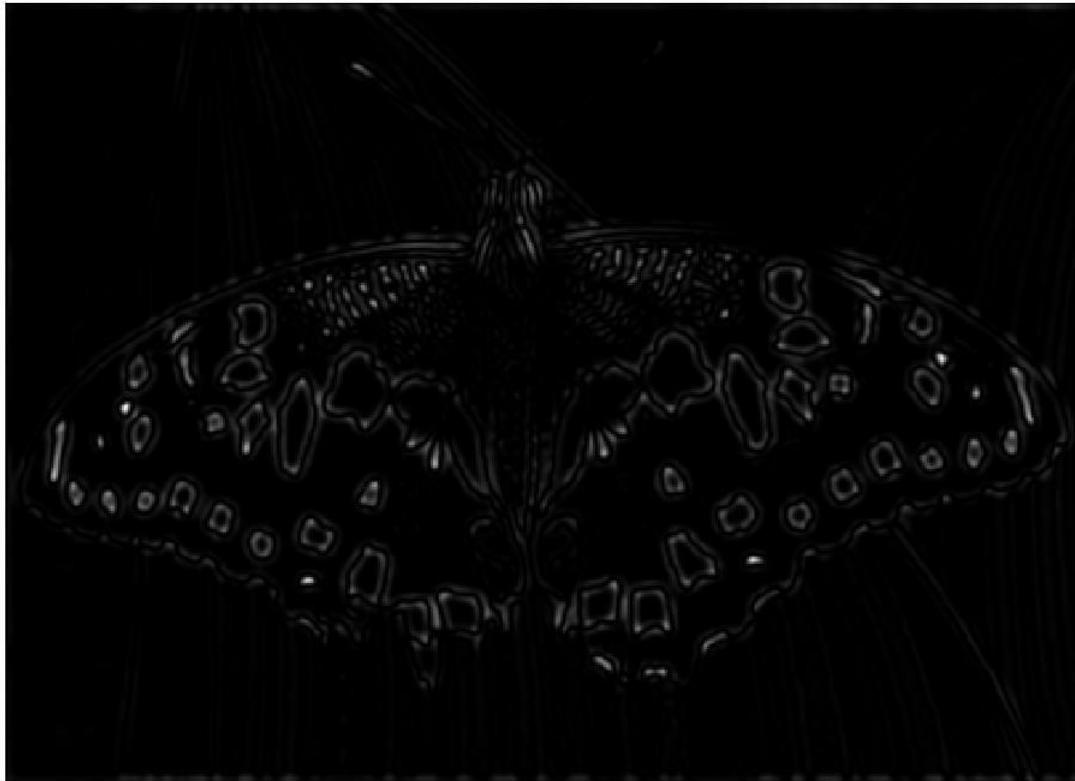
Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales

Scale-space blob detector: Example

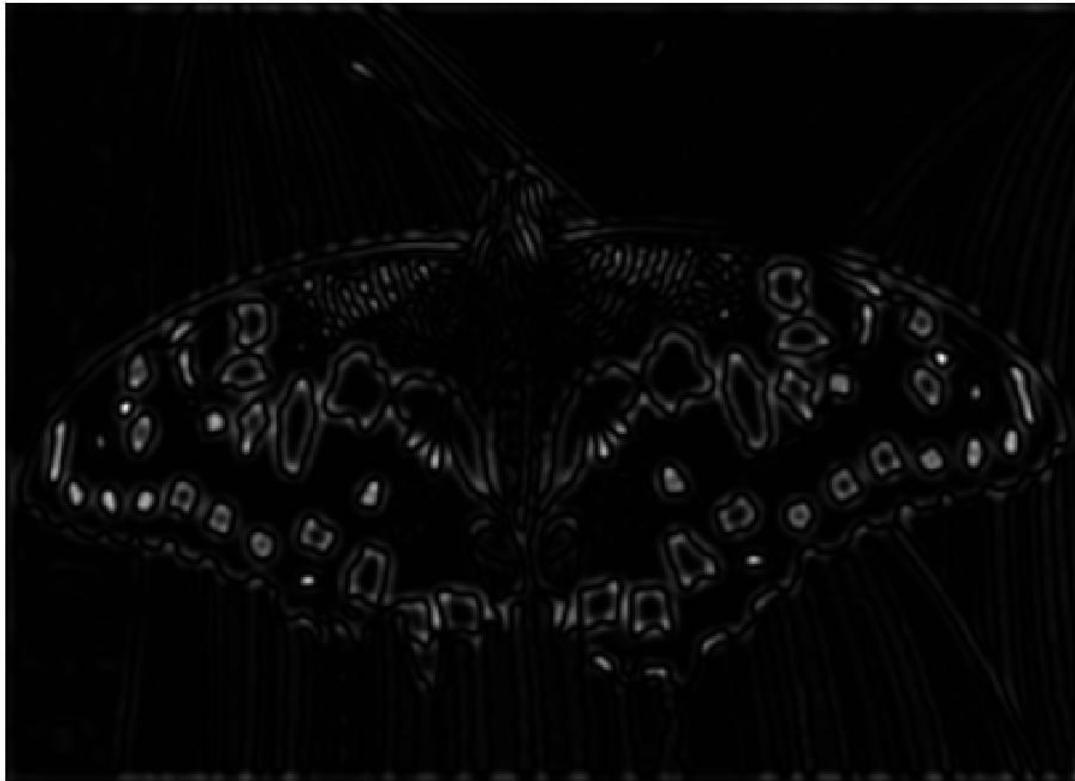


Scale-space blob detector: Example



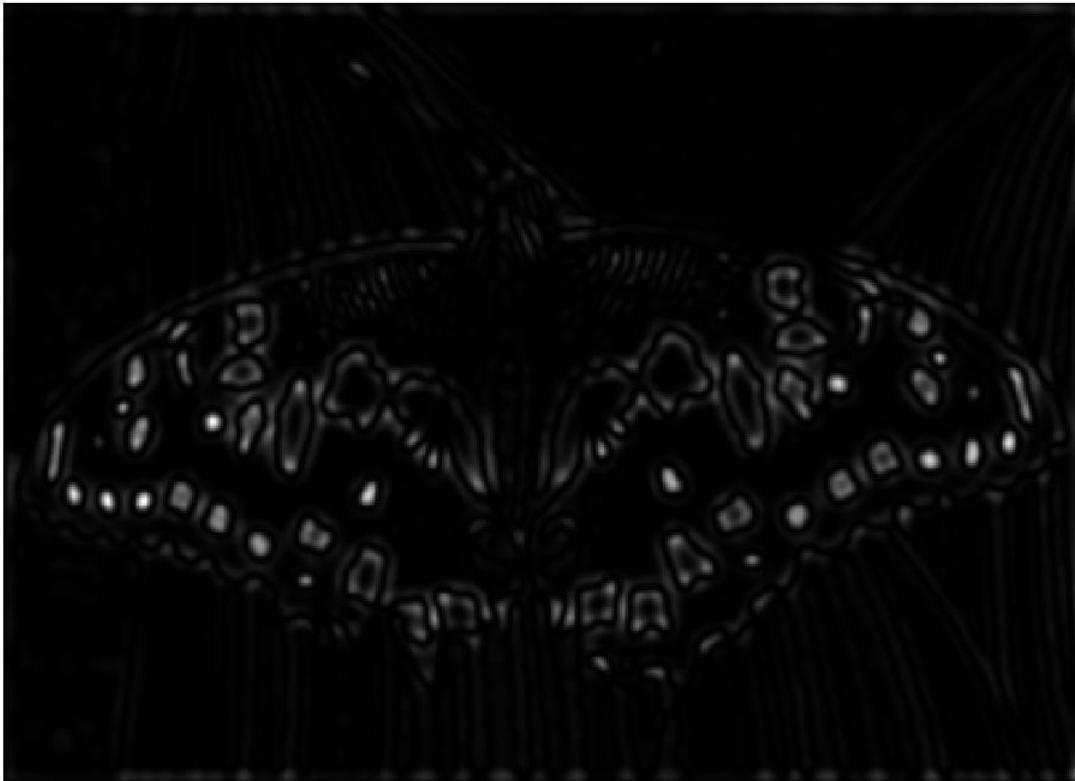
$\sigma = 2$

Scale-space blob detector: Example



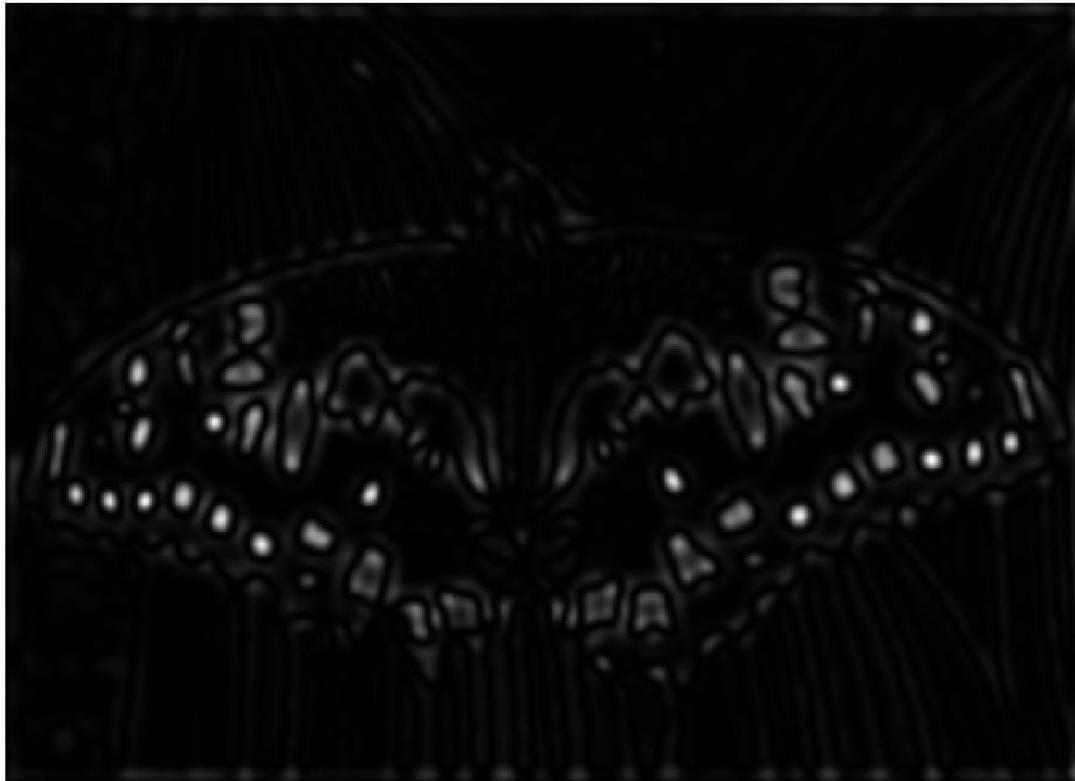
$\sigma = 2.5018$

Scale-space blob detector: Example



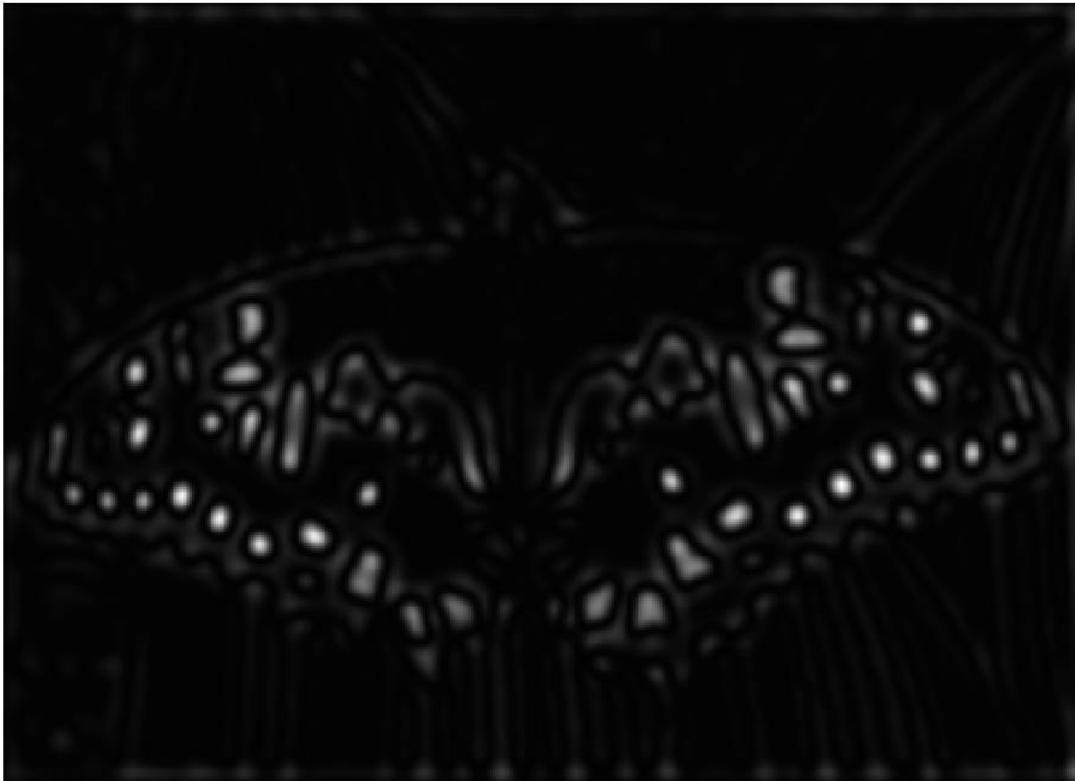
$\sigma = 3.1296$

Scale-space blob detector: Example



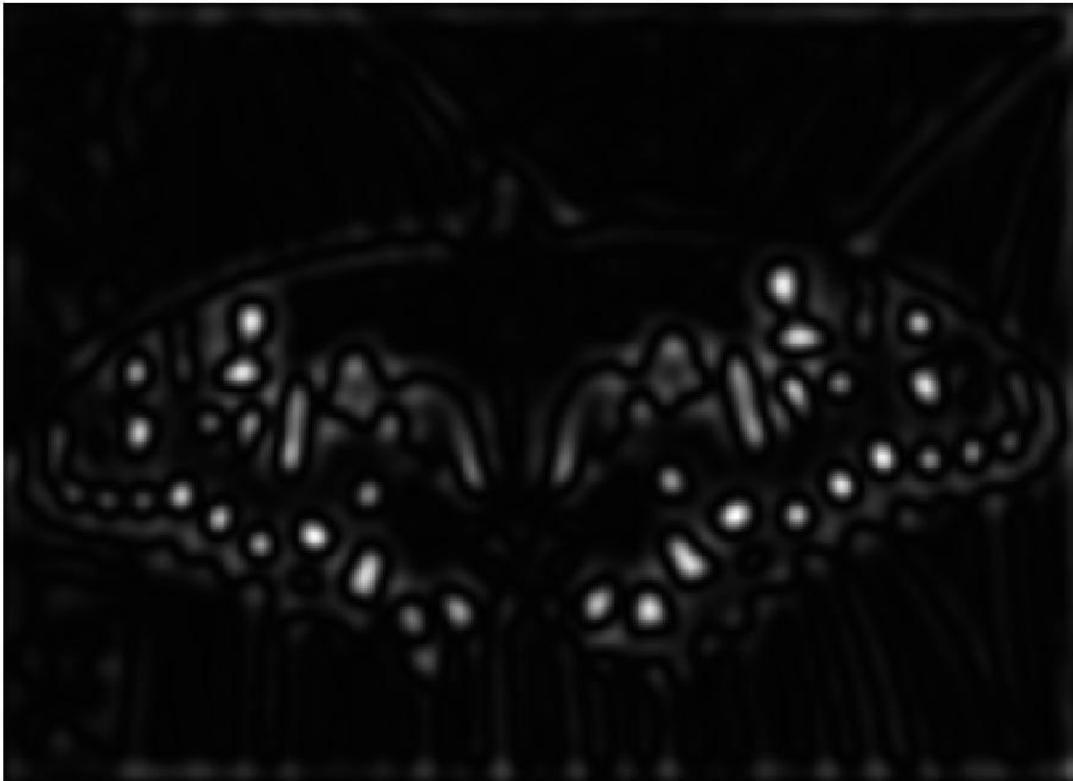
$\sigma = 3.9149$

Scale-space blob detector: Example



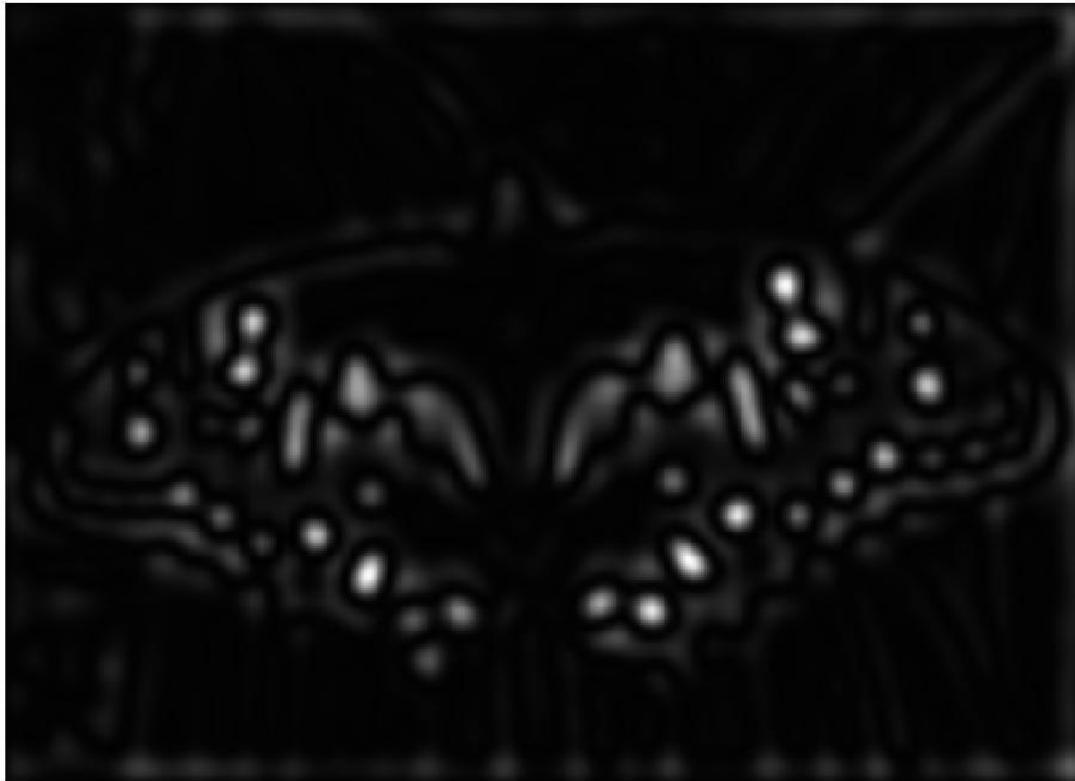
$\sigma = 4.8972$

Scale-space blob detector: Example



$\sigma = 6.126$

Scale-space blob detector: Example



$\sigma = 7.6631$

Scale-space blob detector: Example



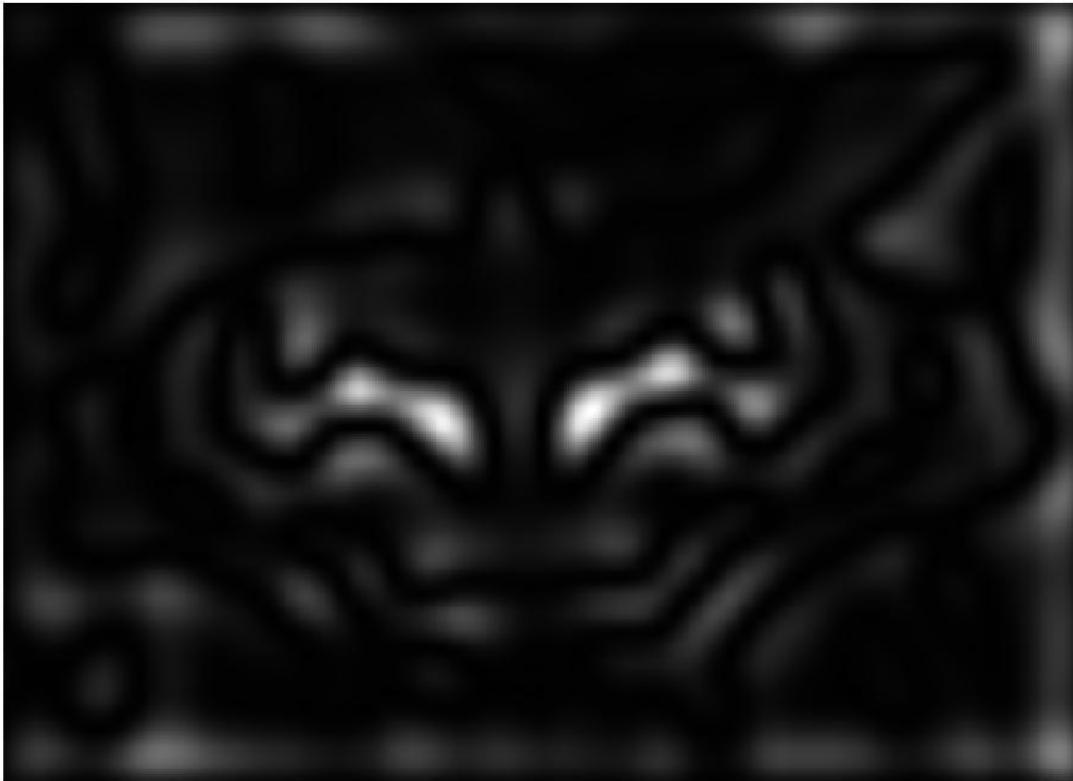
sigma = 9.5859

Scale-space blob detector: Example



$\sigma = 11.9912$

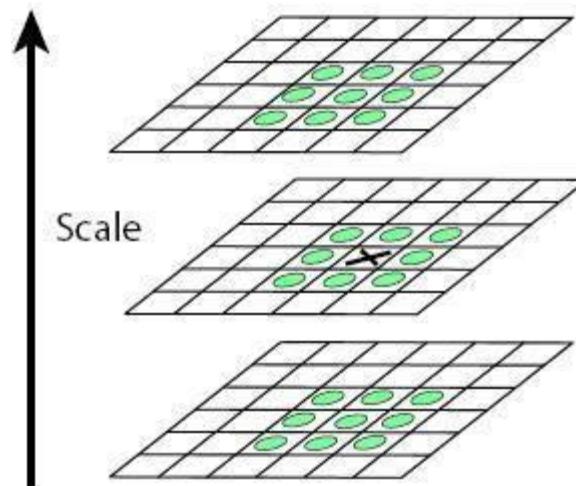
Scale-space blob detector Example



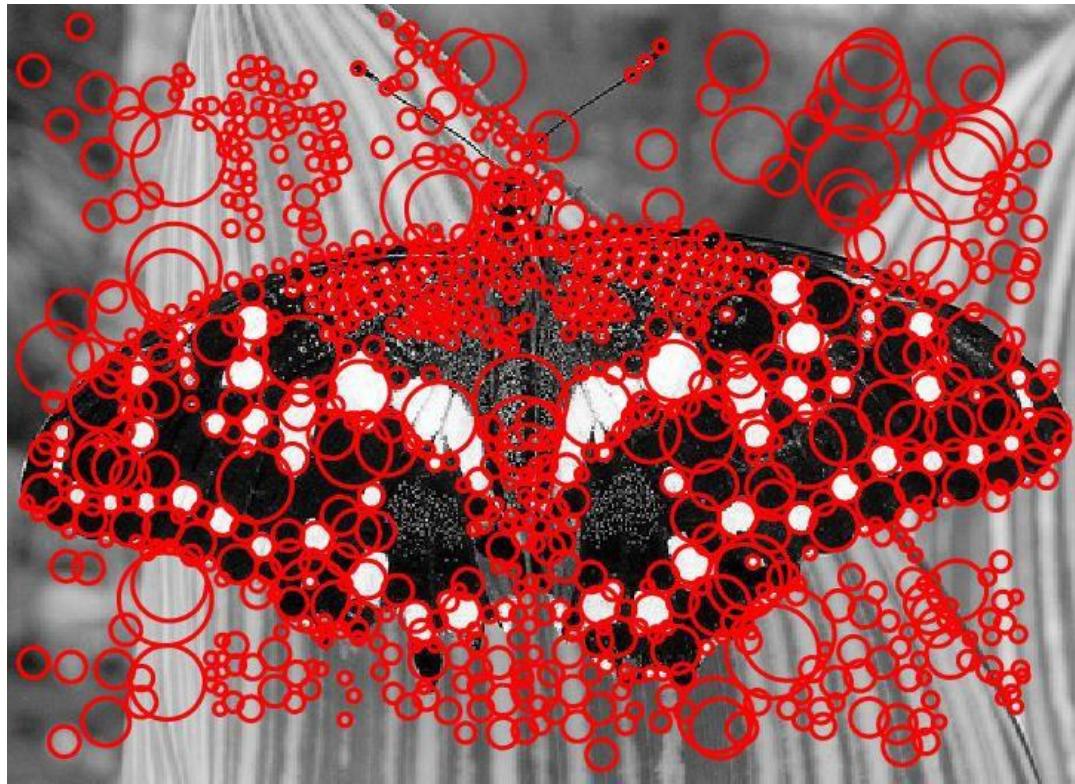
sigma = 15

Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales
2. Find maxima of squared Laplacian response in scale-space



Scale-space blob detector: Example



Efficient implementation (SIFT)

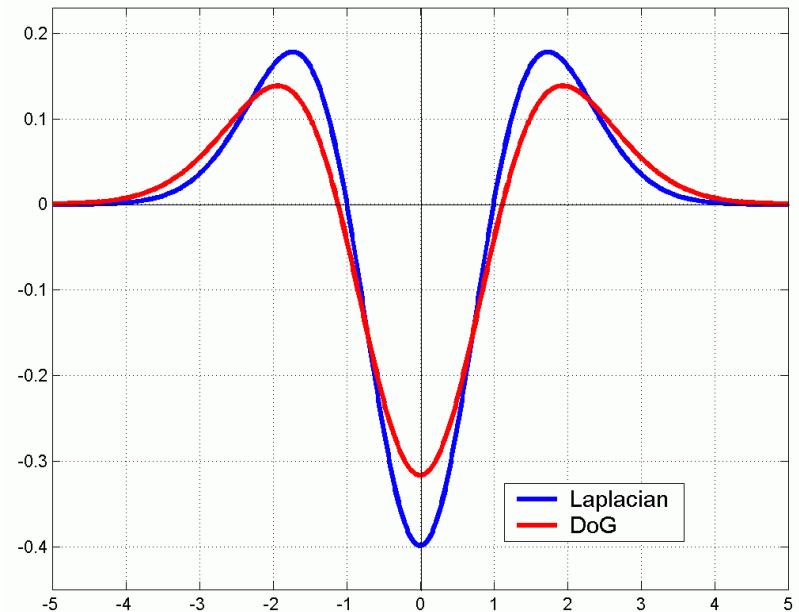
- Approximating the Laplacian with a difference of Gaussians by SIFT detector:

$$L = \sigma^2 \left(G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma) \right)$$

(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



Difference-of-Gaussian (DoG)



-

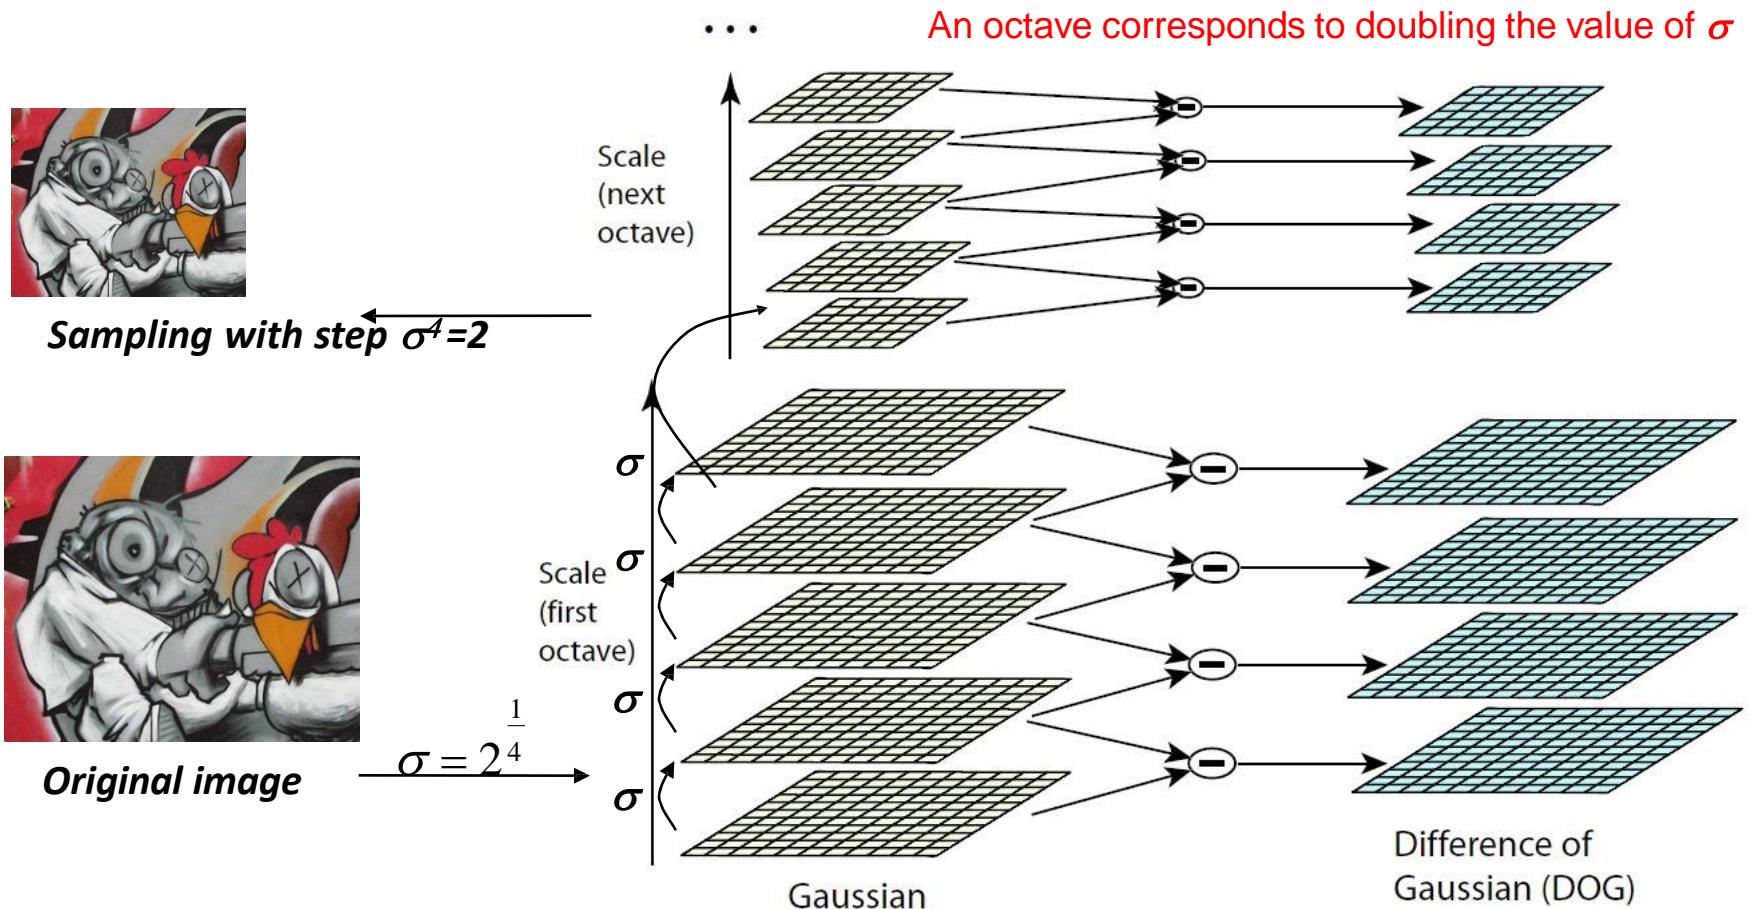


=

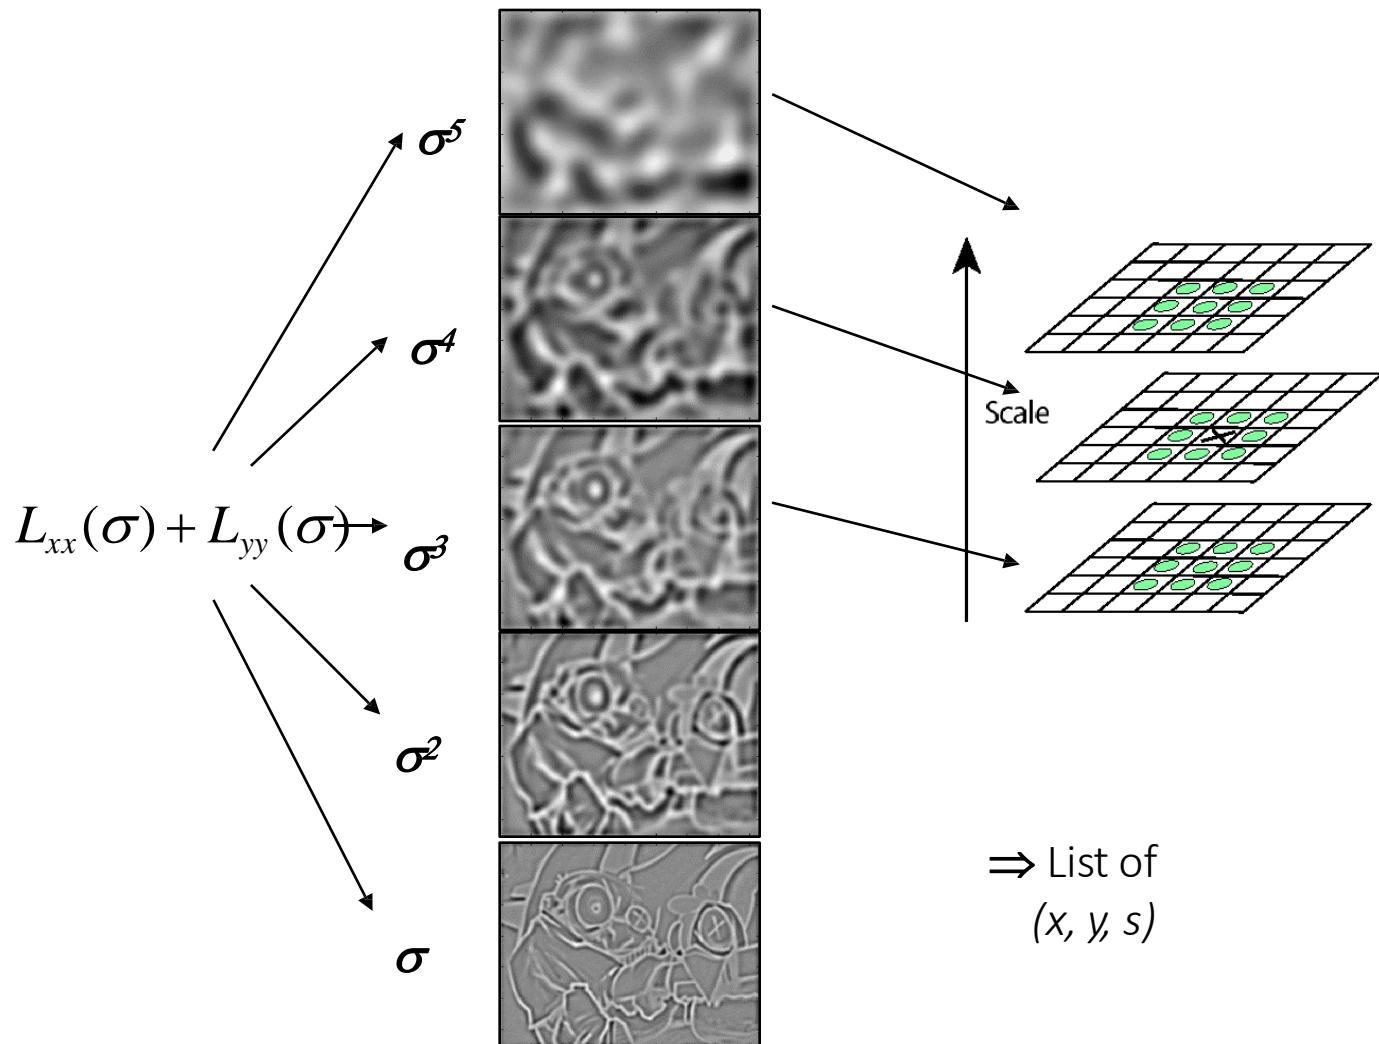
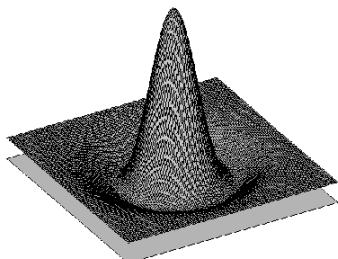


DoG – Efficient Computation

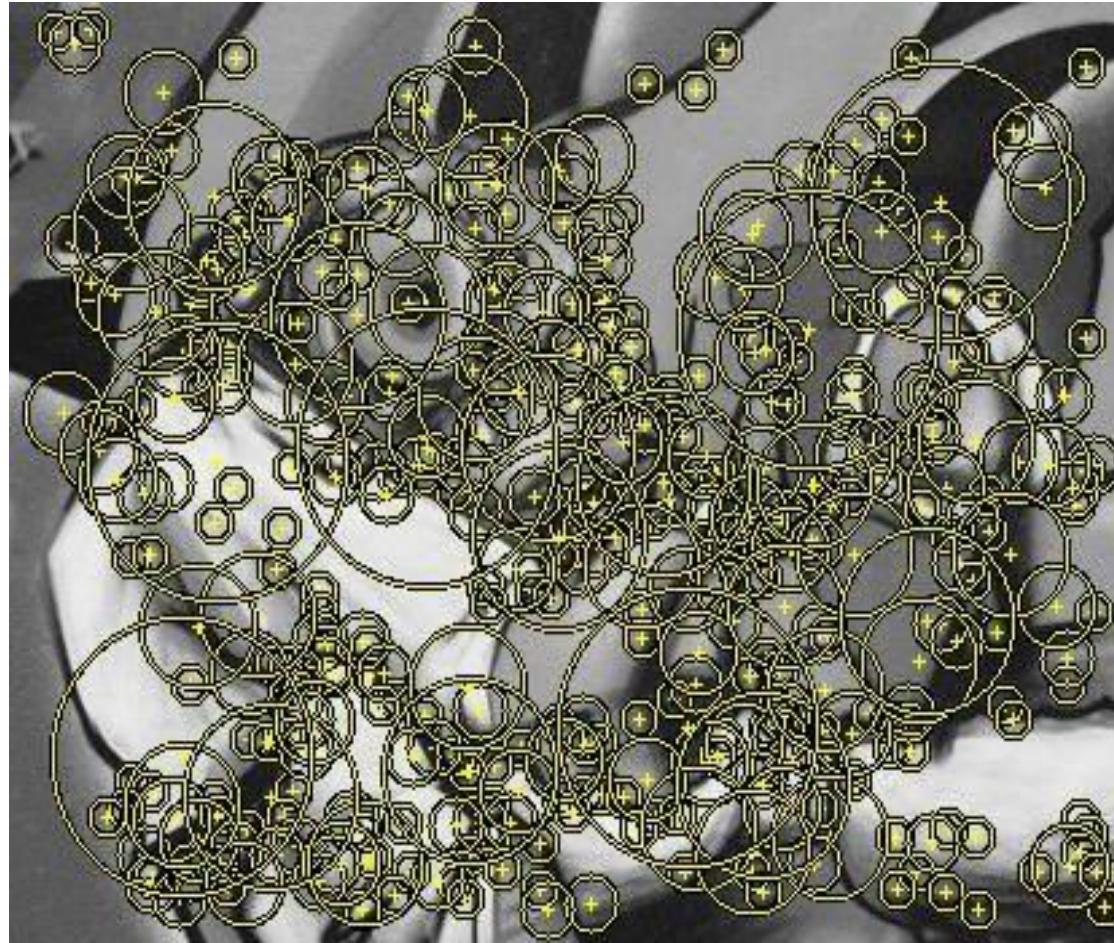
- Computation in Gaussian scale pyramid



Find local maxima in position-scale space of Difference-of-Gaussian



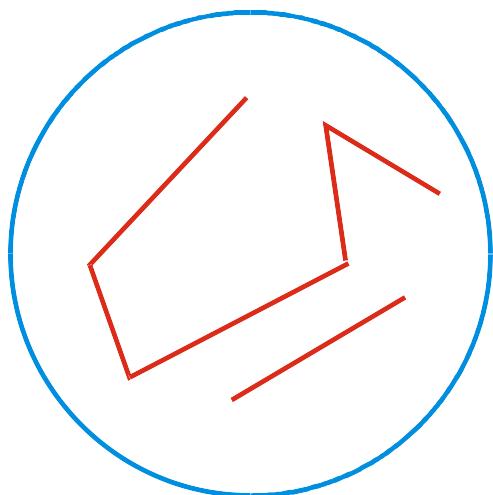
Results: Difference-of-Gaussian



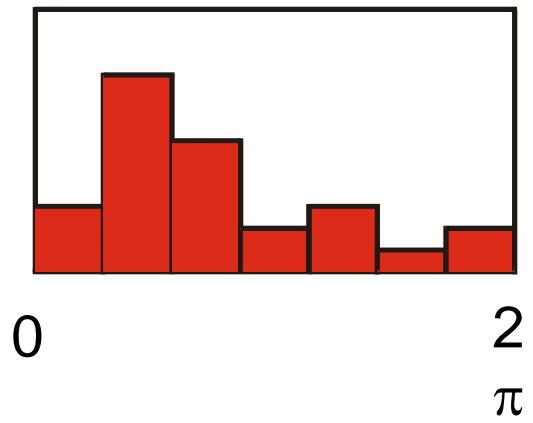
Orientation Normalization

[Lowe, SIFT, 1999]

- Compute orientation histogram



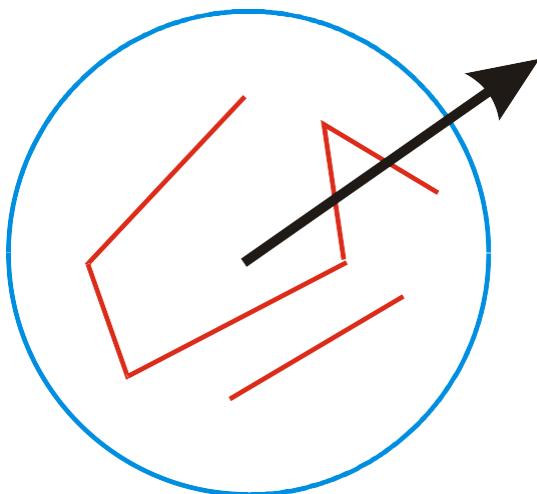
T. Tuytelaars, B. Leibe



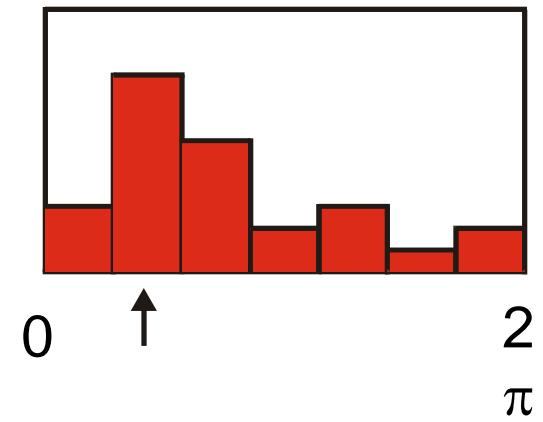
Orientation Normalization

[Lowe, SIFT, 1999]

- Compute orientation histogram
- Select dominant orientation



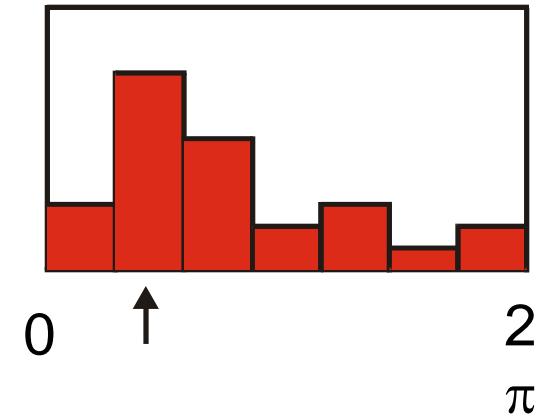
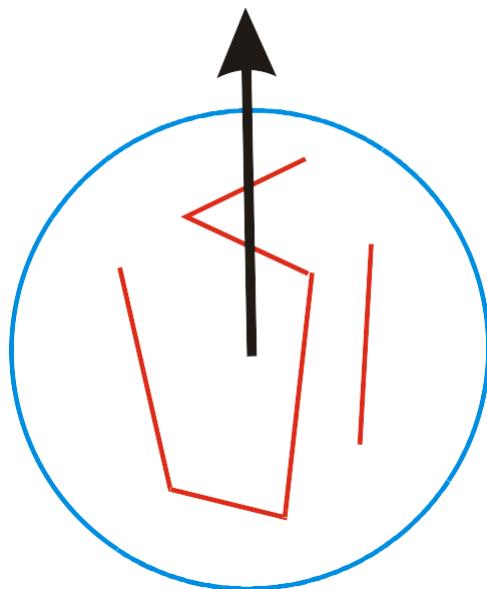
T. Tuytelaars, B. Leibe



Orientation Normalization

[Lowe, SIFT, 1999]

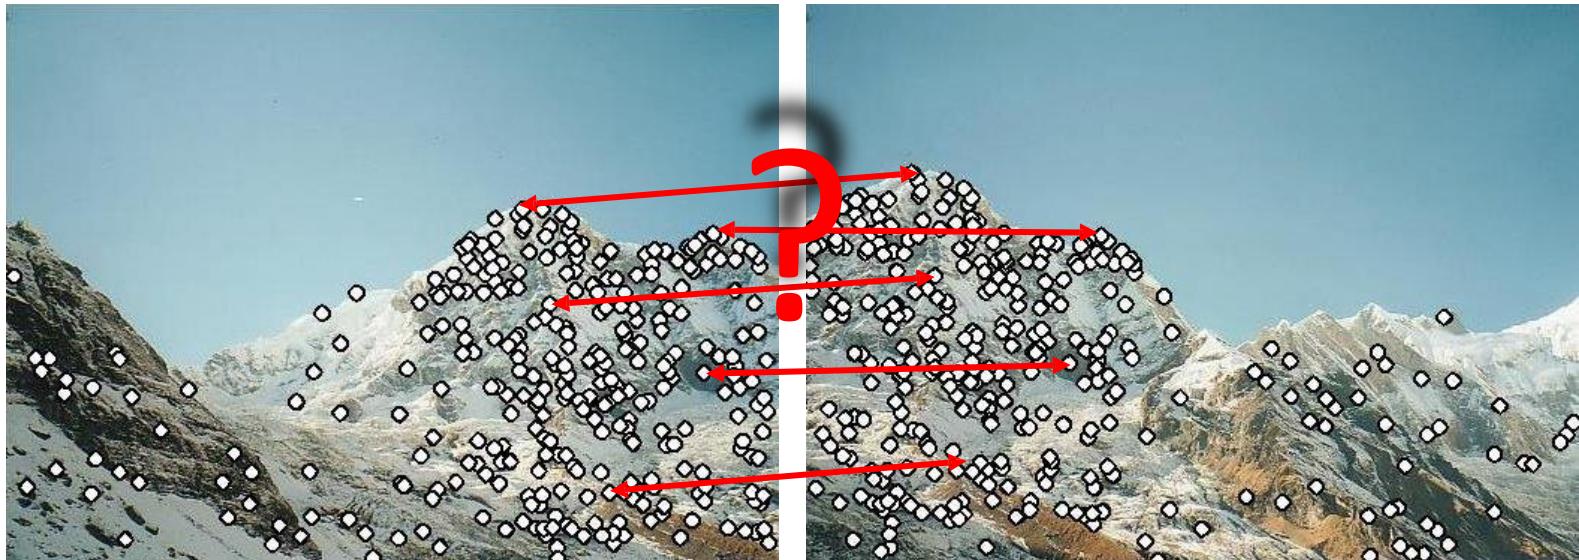
- Compute orientation histogram
- Select dominant orientation
- Normalize: rotate to fixed orientation



Feature descriptors

We know how to detect good points

Next question: **How to match them?**

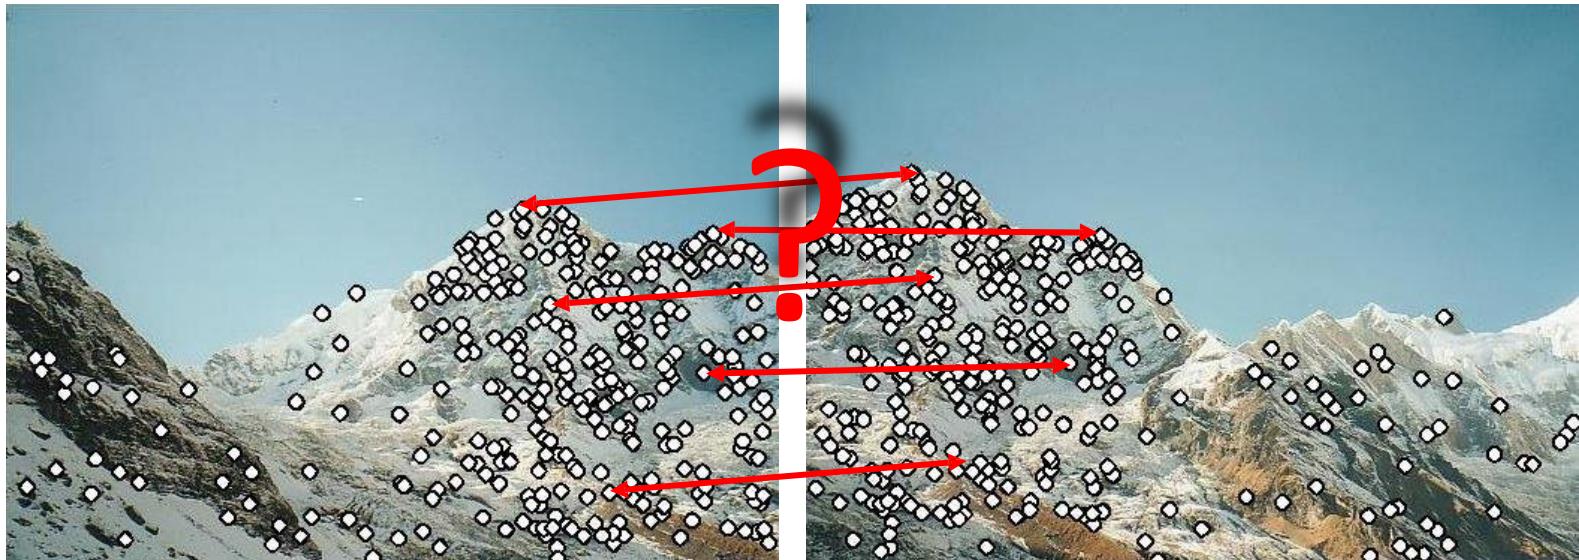


Answer: Come up with a *descriptor* for each point,
find similar descriptors between the two images

Feature descriptors

We know how to detect good points Next question:

How to match them?



Lots of possibilities (this is a popular research area)

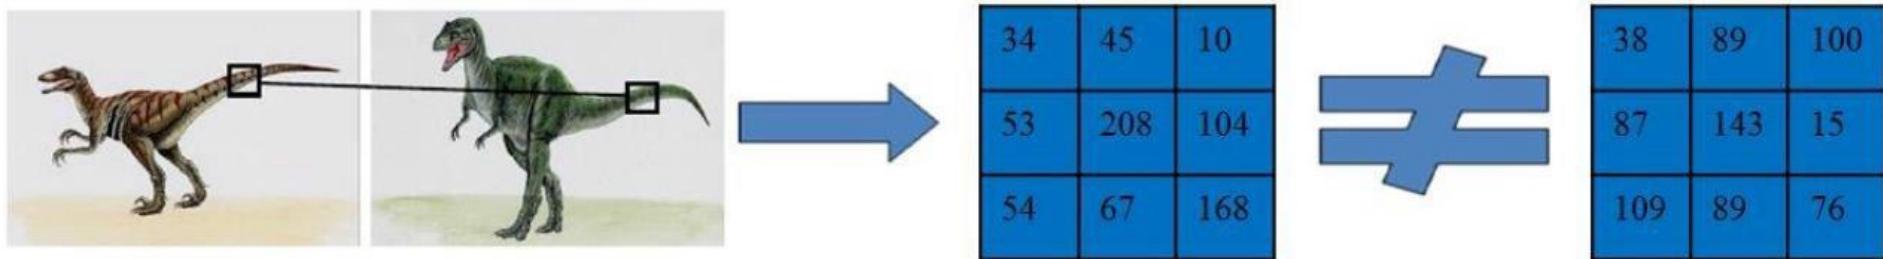
- Simple option: match square windows around the point
- State of the art approach: SIFT
 - David Lowe, UBC <http://www.cs.ubc.ca/~lowe/keypoints/>

Image/Region Matching

- Automatically recognize whether two images/regions contain the similar content.

Image/Region Matching

- Automatically recognize whether two images/regions contain the similar content.
- Comparing the image pixels as they are, will not work.



Image/Region Matching

- Pixel-based distances on high-dimensional data (and images especially) can be very unintuitive.

original



shifted



messed up



darkened



Challenges

Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter

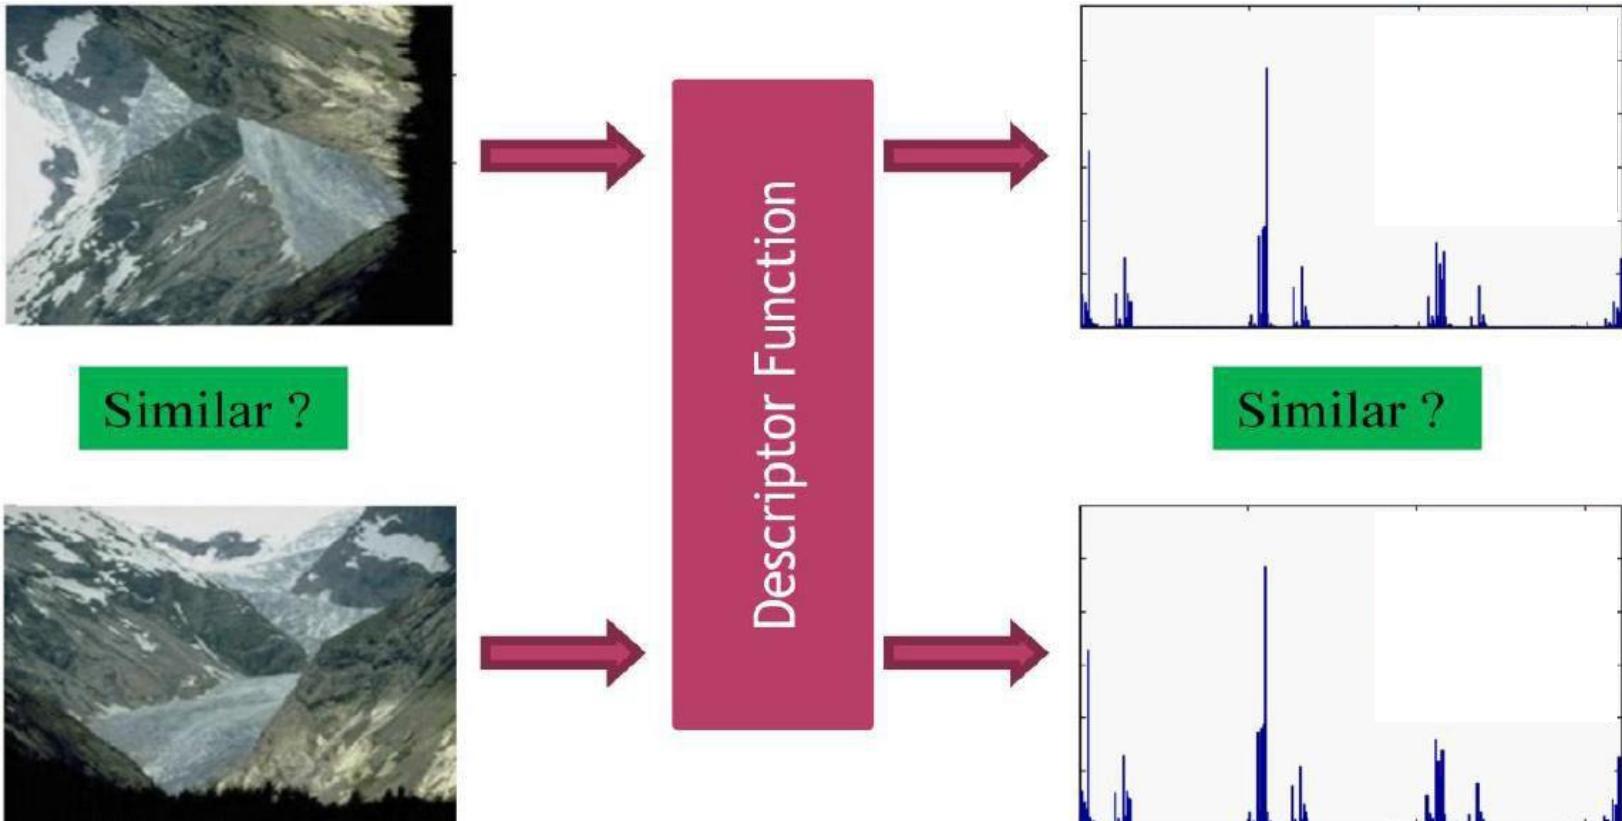


Intra-class variation



Solution

- Descriptors allow certain differences between the images.



Comparing using descriptor function
(images are taken from Corel-database and RSHD descriptor is used)

Applications

- Image Matching
- Image Retrieval
- Biomedical Image Analysis
- Texture Classification
- Image Correspondence
- Face Analysis
- Biometrics
- Building Panorama
- And many more...

Image descriptor

- Descriptions of the visual features
- Described by appearance based characteristics such as color, shape, etc.

Image descriptor

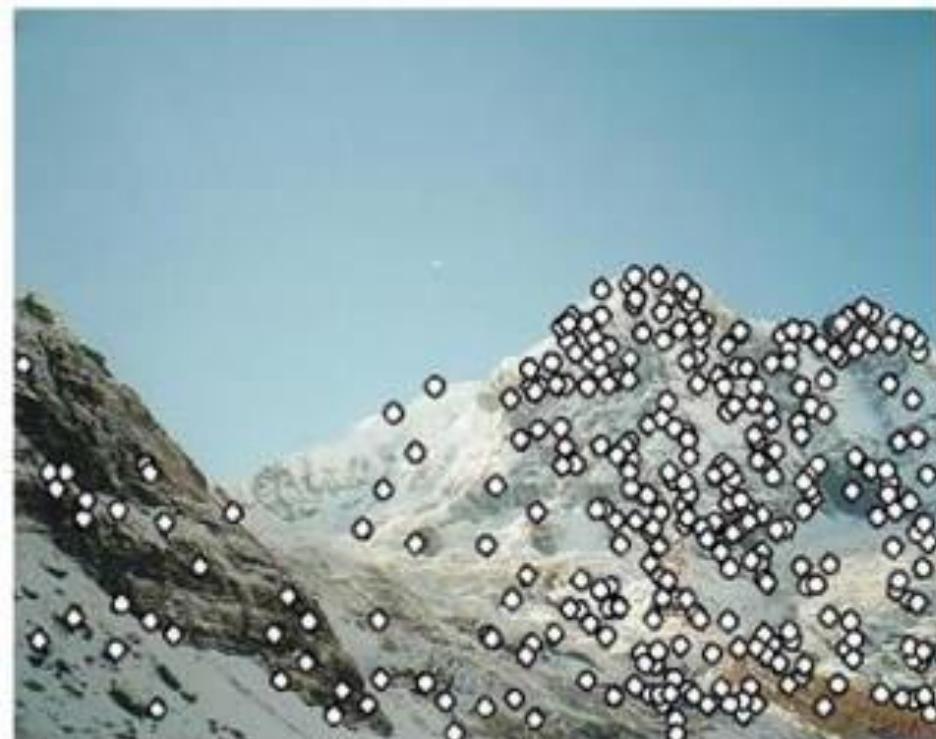
- Descriptions of the visual features
- Described by appearance based characteristics such as color, shape, etc.

A descriptor must be

- Distinctive
- Robust
- Compact
- Low Dimensional

Where to compute the descriptors?

- Over interest regions.
- Interest region may be
 - Key-Points or Global based.



Local Descriptors

- Most available descriptors focus on -
 - Edge/gradient information
 - Capture texture information
 - Exploit local relationship
 - Color also play a vital role
 - Shape features
 - Feature fusion

Widely Used Local Descriptors

- SIFT – Scale Invariant Feature Transform

Distinctive image features from scale-invariant keypoints

DG Lowe - International journal of computer vision, 2004 - Springer

... the assigned orientation, **scale**, and location for each **feature**, thereby providing **invariance** to these ... **Invariant Feature Transform** (SIFT), as it **transforms** image data into **scale-invariant** coordinates relative ... that densely cover the image over the full range of **scales** and locations ...

☆ 99 [Cited by 50252](#) Related articles All 179 versions

- LBP – Local Binary Pattern

Multiresolution gray-scale and rotation invariant texture classification with **local binary patterns**

T Ojala, M Pietikainen... - ... Transactions on pattern ..., 2002 - ieeexplore.ieee.org

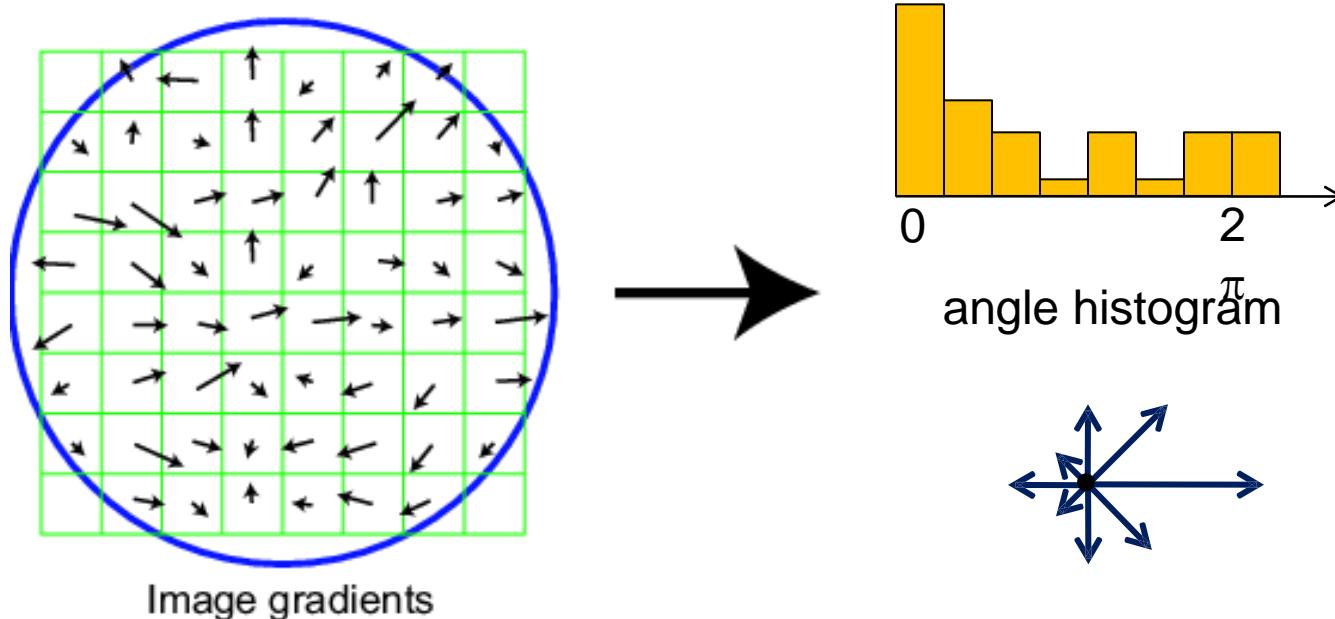
Presents a theoretically very simple, yet efficient, multiresolution approach to gray-scale and rotation invariant texture classification based on **local binary patterns** and nonparametric discrimination of sample and prototype distributions. The method is based on recognizing ...

☆ 99 [Cited by 12251](#) Related articles All 16 versions

Scale Invariant Feature Transform

Basic idea:

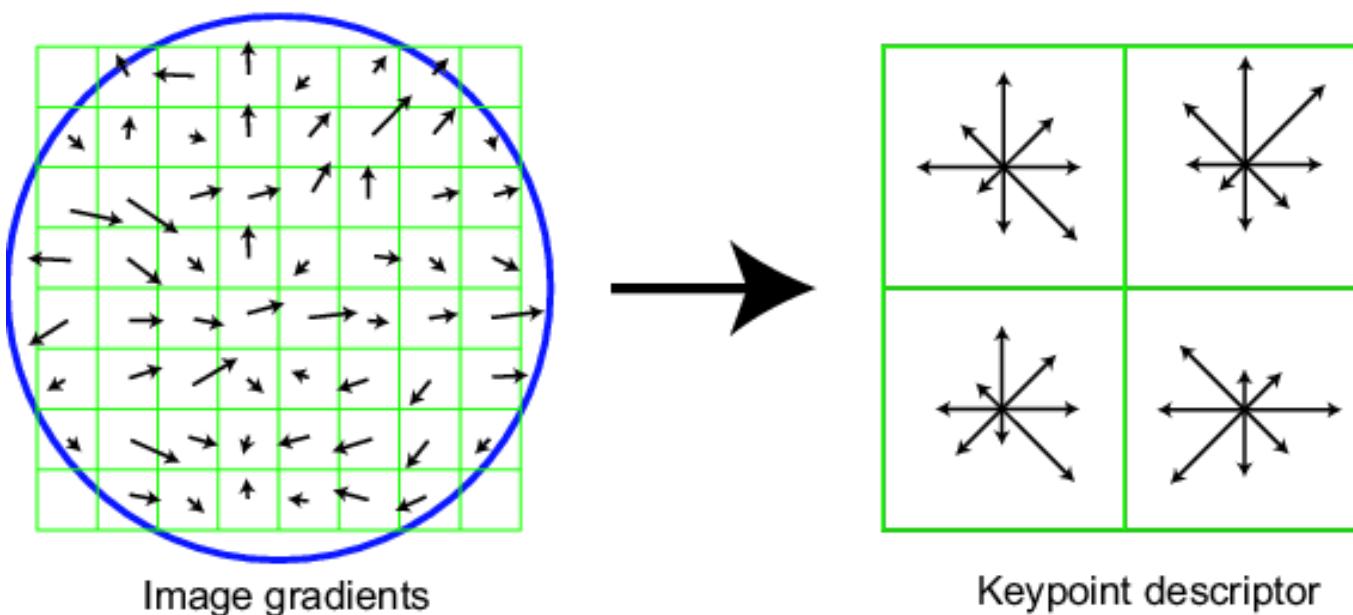
- Take 16x16 square window around detected feature
- Compute edge orientation for each pixel
- Throw out weak edges (threshold gradient magnitude)
- Create histogram of surviving edge orientations



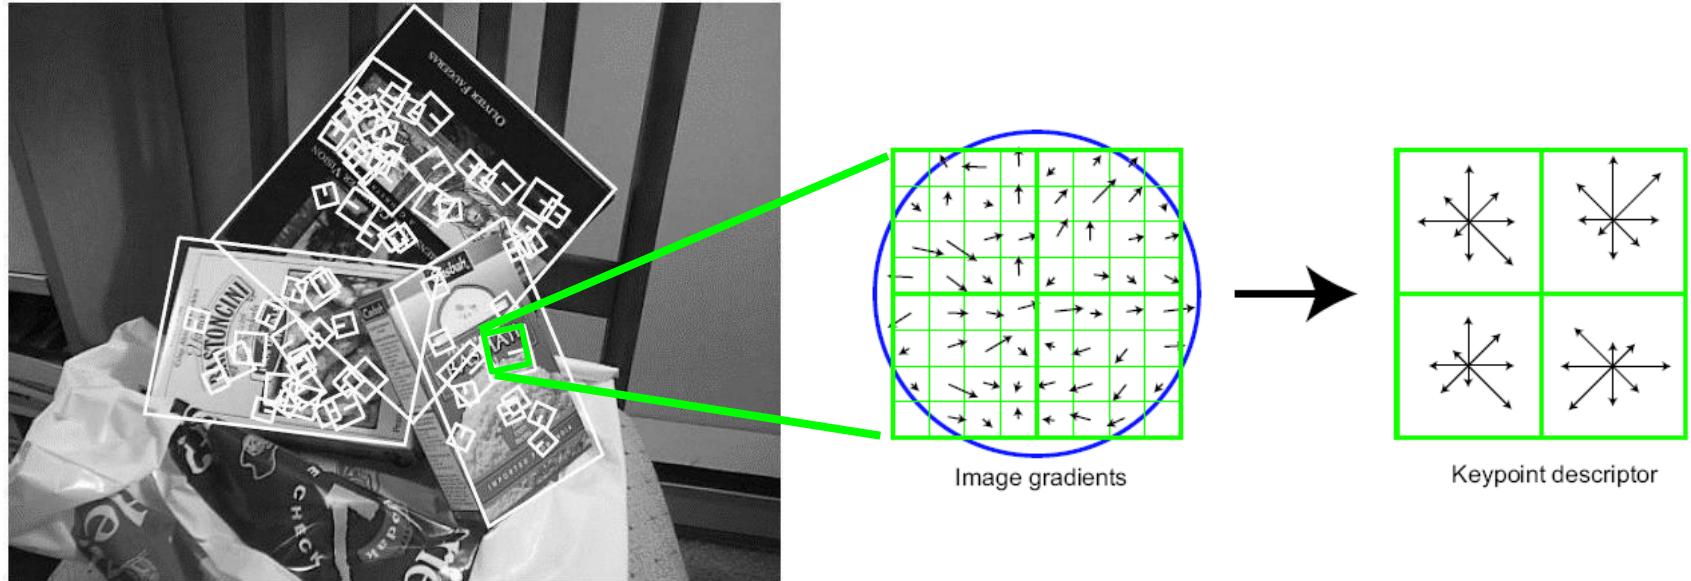
SIFT descriptor

Full version

- Divide the 16x16 window into a 4x4 grid of cells (2x2 case shown below)
- Compute an orientation histogram for each cell
- 16 cells * 8 orientations = 128 dimensional descriptor



Local Descriptors: SIFT Descriptor



**Histogram of oriented
gradients**

- Captures important texture information
- Robust to small translations / affine deformations

[Lowe, ICCV 1999]

Feature matching

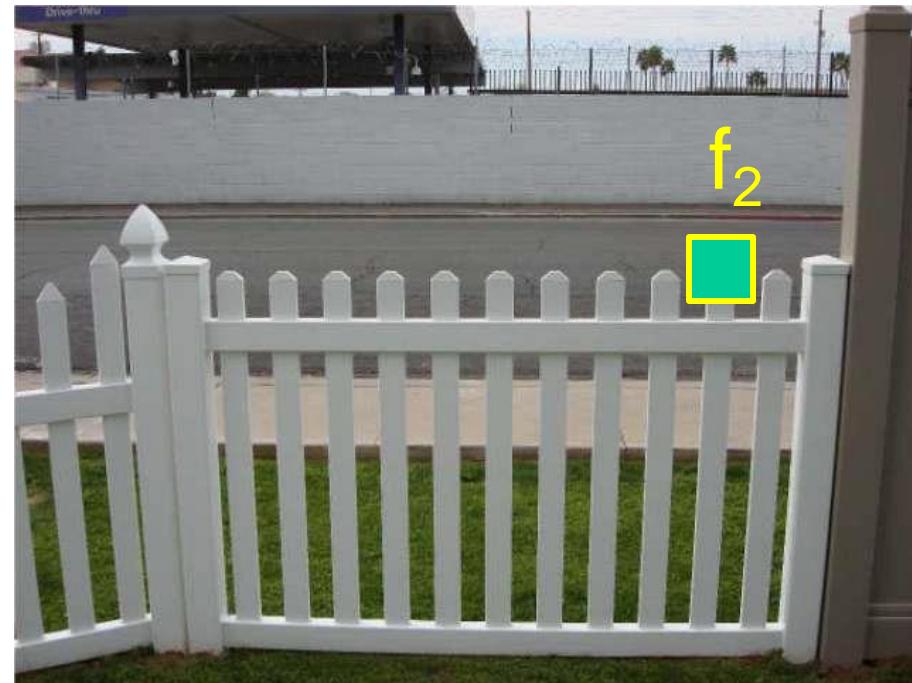
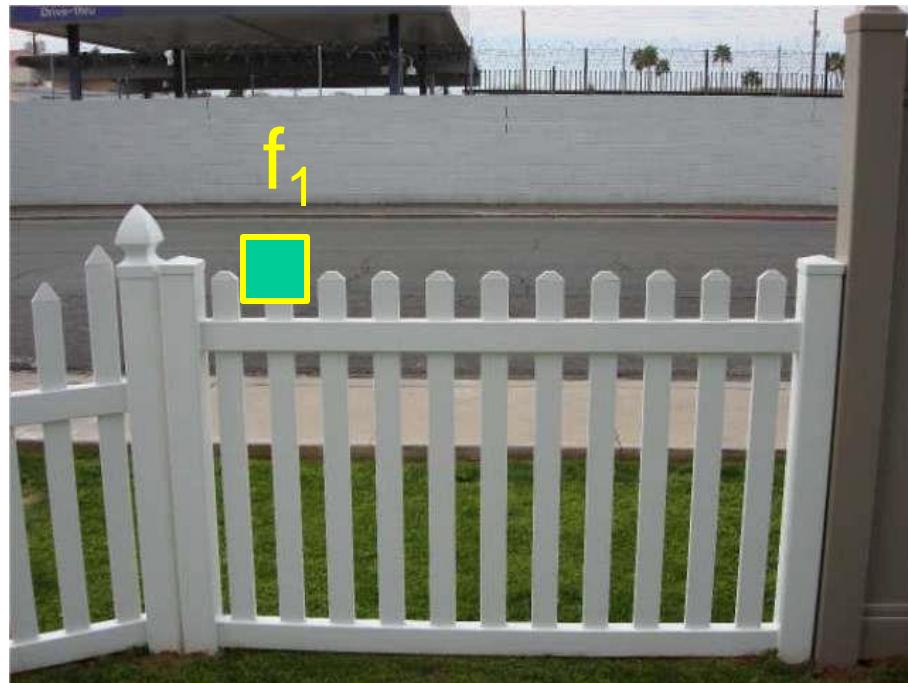
Given a feature in I_1 , how to find the best match in I_2 ?

1. Define distance function that compares two descriptors
2. Test all the features in I_2 , find the one with min distance

Feature distance

How to define the difference between two features f_1, f_2 ?

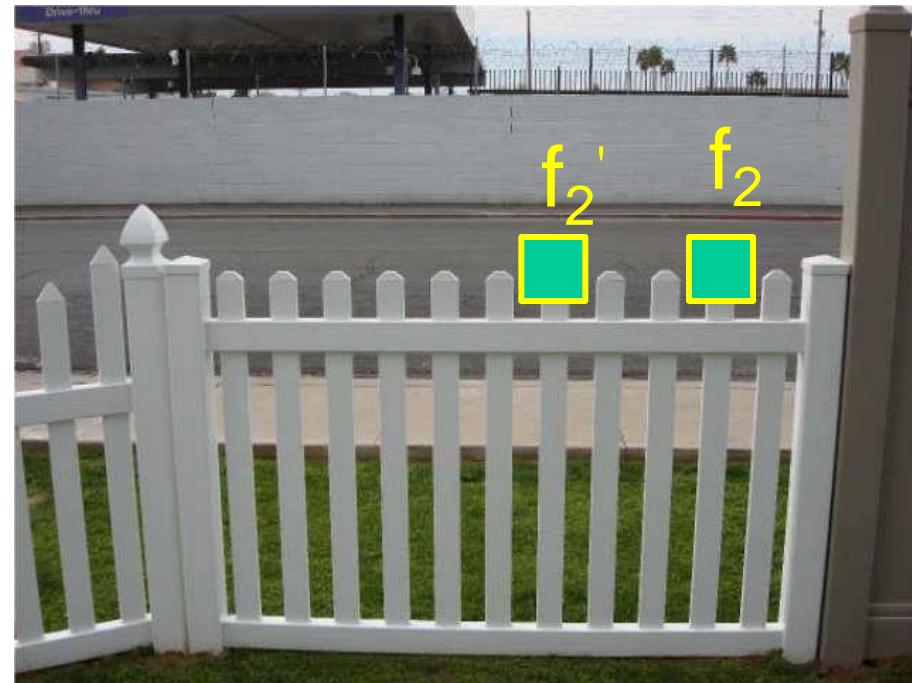
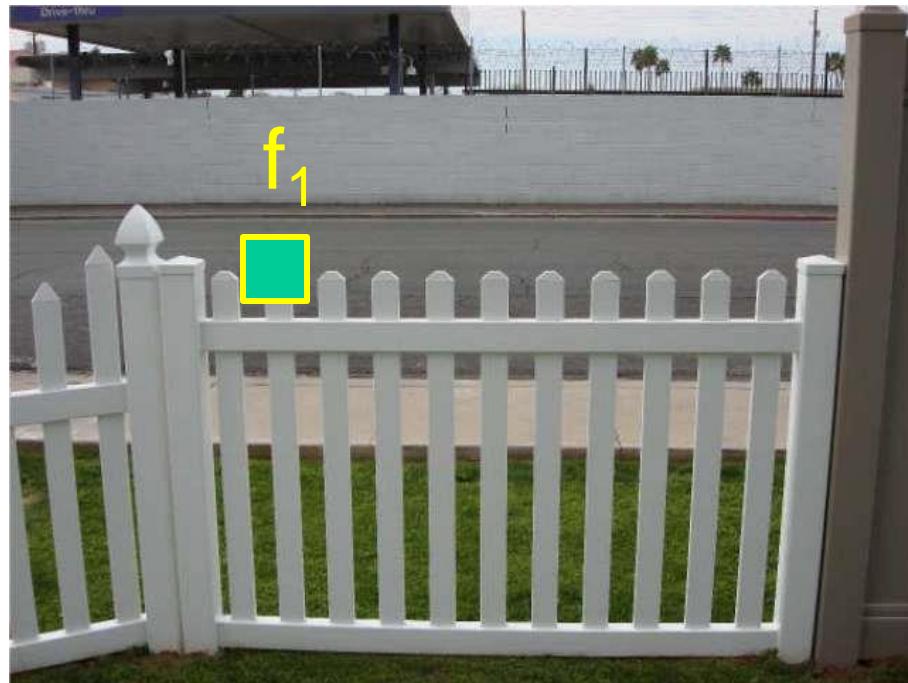
- Simple approach: L_2 distance, $\|f_1 - f_2\|$
- can give good scores to ambiguous (incorrect) matches



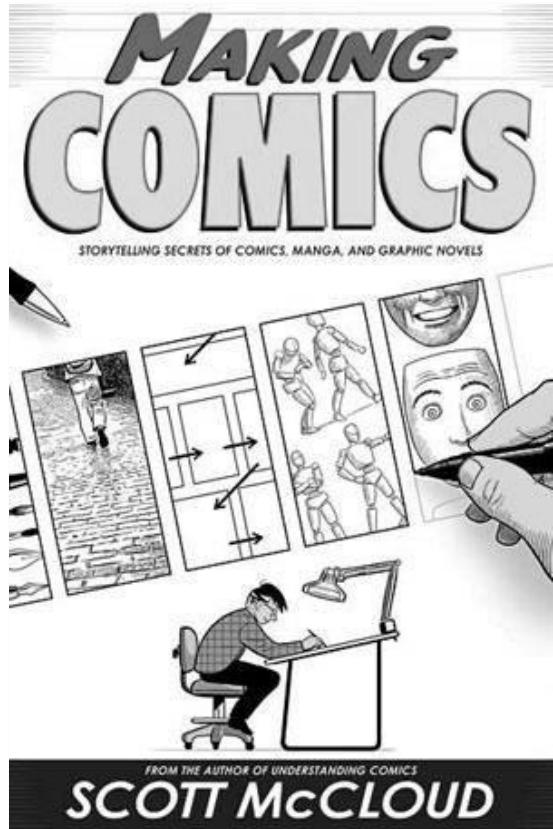
Feature distance

How to define the difference between two features f_1, f_2 ?

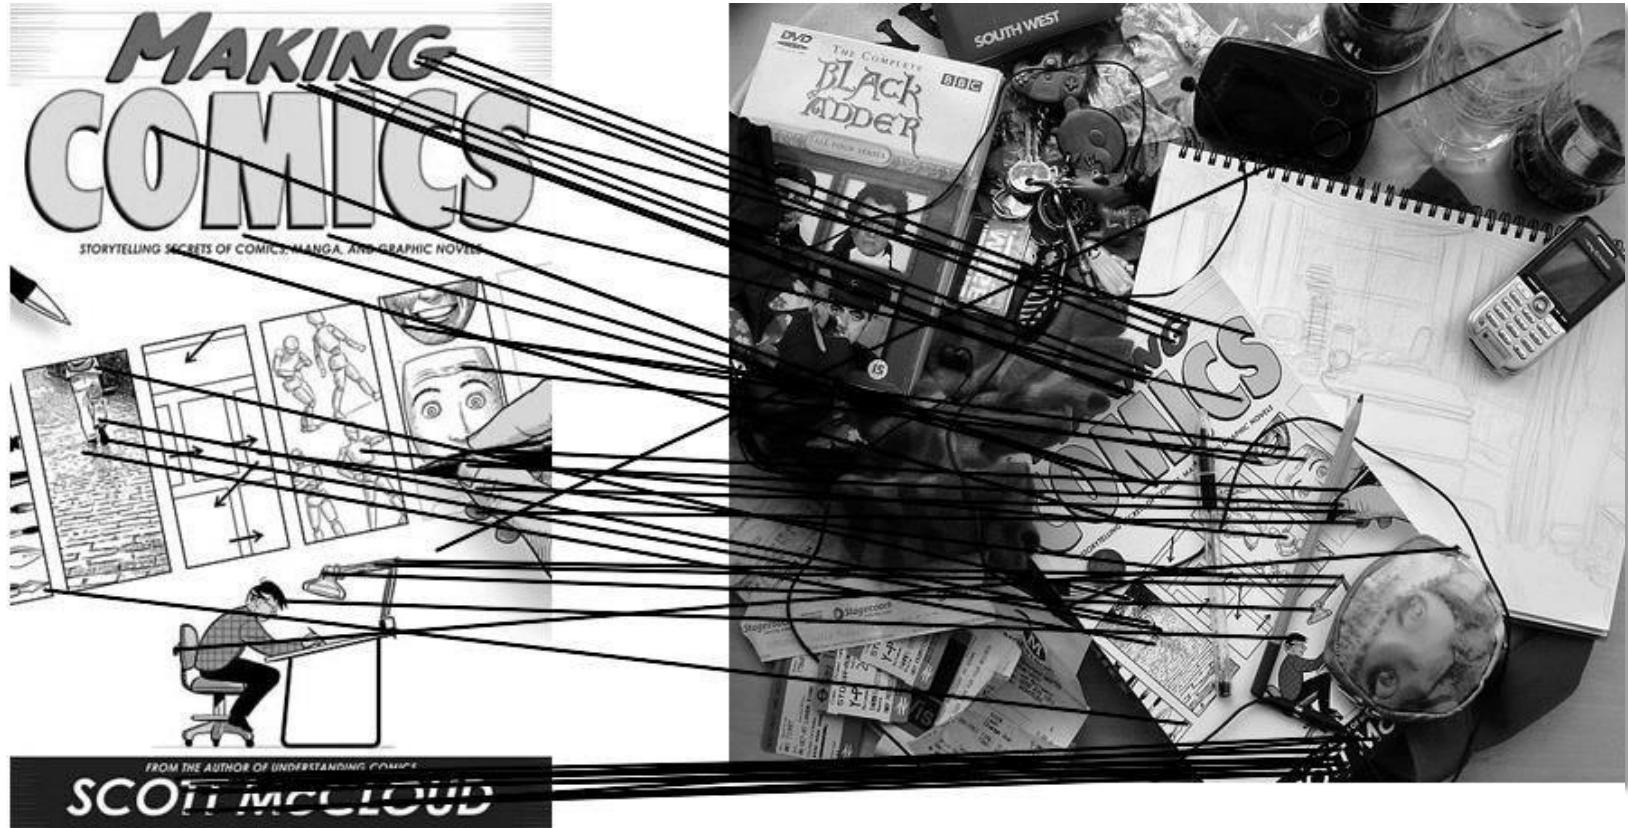
- Better approach: ratio distance = $\|f_1 - f_2\| / \|f_1 - f_2'\|$
 - f_2 is best SSD match to f_1 in I_2
 - f_2' is 2nd best SSD match to f_1 in I_2
 - gives large values for ambiguous matches



Feature matching example

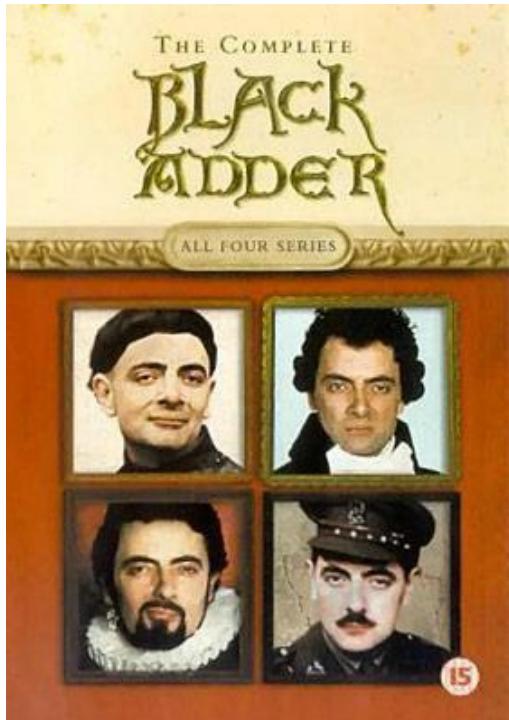


Feature matching example

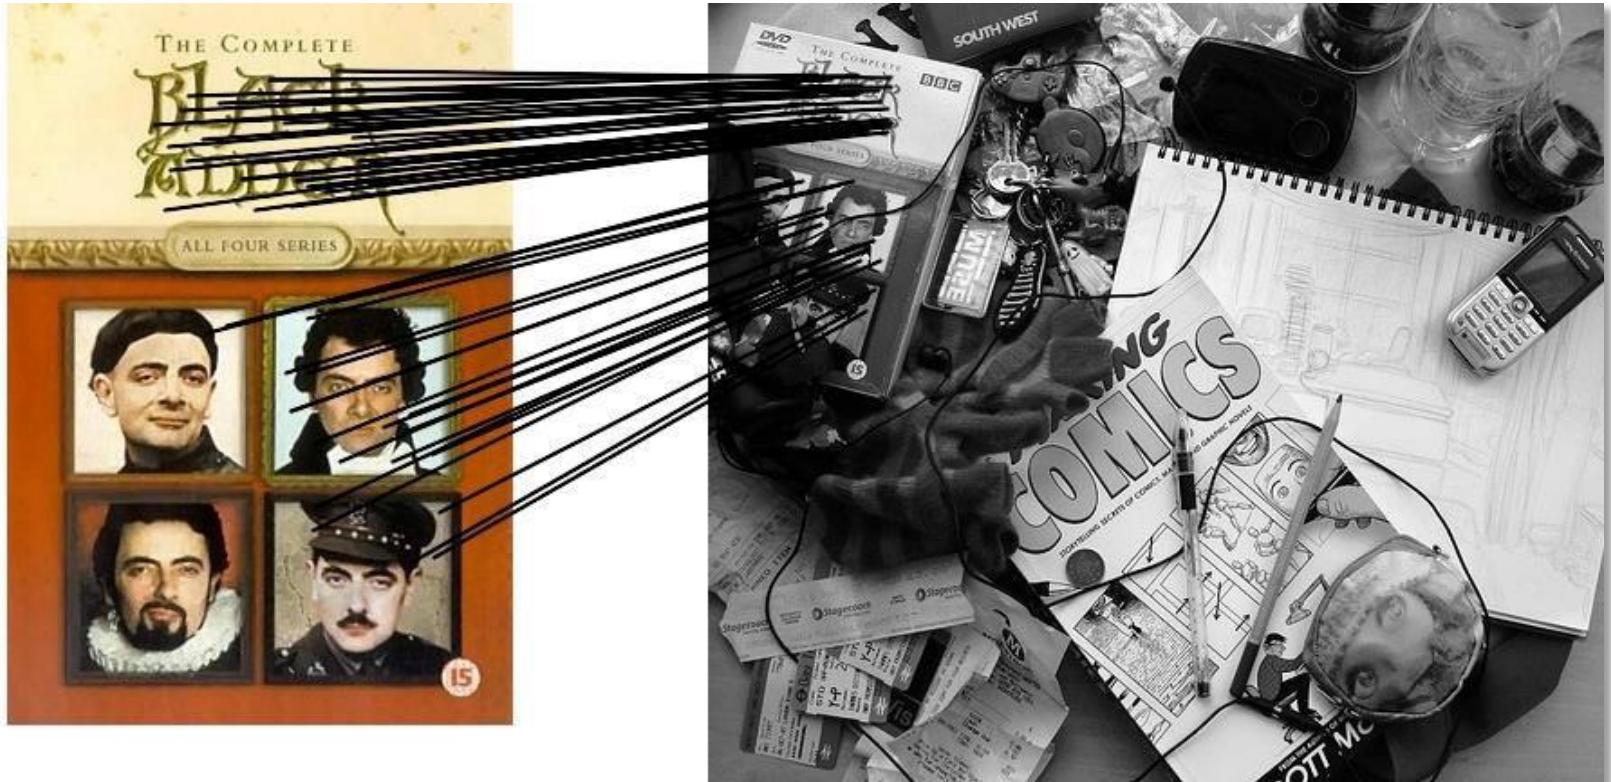


51 matches

Feature matching example



Feature matching example



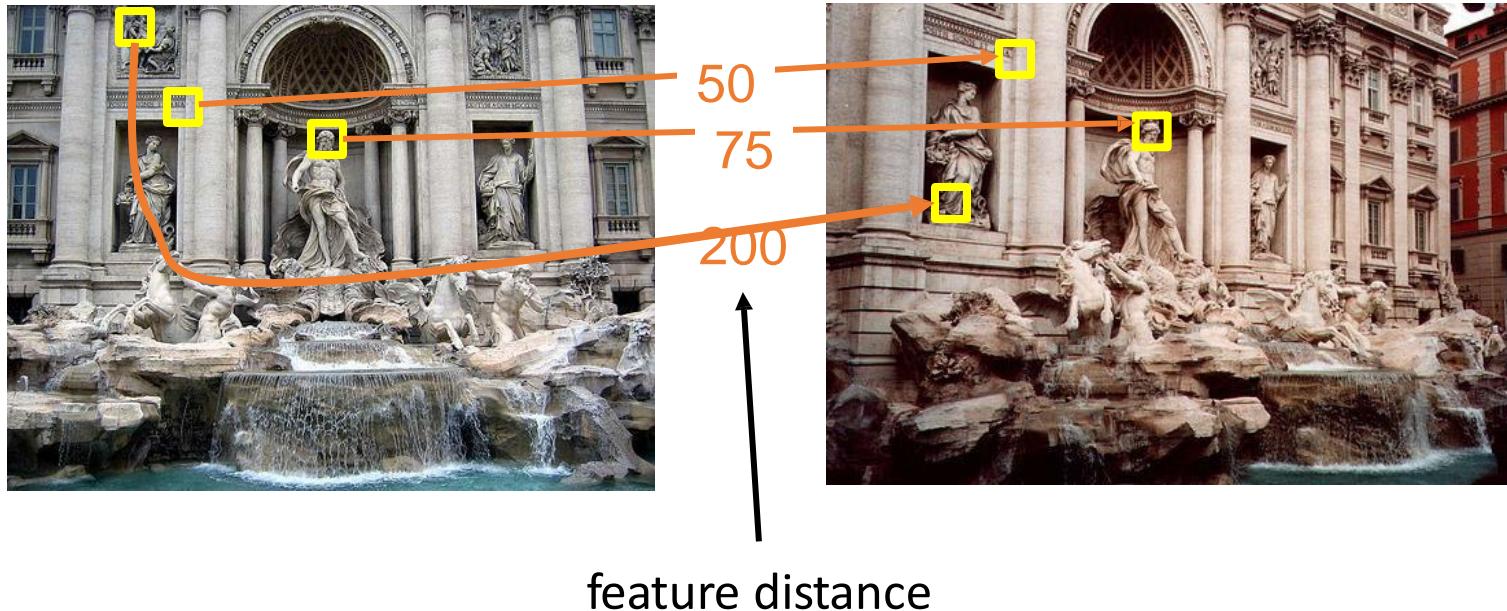
58 matches

Evaluating the results

How can we measure the performance of a feature matcher?

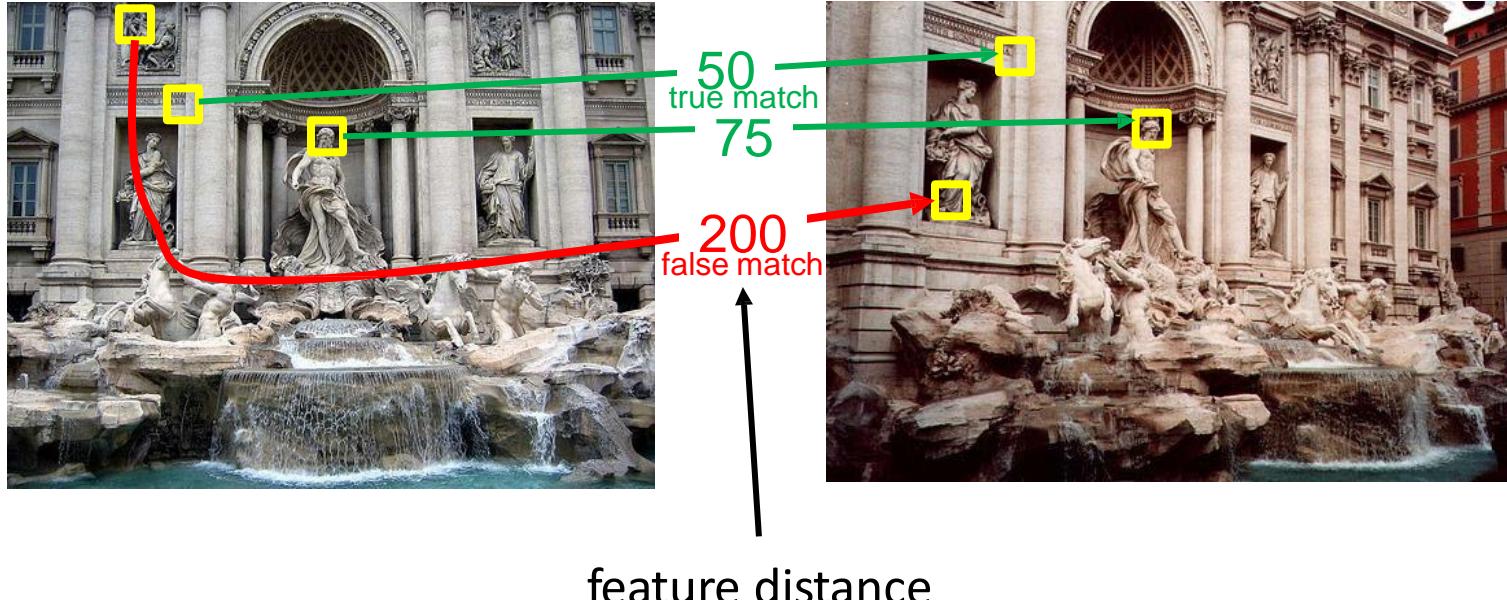
Evaluating the results

How can we measure the performance of a feature matcher?



True/false positives

How can we measure the performance of a feature matcher?



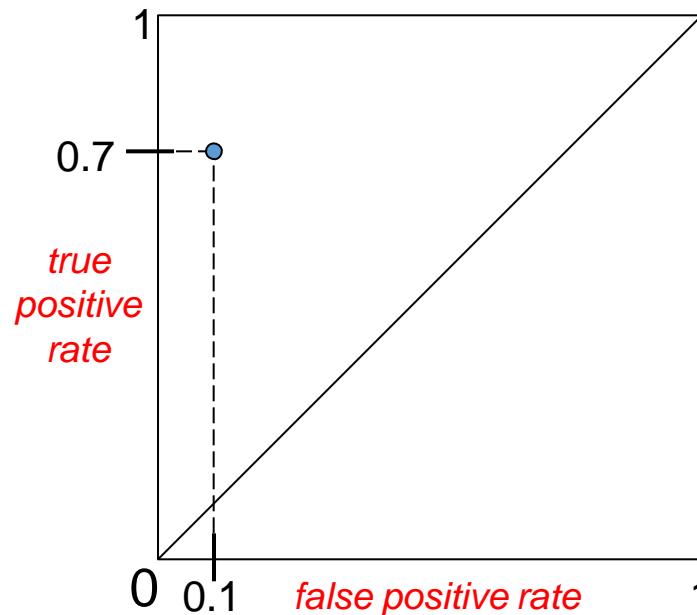
The distance threshold affects performance

- True positives = # of detected matches that are correct
 - Suppose we want to maximize these—how to choose threshold?
- False positives = # of detected matches that are incorrect
 - Suppose we want to minimize these—how to choose threshold?

Evaluating the results

How can we measure the performance of a feature matcher?

$$\frac{\text{\# true positives}}{\text{\# correctly matched features (positives)}} \quad \text{"recall"}$$

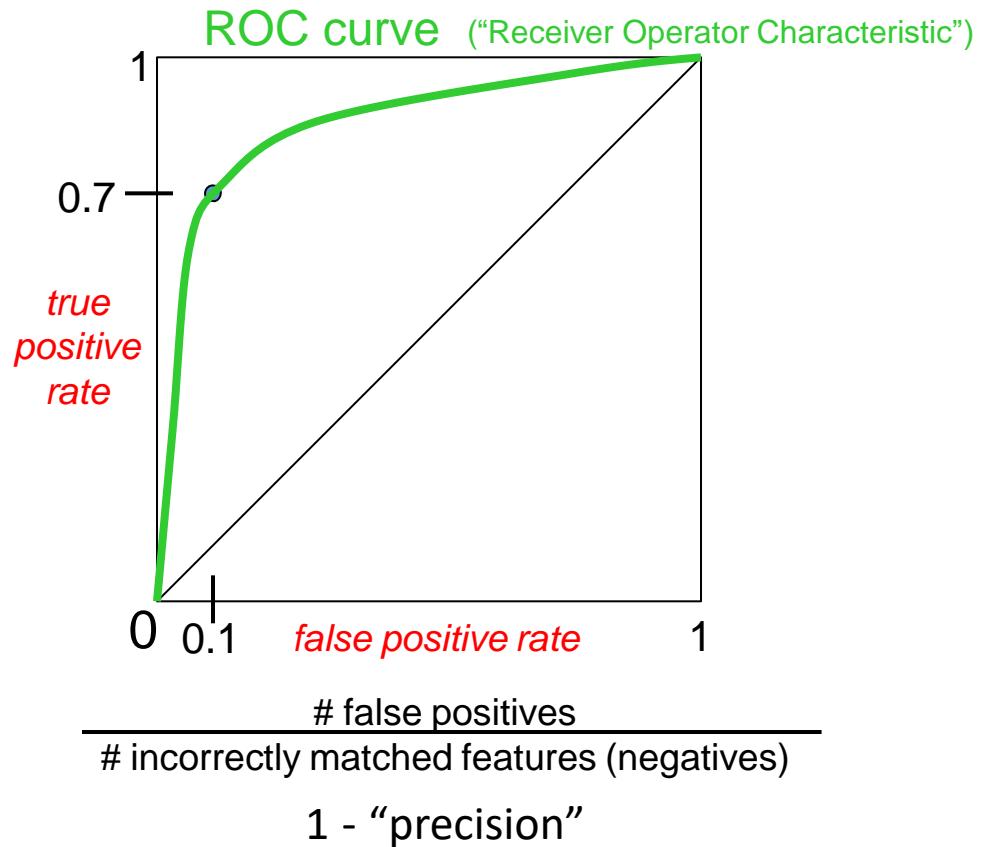


$$\frac{\text{\# false positives}}{\text{\# incorrectly matched features (negatives)}} \quad 1 - \text{"precision"}$$

Evaluating the results

How can we measure the performance of a feature matcher?

$$\frac{\# \text{ true positives}}{\# \text{ correctly matched features} \text{ (positives)}} \text{ "recall"}$$



Evaluating the results

- ROC Curves summarize the trade-off between the true positive rate and false positive rate for a predictive model using different probability thresholds.
- Precision-Recall curves summarize the trade-off between the true positive rate and the positive predictive value for a predictive model using different probability thresholds.
- ROC curves are appropriate when the observations are balanced between each class, whereas precision-recall curves are appropriate for imbalanced datasets.

Local Binary Pattern (LBP)

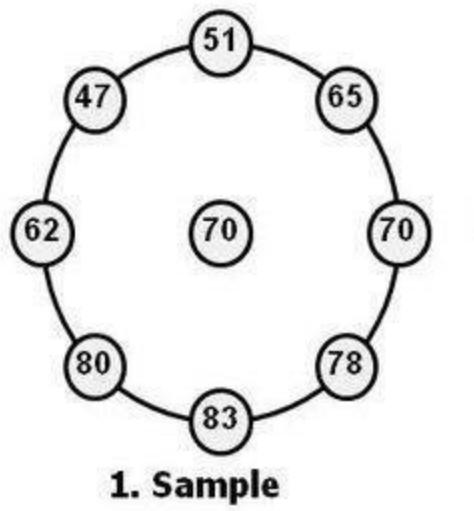


Image Source: scholarpedia

[IEEE TPAMI 2002]

Local Binary Pattern (LBP)

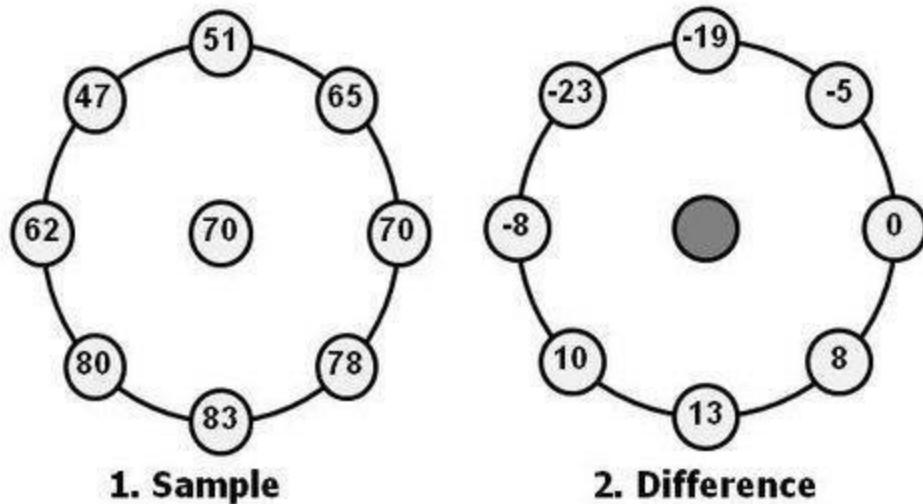


Image Source: scholarpedia

[IEEE TPAMI 2002]

Local Binary Pattern (LBP)

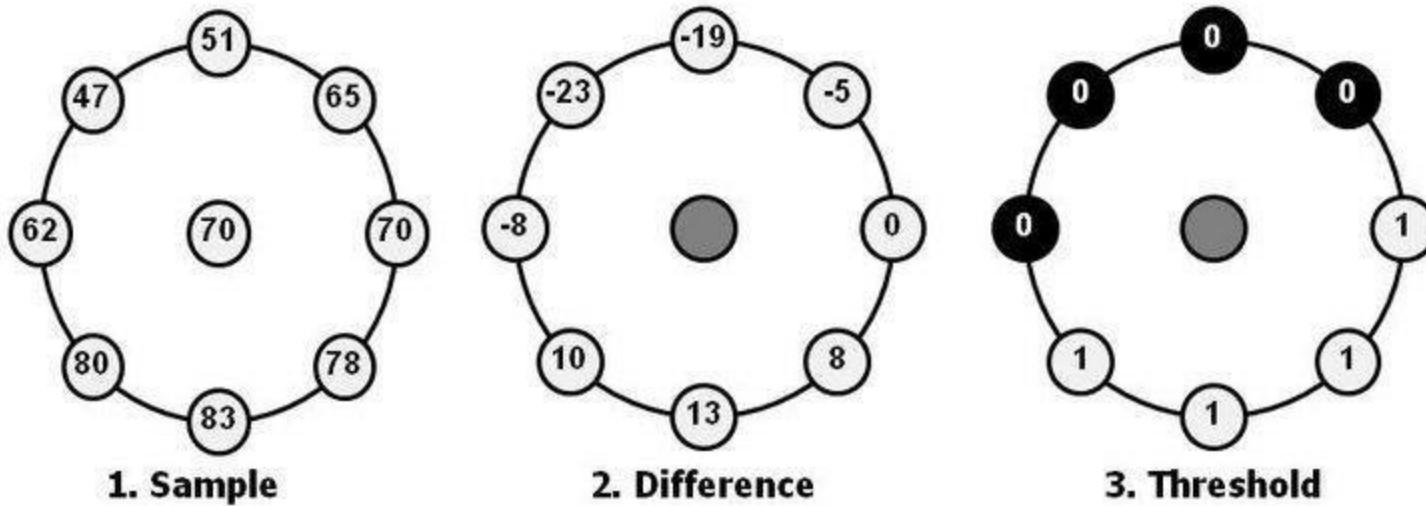


Image Source: scholarpedia

[IEEE TPAMI 2002]

Local Binary Pattern (LBP)

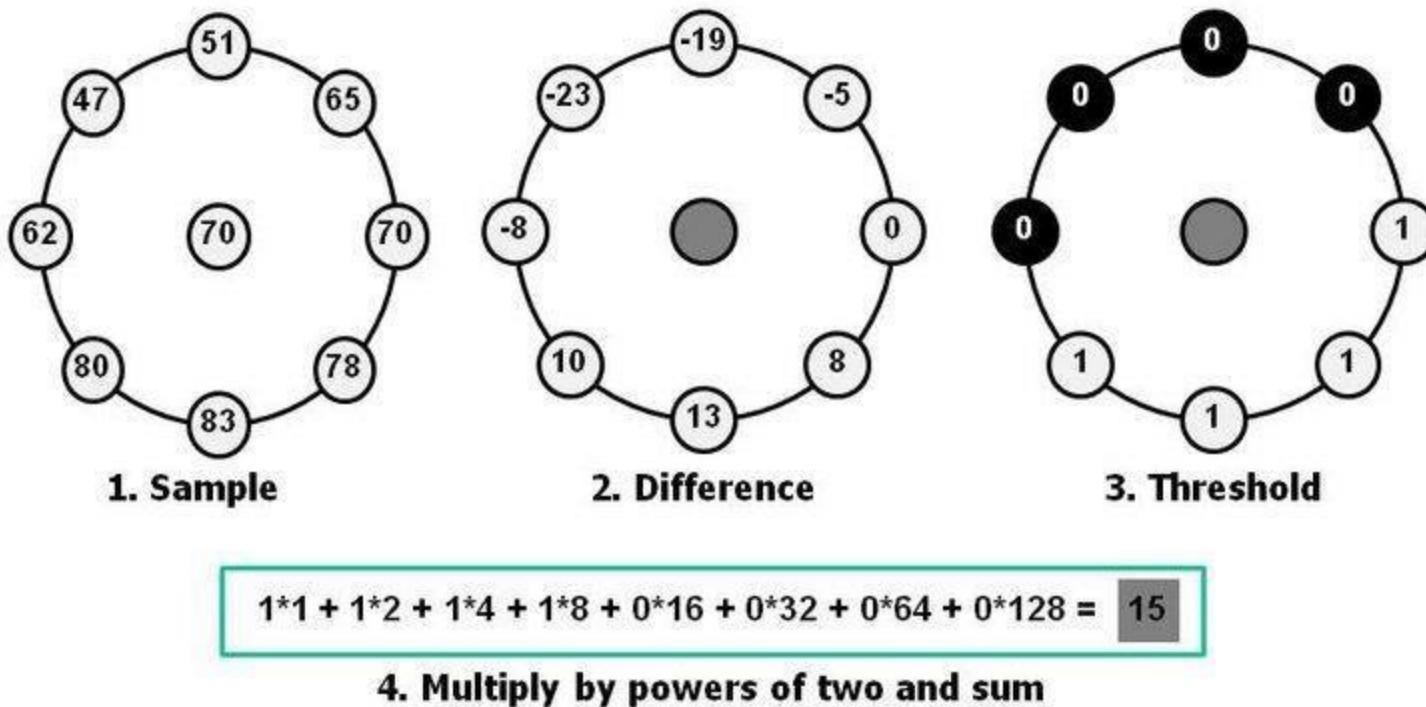


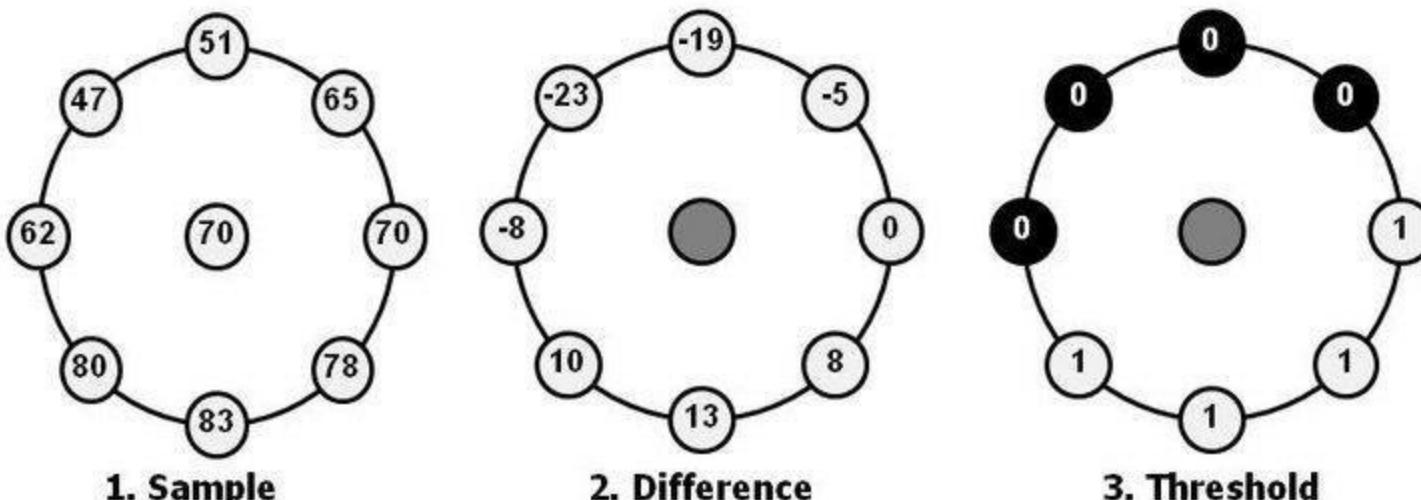
Image Source: scholarpedia

[IEEE TPAMI 2002]

Local Binary Pattern (LBP)

The value of the LBP code of a pixel (x_c, y_c) is given by:

$$LBP_{P,R} = \sum_{p=0}^{P-1} s(g_p - g_c)2^p \quad s(x) = \begin{cases} 1, & \text{if } x \geq 0; \\ 0, & \text{otherwise.} \end{cases}$$



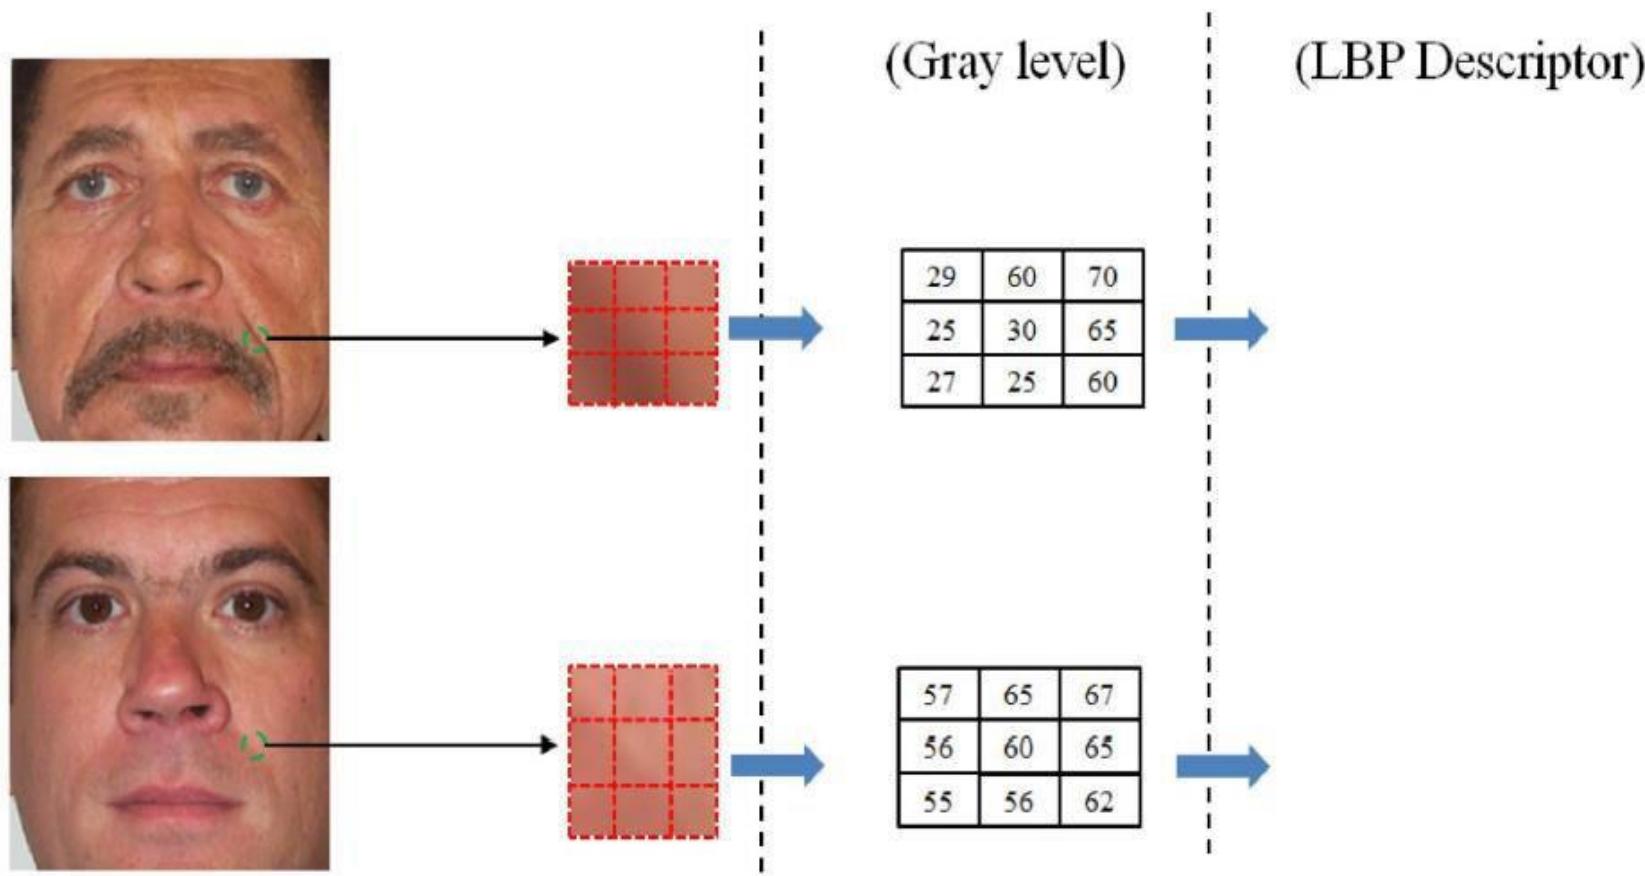
$$1*1 + 1*2 + 1*4 + 1*8 + 0*16 + 0*32 + 0*64 + 0*128 = 15$$

4. Multiply by powers of two and sum

Image Source: scholarpedia

[IEEE TPAMI 2002]

Local Binary Pattern (LBP)



*Image Source: Nguyen, D. T., Cho, S. R., & Park, K. R. (2015). Age Estimation-Based Soft Biometrics Considering Optical Blurring Based on Symmetrical Sub-Blocks for MLBP. *Symmetry*, 7(4), 1882-1913.*

Local Binary Pattern (LBP)

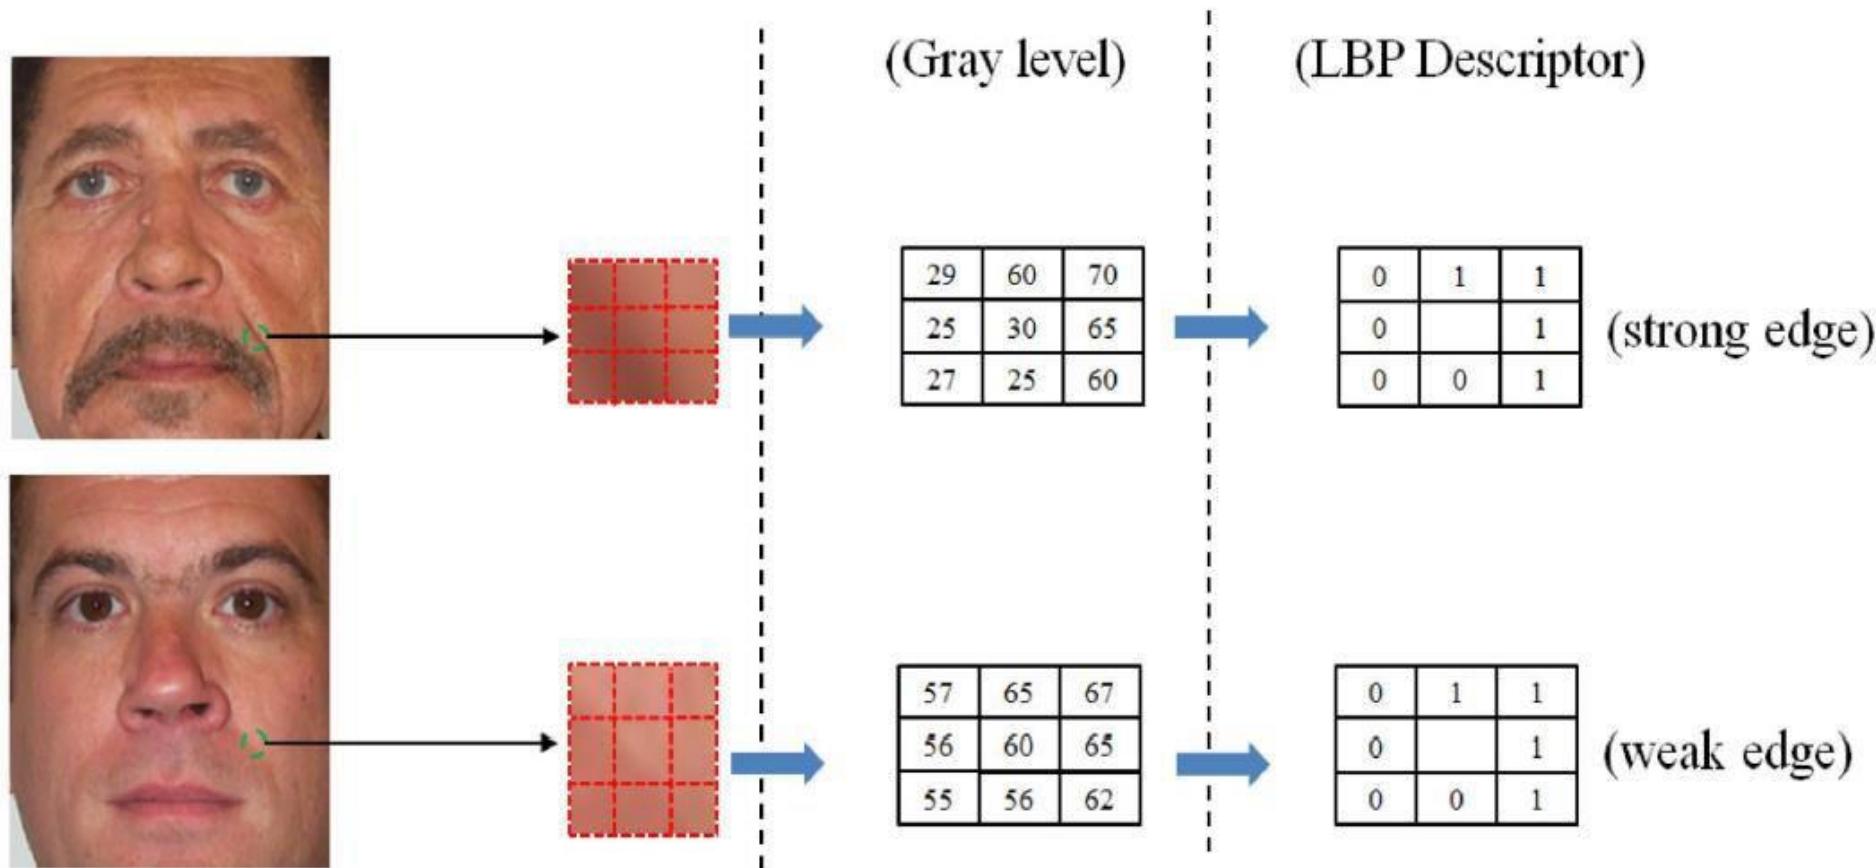


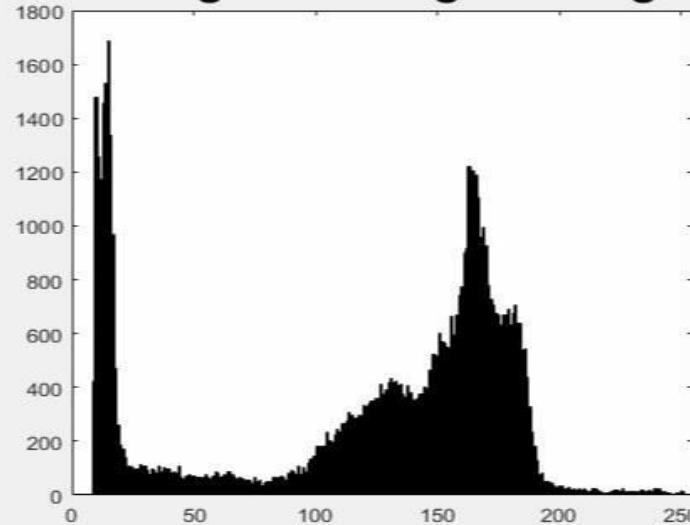
Image Source: Nguyen, D. T., Cho, S. R., & Park, K. R. (2015). Age Estimation-Based Soft Biometrics Considering Optical Blurring Based on Symmetrical Sub-Blocks for MLBP. *Symmetry*, 7(4), 1882-1913.

Local Binary Pattern (LBP)

Original Grayscale Image



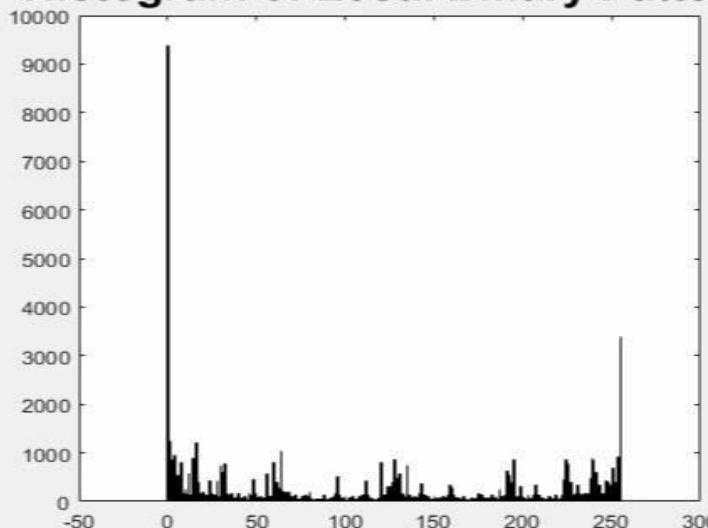
Histogram of original image



Local Binary Pattern



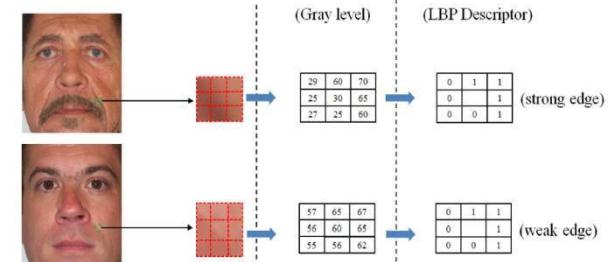
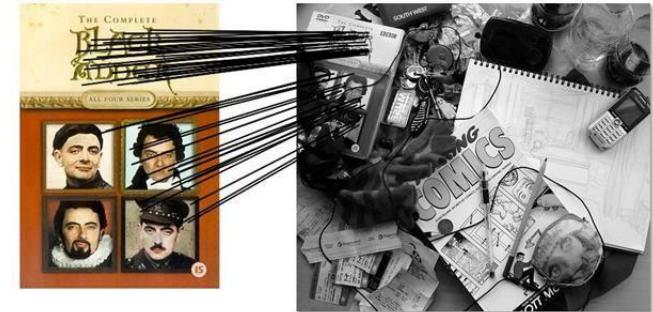
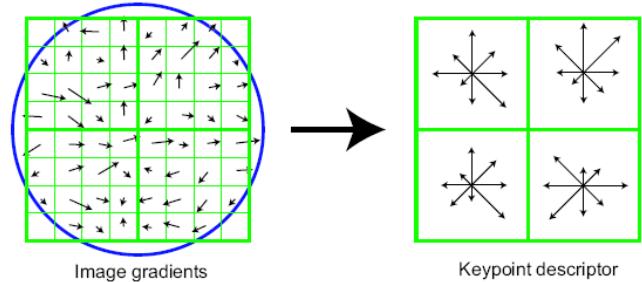
Histogram of Local Binary Pattern



*Image Source:
Mathworks*

Things to remember

- Region detection: repeatable and distinctive
 - Blobs, DoG, stable regions
 - Invariant to affine transformation
- Local Descriptors:
 - Discriminative
 - Robust
 - Compact
 - Efficient
- Scale Invariant Feature Transform
 - Utilizes gradient information
 - Gradient direction binning
- Local Binary Pattern
 - Exploits local relationship
 - Captures edge, spot, corner, etc.

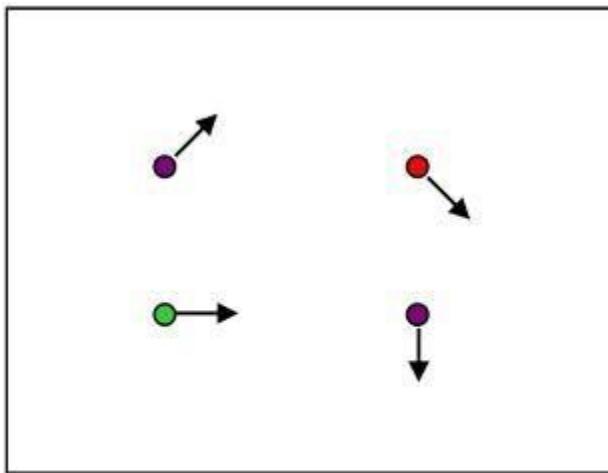


Acknowledgements

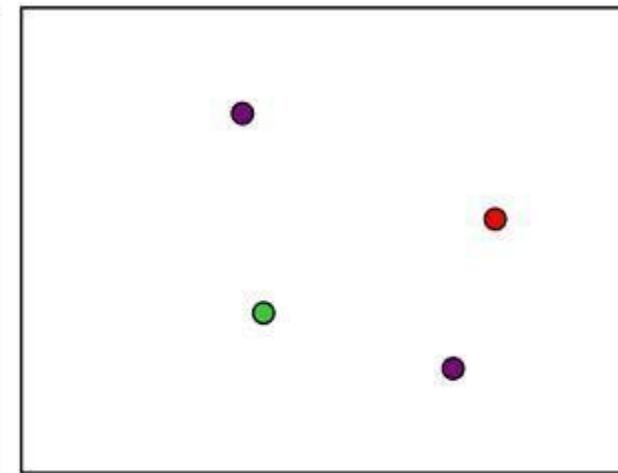
- Thanks to the following researchers for making their teaching/research material online
 - Forsyth
 - Steve Seitz
 - Noah Snavely
 - J.B. Huang
 - Derek Hoiem
 - D. Lowe
 - A. Bobick
 - S. Lazebnik
 - K. Grauman
 - R. Zaleski
 - Leibe

Thank you

- Next class: feature tracking and optical flow



$$I(x,y,t)$$



$$I(x,y,t+1)$$