

Introduction to Parallel Computing

Instructor


Dr. B Krishna Priya

Contents

- Introduction to Parallel computing
- Motivating Parallelism
- Scope of Parallelism
- Parallel Programming Platforms
- Implicit Parallelism: Trends in Microprocessor Architectures
- Limitations of Memory System Performance
- Dichotomy of Parallel Computing Platforms
- Physical Organization of Parallel Platforms

What is parallel computing

- Traditional software has been written for serial computing.
- To be run on single computer with single CPU




Limitations of Serial Computing

- Limits to serial computing - both physical and practical reasons pose significant constraints to simply building ever faster serial computers.
- Transmission speeds - the speed of a serial computer is directly dependent upon how fast data can move through hardware.
- Limits to miniaturization - processor technology is allowing an increasing number of transistors to be placed on a chip. However, even with molecular or atomic-level components, a limit will be reached on how small components can be.
- Economic limitations - it is increasingly expensive to make a single processor faster. Using a larger number of moderately fast commodity processors to achieve the same (or better) performance is less expensive.


What is parallel computing

- In the simplest sense, parallel computing is the simultaneously use of multiple computing resource to solve computational problem.



Parallel computing

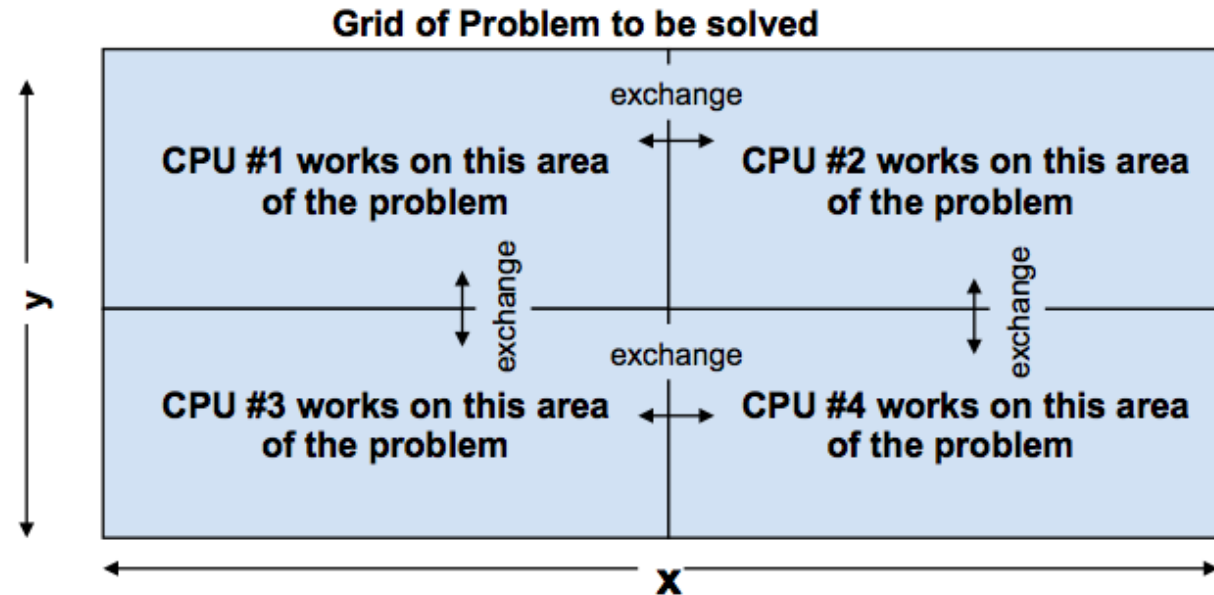
- A parallel computing is a “Collection of processing elements that co-operate and communicate to solve large Problem”
- Basic Principle: A computation can be divided into smaller subproblem each of which can be solved simultaneously.
- Assumption: Parallel computing assumes that a parallel hardware is present at our disposal, which is capable of executing these computations in parallel.



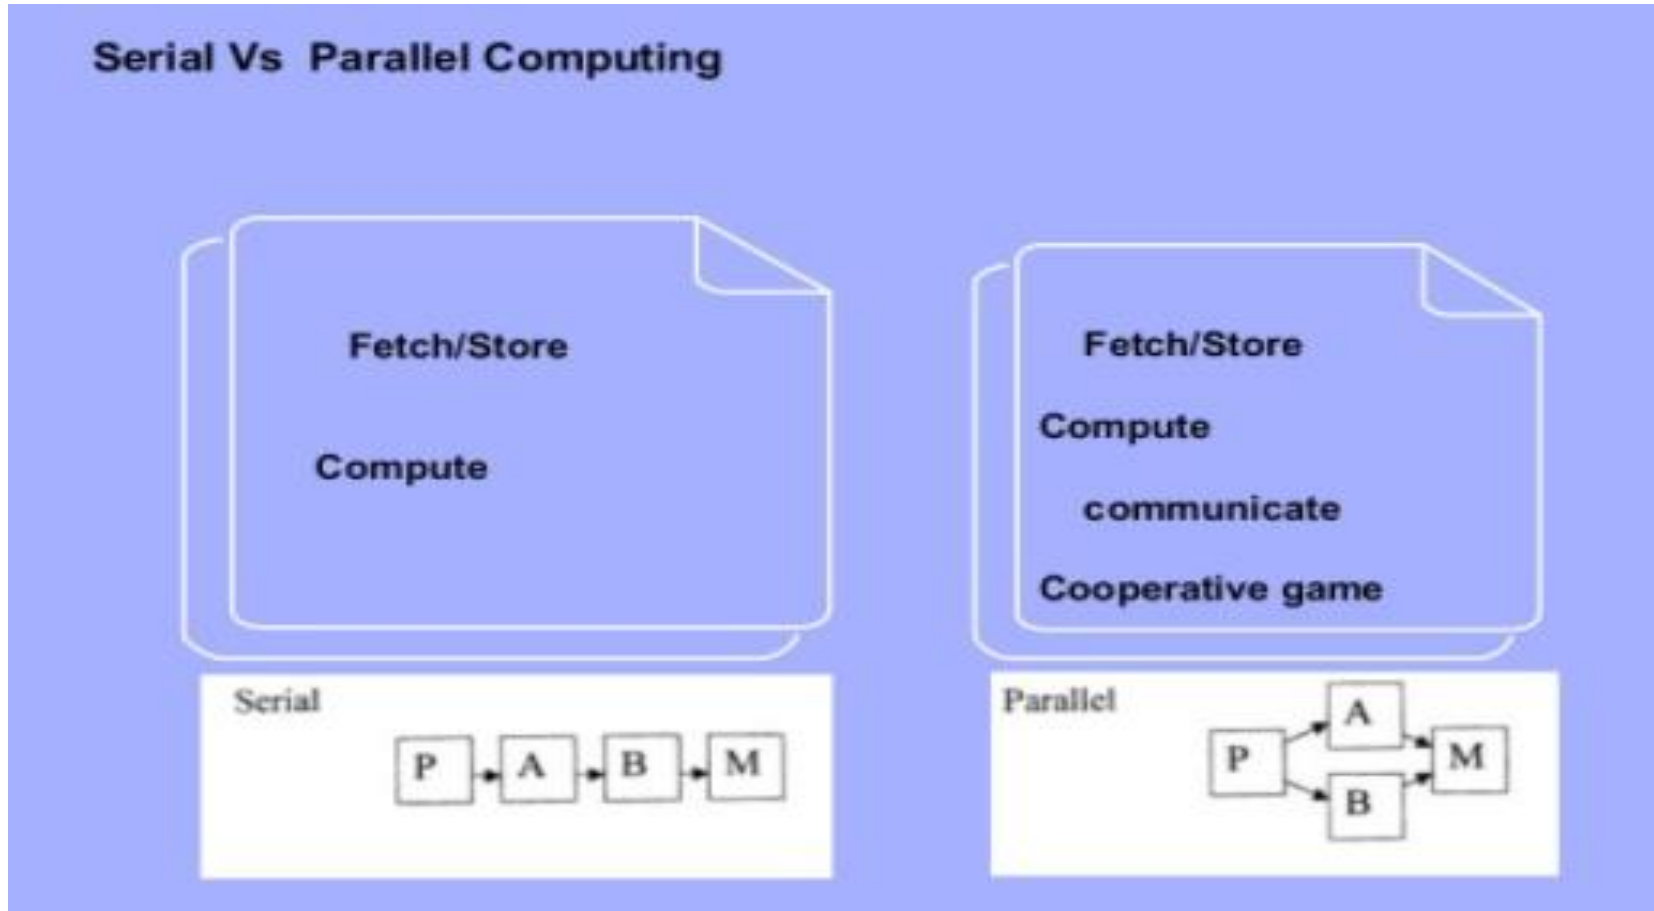
The concept of parallel computing

- One processor in your desktop or laptop can give you output in X hours
- $X > 10$
- Why not use several (Y) processors, and get the results in $X < 1$?

- Each processor works on one part of the problem
- Processors can exchange information with other processors



Serial vs Parallel Computing



Motivating Parallel Computing

The primary reasons for using parallel computing:

- Save time - wall clock time
- Solve larger problems
- Provide concurrency (do multiple things at the same time)


Motivating Parallel Computing

Computational Power argument

Moore's law states [1965]:

- Moore law is a prediction formulated by Moore in 1965 about transistor density on integrated circuit chip(ICs).\
- According to Moore's Law, “ the number of transistors used per square inch in the integrated circuits doubles about every 2 years.”
- As a result the scale gets smaller and smaller and performance of microprocessor has enjoyed an exponential growth.

Technology Growth Graph



Motivating Parallel Computing

The Memory/Disk Speed Argument

- While clock rates of high-end processors have increased at roughly 40% per year over the past decade, DRAM access times have only improved at the rate of roughly 10% per year over this interval.
- This mismatch in speeds causes significant performance bottlenecks.
- Parallel platforms provide increased bandwidth to the memory system.
- Parallel platforms also provide higher aggregate caches.
- Principles of locality of data reference and bulk access, which guide parallel algorithm design also apply to memory optimization.
- Some of the fastest growing applications of parallel computing utilize not their raw computational speed, rather their ability to pump data to memory and disk faster.


Motivating Parallel Computing

The Data Communication Argument

- As the network evolves, the vision of the Internet as one large computing platform has emerged.
- In many other applications (typically databases and data mining) the volume of data is such that they cannot be moved.
- Any analyses on this data must be performed over the network using parallel techniques.

Scope of Parallel computing

- Parallelism finds applications in very diverse application domains for different motivating reasons.
- These range from improved application performance to cost considerations.



Applications

- Commercial computing
 - Weather forecasting
 - Remote sensors, image processing
 - Process optimization, operation research
- Scientific and Engineering applications
 - Computational chemistry
 - Molecular modeling
 - Structure mechanics
- Business application
 - E-governance
 - Medical imaging
- Internet applications
 - Internet servers
 - Digital libraries

