

Topologies and Embedding

Instructor

Dr B Krishna Priya

Topologies and Embedding

- Processes as being arranged in a one-dimensional topology and uses a linear ordering to number the processes.
- In parallel programs:
- processes are naturally arranged in higher-dimensional topologies (e.g., two- or three-dimensional)
- computation and the set of interacting processes are naturally identified by their coordinates in that topology.

Contd..

- In a parallel program in which the processes are arranged in a two-dimensional topology, process (i, j) may need to send message to (or receive message from) process (k, l) .
- An MPI process with rank rank corresponds to process (row, col) in the grid such that $\text{row} = \text{rank}/4$ and $\text{col} = \text{rank}\%4$

Contd..

Figure 6.5. Different ways to map a set of processes to a two-dimensional grid. (a) and (b) show a row- and column-wise mapping of these processes, (c) shows a mapping that follows a space-filling curve (dotted line), and (d) shows a mapping in which neighboring processes are directly connected in a hypercube.

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

(a) Row-major mapping

0	4	8	12
1	5	9	13
2	6	10	14
3	7	11	15

(b) Column-major mapping

0	3	4	5
1	2	7	6
14	13	8	9
15	12	11	10

(c) Space-filling curve mapping

0	1	3	2
4	5	7	6
12	13	15	14
8	9	11	10

(d) Hypercube mapping

Contd..

- MPI provides a set of routines that allows the programmer to arrange the processes in different topologies without having to explicitly specify how these processes are mapped onto the processors.
- It is up to the MPI library to find the most appropriate mapping that reduces the cost of sending and receiving messages.

Creating and Using Cartesian Topologies

- Graphs of processes can be used to specify any desired topology. However, most commonly used topologies in message-passing programs are one-, two-, or higher-dimensional grids, that are also referred to as Cartesian topologies

Contd..

- MPI's function for describing Cartesian topologies is called `MPI_Cart_create` .
- Its calling sequence is as follows.

```
int MPI_Cart_create(MPI_Comm comm_old, int ndims,  
int *dims, int *periods, int reorder, MPI_Comm  
*comm_cart)
```

- This function takes the group of processes that belong to the communicator `comm_old` and creates a virtual process topology.
- The topology information is attached to a new communicator `comm_cart` that is created by `MPI_Cart_create` .

Contd..

- The shape and properties of the topology are specified by the arguments `ndims` , `dims` , and `periods`.
- `ndims`: specifies the number of dimensions of the topology.
- `dims`: specify the size along each dimension of the topology.
- `periods`:specify whether or not the topology has wraparound connections.
- `periods[i]` is true (non-zero in C), then the topology has wraparound connections along dimension `i` , otherwise it does not.

Contd.

- reorder:if the processes in the new group (i.e., communicator) are to be reordered or not.
- If reorder is false, then the rank of each process in the new group is identical to its rank in the old group.
- Total number of processes specified in the dims array is smaller than the number of processes in the communicator specified by comm_old , then some processes will not be part of the Cartesian topology. For this set of processes, the value of comm_cart will be set to MPI_COMM_NULL (an MPI defined constant).
- If it is greater, it will leads to error.

Process naming

- MPI provides two functions, `MPI_Cart_rank` and `MPI_Cart_coord`, for performing coordinate-to-rank and rank-to-coordinate translations, respectively.
- The calling sequences of these routines are the following:
- `int MPI_Cart_rank(MPI_Comm comm_cart, int *coords, int *rank)`
- `int MPI_Cart_coord(MPI_Comm comm_cart, int rank, int maxdims, int *coords)`
- The `MPI_Cart_rank` takes the coordinates of the process as argument in the `coords` array and returns its rank in `rank`.

- The `MPI_Cart_coords` takes the rank of the process and returns its Cartesian coordinates in the array `coords` , of length `maxdims`.
- Note:`maxdims` should be at least as large as the number of dimensions in the Cartesian topology specified by the communicator `comm_cart` .

- The communication performed among processes in a Cartesian topology is that of shifting data along a dimension of the topology.
- `int MPI_Cart_shift(MPI_Comm comm_cart, int dir, int s_step, int *rank_source, int *rank_dest)`
- `dir`: direction of the shift
- `s_step`: size of the shift step
- The computed ranks are returned in `rank_source` and `rank_dest`.