

What is OpenMP?


- **Parallel Computing** gives you more performance to throw at your problems.
 - Parallel computing is when a program uses concurrency to either:
 - Decrease the runtime for the solution to a problem.
 - Increase the size of the problem that can be solved
- **OpenMP** provides a **standard** for **shared memory** programming for scientific applications.
 - Has specific **support for scientific application** needs (*unlike Pthreads*).
 - Rapidly gaining **acceptance** among vendors and developers.
 - See <http://www.openmp.org> for more info.

OpenMP API Overview

- OpenMP - portable **shared memory** parallelism
 - An API for Writing Multithreaded Applications
 - API is a **set** of:
 - 1) **compiler directives** inserted in the source program
 - in addition to some 2) **library functions** and 3) **environment variables**
- OpenMP: **Programming Model**
 - **Fork-Join** Parallelism:
 - Master thread **spawns/generates** a team of threads as needed.
 - Parallelism is added incrementally:
 - i.e. the **sequential** program **evolves** into a **parallel** program

API Semantics

- **Master** thread executes sequential code.
- **Master** and **slaves** execute parallel code.
- **Note:**
 - very similar to **fork-join** semantics of *Pthreads* create/join primitives.



How is OpenMP typically used?

- OpenMP is usually used to parallelize loops:
 - Find your most **time-consuming loops**.
 - Split them up between threads

**Split-up this
loop between
multiple threads**

```
void main()  
{  
    double Res[1000];  
  
    for( int i=0;i<1000;i++ ) {  
        do_huge_comp(Res[i]);  
    }  
}
```

Sequential Program

```
void main()  
{  
    double Res[1000];  
    #pragma omp parallel for  
    for( int i=0;i<1000;i++ ) {  
        do_huge_comp(Res[i]);  
    }  
}
```

Parallel Program

OpenMP: How do threads interact?

- **OpenMP** is a **shared memory** model.
 - Threads communicate by sharing variables.
- **Unintended** sharing of data **can lead to race conditions**:
 - **race condition**: when the *program's outcome* changes as the threads are scheduled differently.
- To control race conditions:
 - Use **synchronization** to protect data conflicts.
- **Synchronization** is **expensive** so:
 - Often, we intend to change/control how data is stored to minimize the need for synchronization.

OpenMP Directives

- OpenMP implementation
 - Compiler directives, Library, and Environment, Unlike Pthreads (*purely a library*).
- **Parallelization** directives:
 - `parallel region`
 - `parallel for`
- **Data environment** directives:
 - `shared`, `private`, `threadprivate`, `reduction`, etc.
- **Synchronization** directives:
 - `barrier`, `critical`, `atomic`
- General Rules about Directives
 - They always apply to the ***next statement*** (or block of statements), which must be a structured block.

```
#pragma omp ...  
    Statement  
#pragma omp ...  
{  
    statement1;  
    statement2;  
    statement3;  
}
```

OpenMP: Contents

- OpenMP's **constructs** fall into **5 categories**:



- **Parallel Regions**

- Worksharing

- Data Environment

- Synchronization

- Runtime functions/environment variables

- Some Advanced Features

OpenMP Parallel Region

```
#pragma omp parallel
```

- A number of threads are **spawned/created** at the **entry**.
- Each thread executes the **same code** (SPMD model).
- The master thread ***waits all other threads at the end***. (join)
- Very similar to a number of **create/join's** with the **same function** in *Pthreads*.

OpenMP: Parallel Regions

- You create threads in OpenMP with the “**omp parallel**” pragma.
- For example, to create a 4 thread Parallel Region:

Each thread
redundantly
executes the
code within
the structured
block

```
double A[1000];  
omp_set_num_threads(4);  
#pragma omp parallel  
{  
    int ID = omp_thread_num() ;  
    foo(ID,A) ;  
}
```

Each thread calls **foo(ID,A)** for **ID = 0 to 3**

OpenMP: Parallel Regions

- Each thread executes the **same code redundantly**.

Master Thread

double A[1000];

omp_set_num_threads(4)

a single copy of **A** is shared between all threads.

foo(0,A)

foo(1,A)

foo(2,A)

foo(3,A)

printf("all done\n");

```
double A[1000];  
omp_set_num_threads(4);  
#pragma omp parallel  
{  
    int ID = omp_get_thread_num();  
    foo(ID, A);  
}  
printf("all done\n");
```

Threads wait here for **all threads** to finish before proceeding (i.e. a **barrier**)


#pragma omp parallel [clause list] Typical clauses in [clause list]

- Conditional parallelization
 - if (scalar expression)
 - Determine whether the parallel construct creates threads
- Degree of concurrency
 - num_threads (integer expression)
- number of threads to create
- Data Scoping
 - private (variable list) • Specifies variables local to each thread
 - firstprivate (variable list) • Similar to the private • Private variables are initialized to variable value before the parallel directive
 - shared (variable list) • Specifies variables that are shared among all the threads
 - default (data scoping specifier) • Default data scoping specifier may be shared or none

Example: `#pragma omp parallel if (is_parallel == 1) num_threads(8)
shared (var_b) private (var_a) firstprivate (var_c) default (none) { /*
structured block */ }`

- `if (is_parallel == 1) num_threads(8)` – If the value of the variable `is_parallel` is one, create 8 threads
- `shared (var_b)` – Each thread shares a single copy of variable `b`
- `private (var_a) firstprivate (var_c)` – Each thread gets private copies of variable `var_a` and `var_c` – Each private copy of `var_c` is initialized with the value of `var_c` in main thread when the parallel directive is encountered
- `default (none)` – Default state of a variable is specified as none (rather than shared) – Signals error if not all variables are specified as shared or private

OpenMP: Contents

- OpenMP's **constructs** fall into 5 categories:
 - Parallel Regions
 -  – **Worksharing**
 - Data Environment
 - Synchronization
 - Runtime functions/environment variables
 - Some Advanced Features

OpenMP: Work-Sharing Constructs

- The “**for**” Work-Sharing construct splits up loop iterations among the threads in a team

```
#pragma omp parallel
#pragma omp for
for ( I=0; I<N; I++ )
{
    NEAT_STUFF(I) ;
}
```

```
#pragma omp parallel for
for ( I=0; I<N; I++ )
{
    NEAT_STUFF(I) ;
}
```

... OR ...

By default, there is a **barrier** at the end of the “**omp for**”.

Use the “**nowait**” clause to **turn off** (*disable it*) the barrier.

Work Sharing Constructs

A motivating example

Sequential code

```
for(i=0;i<N;i++) { a[i]=a[i] + b[i]; }
```

OpenMP parallel region

```
#pragma omp parallel
{
    int id, i, Nthrds, istart, iend;
    id = omp_get_thread_num();
    Nthrds = omp_get_num_threads();
    istart = id * N / Nthrds;
    iend = (id+1) * N / Nthrds;
    for(i=istart; i<iend; i++){
        a[i] = a[i] + b[i];
    }
}
```



OpenMP parallel region
and a work-sharing
for-construct

```
#pragma omp parallel
#pragma omp for schedule(static)
for(i=0;i<N;i++) { a[i] = a[i] + b[i]; }
```

OpenMP Directive: `parallel for`

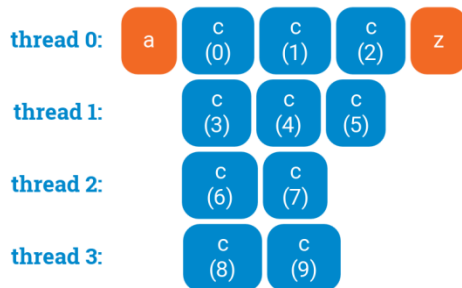
```
a();  
  
for(int i=0; i<10; ++i)  
{  
    c(i);  
}  
  
z();
```

```
omp_set_num_threads(4);  
a();  
#pragma omp parallel  
for(int i = 0; i < 8; ++i)  
{  
    c(i);  
}  
z();
```




OpenMP Directive: `parallel for`

```
omp_set_num_threads(4);  
a();  
#pragma omp parallel  
{  
    #pragma omp for  
    for( int i= 0; i<10; ++i)  
    {  
        c(i);  
    }  
}  
z();
```



```
omp_set_num_threads(4);  
a();  
  
#pragma omp parallel for  
for( int i=0; i<10; ++i)  
{  
    c(i);  
}  
  
z();
```



It is just a shorthand

Work Sharing Directives

- Always occur **within** a parallel region directive
- Two principal:
 - parallel for
 - parallel section

OpenMP Parallel For

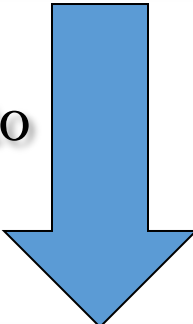
- Each thread executes a subset of the iterations
- All threads wait at the end of the **parallel for**

```
#pragma omp parallel
#pragma omp for
for( ... ) { ... }
```

Note that: Same for
parallel sections

```
#pragma omp parallel
{
    #pragma omp sections
    {
        .....
        #pragma omp section
        { ..... }
        #pragma omp section
        { ..... }
        .....
    }
}
```

is equivalent to



```
#pragma omp parallel for
for ( ... ) { ... }
```

Example: Matrix Multiply

Sequential Approach

```
for( i=0; i<n; i++ )  
    for( j=0; j<n; j++ ) {  
        c[i][j] = 0.0;  
        for( k=0; k<n; k++ )  
            c[i][j] += a[i][k] * b[k][j];  
    }
```

OpenMP Based Parallel Approach

```
#pragma omp parallel for  
for( i=0; i<n; i++ )  
    for( j=0; j<n; j++ ) {  
        c[i][j] = 0.0;  
        for( k=0; k<n; k++ )  
            c[i][j] += a[i][k] * b[k][j];  
    }
```

OpenMP Directive: parallel for

Common mistakes in the use `omp parallel` or `omp for`

```
a();  
#pragma omp for  
for (int i=0; i<8; ++i)  
{  
    c(i);  
}  
z();
```

thread 0: a c(0) c(1) c(2) c(3) c(4) c(5) c(6) c(7) z

thread 1:

thread 2:

thread 3:

```
a();  
#pragma omp parallel  
{  
    b();  
    #pragma omp for  
    for (int i=0; i<10; ++i) {  
        c(i);  
    }  
    d();  
}  
z();
```

thread 0: a b c(0) c(1) c(2) d z

thread 1: b c(3) c(4) c(5) d

thread 2: b c(6) c(7) d

thread 3: b c(8) c(9) d

OpenMP **parallel for**
Waiting / No Waiting

Multiple Work Sharing Directives

May occur within a **single parallel region**
All threads **wait at the end** of the **first for**.

```
#pragma omp parallel
{
    .....
    #pragma omp for
    for( ; ; ) { ... }
    .....
    #pragma omp for
    for( ; ; ) { ... }
    .....
}
```

The **nowait** Qualifier



Threads proceed to second **for** **without waiting**.

```
#pragma omp parallel
{
    #pragma omp for nowait
    for( ; ; ) { ... }
    #pragma omp for
    for( ; ; ) { ... }
}
```

Note the Difference between ...

```
#pragma omp parallel
{
    #pragma omp for
    for( ; ; ) { ... }
    foo();
    #pragma omp for
    for( ; ; ) { ... }
}
```


... and ...

```
#pragma omp parallel for
for( ; ; ) { ... }
foo();
#pragma omp parallel for
for( ; ; ) { ... }
```

OpenMP parallel for loops: waiting

- In a **parallel** region, OpenMP will **automatically wait** for all threads to finish before execution continues.
- There is also a **synchronization point** after each **omp for** loop;
 - here no thread will execute **d()** until all threads are done with the loop


```
a();  
#pragma omp parallel  
{  
    b();  
    #pragma omp for  
    for (int i=0; i<10; ++i)  
    {  
        c(i);  
    }  
    d();  
}  
z();
```



OpenMP parallel for loops: waiting

- However, if you do not need synchronization after the loop, you can **disable** it with **nowait**:

```
a();  
#pragma omp parallel  
{  
    b();  
    #pragma omp for nowait  
    for (int i=0; i<10; ++i)  
    {  
        c(i);  
    }  
    d();  
}  
z();
```




OpenMP parallel for loops:

Interaction with **critical** sections

If you need a critical section after a loop, note that normally OpenMP will first wait for all threads to finish their loop iterations before letting any of the threads to enter a critical section:

```
a();  
#pragma omp parallel  
{  
    b();  
    #pragma omp for  
    for (int i=0; i<10; ++i)  
    {  
        c(i);  
    }  
    #pragma omp critical  
    {  
        d();  
    }  
}  
z();
```




OpenMP parallel for loops:

Interaction with critical sections

You can disable **waiting**, so that some threads can start doing post-processing early. This would make sense if, e.g., **d()** updates some global data structure based on what the thread computed in its own part of the parallel for loop:

```
a();  
#pragma omp parallel  
{  
    b();  
    #pragma omp for nowait  
    for (int i=0; i<10; ++i)  
    {  
        c(i);  
    }  
    #pragma omp critical  
    {  
        d();  
    }  
}  
z();
```

Notice




OpenMP parallel for loops:

No waiting before a loop

Now, note that there is **no** synchronization point before the loop starts.

If threads reach the for loop at different times, they can start their own part of the work as soon as they are there, without waiting for the other threads:

```
a();  
#pragma omp parallel  
{  
    #pragma omp critical  
    {  
        b();  
    }  
    #pragma omp for  
    for (int i=0; i<10; ++i)  
    {  
        c(i);  
    }  
    d();  
}  
z();
```



References

- OpenMP topic: Loop parallelism
 - <https://pages.tacc.utexas.edu/~eijkhout/pcse/html/omp-loop.html>
- A “Hands-on” Introduction to OpenMP
 - <https://www.openmp.org/wp-content/uploads/omp-hands-on-SC08.pdf>
- OpenMP in a nutshell
 - <http://www.bowdoin.edu/~ltoma/teaching/cs3225-GIS/fall17/Lectures/openmp.html>
- #pragma omp parallel (IBM)
 - https://www.ibm.com/support/knowledgecenter/SSGH3R_13.1.3/com.ibm.xlcpp1313.aix.doc/compiler_ref/prag_omp_parallel.html
- OpenMP Directives (Microsoft)
 - <https://docs.microsoft.com/en-us/cpp/parallel/openmp/reference/openmp-directives?view=vs-2019>
- Guide into OpenMP: Easy multithreading programming for C++
 - <https://bisqwit.iki.fi/story/howto/openmp/>